885. VIRGO is an operating system kernel forked off from Linux kernel mainline to add cloud functionalities (system calls,
modules etc.,) within kernel itself with machine learning, analytics, debugging, queueing support in the deepest layer of
OSI stack i.e AsFer, USBmd, KingCobra together with VIRGO constitute the previous functionalities. Presently there seems to
be no cloud implementation with fine-grained cloud primitives (system calls, modules etc.,) included in kernel itself
though there are coarse grained clustering and SunRPC implementations available. VIRGO complements other Clustering and
application layer cloud OSes like cloudstack, openstack etc., in these aspects - CloudStack and OpenStack can be deployed
on a VIRGO Linux Kernel Cloud - OpenStack nova compute, neutron network, cinder/swift storage subsystems can be augmented
to have additional drivers that invoke lowlevel VIRGO syscall and kernel module primitives (assuming there are no
coincidental replications of functionalities) thereby acting as a foundation to application layer cloud.

886. Remote Device Invocation , which is an old terminlogy for Internet-Of-Things has already been experimented in SunRPC
and KOrbit CORBA-in-linux-kernel kernel modules in old linux kernels (http://www.csn.ul.ie/~mark/fyp/fypfinal.html - with
Solaris MC and example Remote Device Client-Server Module implementation). VIRGO Linux with the larger encompassing
NeuronRain suite is an effort to provide a unified end-to-end application-to-kernel machine-learning propelled cloud and
internet-of-things framework.

887. Memory pooling:
---------------
Memory pooling is proposed to be implemented by a new virgo_malloc() system call that transparently allocates a block of
virtual memory from memory pooled from virtual memory scattered across individual machines part of the cloud.

888. CPU pooling or cloud ability in a system call:
-----------------------------------------------

Clone() system call is linux specific and internally it invokes sys_clone(). All fork(),vfork() and clone() system calls internally invoke do_fork(). A new system call virgo_clone() is proposed to create a thread transparently on any of the available machines on the cloud.This creates a thread on a free or least-loaded machine on the cloud and returns the results.

virgo_clone() is a wrapper over clone() that looks up a map of machines-to-loadfactor and get the host with least load and invokes clone() on a function on that gets executed on the host. Usual cloud implementations provide userspace API that have something similar to this - call(function,host). Loadfactor can be calculated through any of the prominent loadbalancing algorithm. Any example userspace code that uses clone() can be replaced with virgo_clone() and all such threads will be running in a cloud transparently.Presently Native POSIX threads library(NPTL) and older LinuxThreads thread libraries internally use clone().

Kernel has support for kernel space sockets with kernel_accept(), kernel_bind(), kernel_connect(), kernel_sendmsg() and kernel_recvmsg() that can be used inside a kernel module. Virgo driver implements virgo_clone() system call that does a kernel_connect() to a remote kernel socket already __sock_create()-d, kernel_bind()-ed and kernel_accept()-ed and does kernel_sendmsg() of the function details and kernel_recvmsg() after function has been executed by clone() in remote machine. After kernel_accept() receives a connection it reads the function and parameter details. Using these kthread_create() is executed in the remote machine and results are written back to the originating machine. This is somewhat similar to SunRPC but adapted and made lightweight to suit virgo_clone() implementation without any external data representation.

Experimental Prototype
----------------------
virgo_clone() system call and a kernel module virgocloudexec which implements Sun RPC interface have been implemented.

VIRGO - loadbalancer to get the host:ip of the least loaded node
------------------------------------------------------------------
889. Loadbalancer option 1 - Centralized loadbalancer registry that tracks load:
------------------------------------------------------------------------------

Virgo_clone() system call needs to lookup a registry or map of host-to-load and get the least loaded host:ip from it. This requires a  load monitoring code to run periodically and update the map. If this registry is located on a single machine then simultaneous virgo_clone() calls from many machines on the cloud could choke the registry. Due to this, loadbalancer registry needs to run on a high-end machine. Alternatively,each machine can have its own view of the load and multiple copies of load-to-host registries can be stored in individual machines. Synchronization of the copies becomes a separate task in itself(Cache coherency). Either way gives a tradeoff between accuracy, latency and efficiency.

Many application level userspace load monitoring tools are available but as virgo_clone() is in kernel space, it needs to be investigated if kernel-to-kernel loadmonitoring can be done without userspace data transport.Most Cloud API explicitly invoke a function on a host. If this functionality is needed, virgo_clone() needs to take host:ip address as extra argument,but it reduces transparent execution.

(Design notes for LB option 1 handwritten by myself are at :http://sourceforge.net/p/virgo-linux/code-0/HEAD/tree/trunk/virgo-docs/MiscellaneousOpenSourceDesignAndAcademicResearchNotes.pdf)

890. Loadbalancer option 2 - Linux Psuedorandom number generator based load balancer(experimental) instead of centralized registry that tracks load:
---------------------------------------------------------------------------------------

Each virgo_clone() client has a PRG which is queried (/dev/random or /dev/urandom) to get the id of the host to send the next virgo_clone() function to be executed
Expected number of requests per node is derived as:

expected number of requests per node = summation(each_value_for_the_random_variable_for_number_of_requests * probability_for_each_value) where random variable ranges from 1 to k where N is number of processors and k is the number of requests to be distributed on N nodes

=expected number of requests per node = (math.pow(N, k+2) - k*math.pow(N,2) + k*math.pow(N,1) - 1) / (math.pow(N, k+3) - 2*math.pow(N,k+2) + math.pow(N,k+1))

This loadbalancer is dependent on efficacy of the PRG and since each request is uniformly, identically, independently distributed use of PRG
would distribute requests evenly. This obviates the need for loadtracking and coherency of the load-to-host table.

(Design notes for LB option 2 handwritten by myself at :http://sourceforge.net/p/virgo-linux/code-0/HEAD/tree/trunk/virgo-docs/MiscellaneousOpenSourceDesignAndAcademicResearchNotes.pdf)


(python script in virgo-python-src/)

*********************************************************************************************
891. Implemented VIRGO Linux components (as on 7 March 2016)
*********************************************************************************************
1. cpupooling virtualization - VIRGO_clone() system call and VIRGO cpupooling driver by which a remote procedure can be invoked in kernelspace.(port: 10000)
2. memorypooling virtualization - VIRGO_malloc(), VIRGO_get(), VIRGO_set(), VIRGO_free() system calls and VIRGO memorypooling driver by which kernel memory can be allocated in remote node, written to, read and freed - A kernelspace memcache-ing.(port: 30000)
3. filesystem virtualization - VIRGO_open(), VIRGO_read(), VIRGO_write(), VIRGO_close() system calls and VIRGO cloud filesystem driver by which file IO in remote node can be done in kernelspace.(port: 50000)
4. config - VIRGO config driver for configuration symbols export.
5. queueing - VIRGO Queuing driver kernel service for queuing incoming requests, handle them with workqueue and invoke KingCobra service routines in kernelspace. (port: 60000)
6. cloudsync - kernel module for synchronization primitives (Bakery algorithm etc.,) with exported symbols that can be used in other VIRGO cloud modules for critical section lock() and unlock()
7. utils - utility driver that exports miscellaneous kernel functions that can be used across VIRGO Linux kernel
8. EventNet - eventnet kernel driver to vfs_read()/vfs_write() text files for EventNet vertex and edge messages (port: 20000)

9. Kernel_Analytics - kernel module that reads machine-learnt config key-value pairs set in /etc/virgo_kernel_analytics.conf. Any machine learning software can be used to get the key-value pairs for the config. This merges three facets - Machine Learning, Cloud Modules in VIRGO Linux-KingCobra-USBmd , Mainline Linux Kernel
10. Testcases and kern.log testlogs for the above
11. SATURN program analysis wrapper driver.

Thus VIRGO Linux at present implements a minimum cloud OS (with cloud-wide cpu, memory and file system management) over Linux and potentially fills in a gap to integrate both software and hardware into cloud with machine learning and analytics abilities that is absent in application layer cloud implementations. Thus VIRGO cloud is an IoT operating system kernel too that enables any hardware to be remote controlled. Data analytics using AsFer can be done by just invoking requisite code from a kernelspace driver above and creating an updated driver binary (or) by kernel_analytics module which reads the userland machine-learnt config.

*************************************************************************************************
****
VIRGO Features (list is quite dynamic and might be rewritten depending on feasibility - longterm with no deadline)
*************************************************************************************************
****
892. (FEATURE - DONE-minimum separate config file support in client and kernel service )1. More Sophisticated VIRGO config file and read_virgo_config() has to be invoked on syscall clients virgo_clone and virgo_malloc also. Earlier config was being read by kernel module only which would work only on a single machine. A separate config module kernel service has been added for future use while exporting kernel-wide configuration related symbols. VIRGO config files have been split into /etc/virgo_client.conf and /etc/virgo_cloud.conf to delink the cloud client and kernel service config parameters reading and to do away with oft occurring symbol lookup errors and multiple definition errors for num_cloud_nodes and node_ip_addrs_in_cloud - these errors are frequent in 3.15.5 kernel than 3.7.8 kernel. Each VIRGO module and system call now reads the config file independent of others - there is a read_virgo_config_<module>_<client_or_service>() function variant for each driver and system call. Though at present smacks of a replicated code, in future the config reads for each component (system call or module) might vary significantly depending on necessities.  New kernel module config has been added in drivers/virgo. This is for future prospective use as a config export driver that can be looked up by any other VIRGO module for config parameters.  include/linux/virgo_config.h has the declarations for all the config variables declared within each of the VIRGO kernel modules.  Config variables in each driver and system call have been named with prefix and suffix to differentiate the module and/or system call it serves.  In geographically distributed cloud virgo_client.conf has to be in client nodes and virgo_cloud.conf has to be in cloud nodes. For VIRGO Queue - KingCobra REQUEST-REPLY peer-to-peer messaging system same node can have virgo_client.conf and virgo_cloud.conf.  Above segregation largely simplifies the build process as each module and system call is independently built without need for a symbol to be exported from other module by pre-loading it.(- from commit comments done few months ago)


893. (FEATURE - Special case implementation DONE) 2. Object Marshalling and Unmarshalling (Serialization) Features - Feature 4 is a marshalling feature too as Python world PyObjects are serialized into VIRGO linux kernel and unmarshalled back bottom-up with CPython and Boost::Python C++ invocations - CPython and Boost internally take care of serialization.

894. (FEATURE - DONE) Virgo_malloc(), virgo_set(), virgo_get() and virgo_free() syscalls that virtualize the physical memory across all cloud nodes into a single logical memory behemoth (NUMA visavis UMA). (There are random crashes in

copy_to_user and copy_from_user in syscall path for VIRGO memory pooling commands that were investigated but turned out to be mystery). These crashes have either been resolved or occur less in 3.15.5 and 4.1.5 kernels.
Initial Design Handwritten notes committed at: http://sourceforge.net/p/virgo-linux/code-0/210/tree/trunk/virgo-docs/VIRGO_Memory_Pooling_virgomalloc_initial_design_notes.pdf

895. (FEATURE - DONE) Integrated testing of AsFer-VIRGO Linux Kernel request roundtrip - invocation of VIRGO linux kernel system calls from AsFer Python via C++ or C extensions - Commits for this have been done on 29 January 2016. This unifies userlevel applications and kernelspace modules so that AsFer Python makes VIRGO linux kernel an extension. Following is schematic diagram and More details in commit notes below.

895.1 Schematic Diagram:
-----------------------
```
        AsFer Python -----> Boost::Python C++ Extension ------> VIRGO memory system calls --------> VIRGO Linux Kernel
Memory Drivers
        /\                                                                                          V
         |                                                                                          |
          ---------------------------------------------<----------------------------------------------

        AsFer Python -----> CPython Extensions ------> VIRGO memory system calls --------> VIRGO Linux Kernel Memory
Drivers
        /\                                                                                          V
         |                                                                                          |
          ---------------------------------------------<----------------------------------------------
```
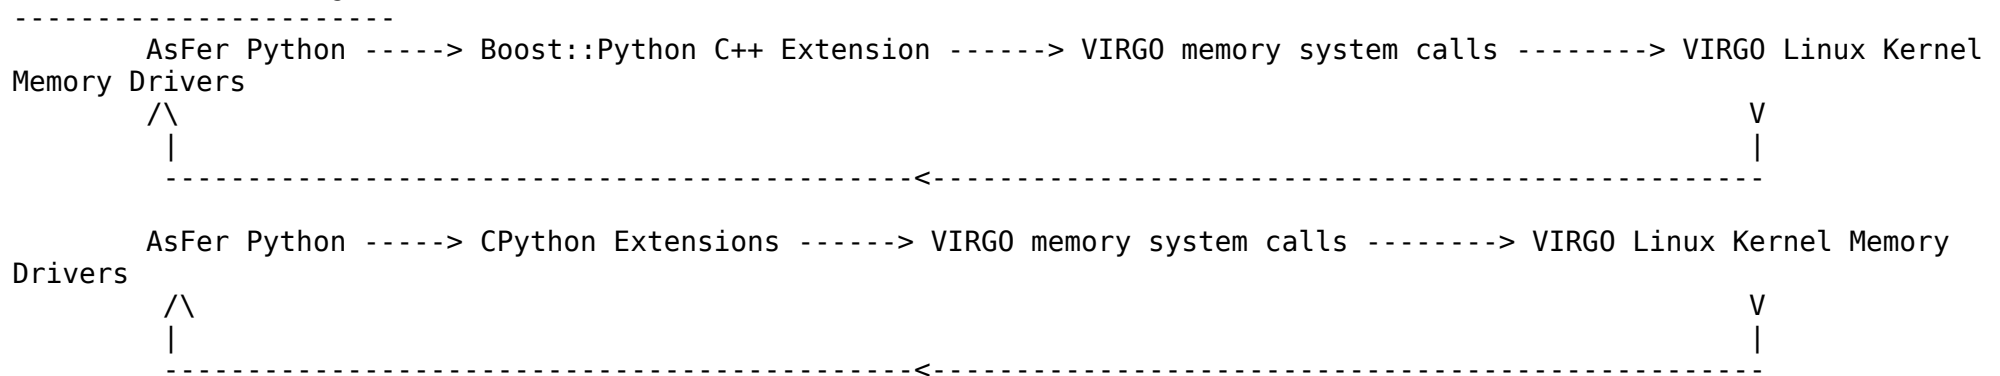
896. (FEATURE - DONE) Multithreading of VIRGO cloudexec kernel module (if not already done by kernel module subsystem internally)

897. (FEATURE - DONE) Sophisticated queuing and persistence of CPU and Memory pooling requests in Kernel Side (by possibly improving already existing kernel workqueues). Either open source implementations like ZeroMQ/ActiveMQ can be used or Queuing implementation has to be written from scratch or both. ActiveMQ supports REST APIs and is JMS implementation. This feature has been marked completed because recently NeuronRain AsFer backend has been updated to support KingCobra REQUEST_REPLY.queue as a datasource for Streaming Abstract Generator. By enabling use_as_kingcobra_service=1 in cpupooling and memorypooling VIRGO drivers, any incoming CPU and Memory related request can be routed to KingCobra by linux workqueue in VIRGO queue and disk persisted (/var/log/REQUEST_REPLY.queue) by KingCobra servicerequest recipient. Also Kafka Publisher/Subscriber have been implemented in NeuronRain AsFer which invoke Streaming Abstract Generator with KingCobra REQUEST_REPLY.queue as datasource to publish persisted already received CPU and Memory requests to Kafka Message Queue. Thus queuing and persistence for VIRGO CPU and Memory is in place. ZeroMQ does not have persistence and is used for NeuronRain client side Router-Dealer concurrent request servicing pattern.

898. (FEATURE - DONE-Minimum Functionality - this section is an extended draft on respective topics in NeuronRain AstroInfer design - https://github.com/shrinivaasanka/asfer-github-code/blob/master/asfer-docs/AstroInferDesign.txt) Integration of Asfer(AstroInfer) algorithm codes into VIRGO which would add machine learning capabilities into VIRGO - That is, VIRGO cloud subsystem which is part of a linux kernel installation "learns" and "adapts" to the processes that are exexcuted on VIRGO. This catapults the power of the Kernel and Operating System into an artificially (rather approximately

naturally) intelligent computing platform (a software "brain"). For example VIRGO can "learn" about "execution times" of processes and suitably act for future processes. PAC Learning of functions could be theoretical basis for this. Initial commits for Kernel Analytics Module which reads the /etc/virgo_kernel_analytics.conf config have been done. This config file virgo_kernel_analytics.conf having csv(s) of key-value pairs of analytics variables is set by AsFer or any other Machine Learning code. With this VIRGO Linux Kernel is endowed with abilities to dynamically evolve than being just a platform for user code. Implications are huge - for example, a config variable "MaxNetworkBandwidth=255" set by the ML miner in userspace based on a Perceptron or Logistic Regression executed on network logs can be read by a kernel module that limits the network traffic to 255Mbps. Thus kernel is no longer a static predictable blob behemoth. With this, VIRGO is an Internet-of-Things kernel that does analytics and based on analytics variable values integrated hardware can be controlled across the cloud through remote kernel module function invocation. This facility has been made dynamic with Boost::Python C++ and CPython extensions that permit flow of objects from machine learnt AsFer kernel analytics variables to VIRGO Linux Kernel memory drivers via VIRGO system calls directly and back - Commits on 29 January 2016 - this should obviate re-reading /etc/virgo_kernel_analytics.conf and is an exemplary implementation which unifies C++/Python into C/Kernel.

-------------------------------------------
Example scenario 898.1 without implementation:
-------------------------------------------
- Philips Hue IoT mobile app controlled bulb - http://www2.meethue.com/en-xx/
- kernel_analytics module learns key-value pairs from the AsFer code and exports it VIRGO kernel wide
- A driver function with in bulb embedded device driver can be invoked through VIRGO cpupooling (invoked from remote virgo_clone() system_call)
based on if-else clause of the kernel_analytics variable i.e remote_client invokes virgo_clone() with function argument "lights on" which is routed to another cloud node. The recipient cloud node "learns" from AsFer kernel_analytics that Voltage is low or Battery is low from logs and decides to switch in high beam or low beam.
-------------------------------------------
Example scenario 898.2 without implementation:
-------------------------------------------
- A swivel security camera driver is remotely invoked via virgo_clone() in the VIRGO cloud.
- The camera driver uses a machine learnt variable exported by kernel_analytics-and-AsFer to pan the camera by how much degrees.
--------------------------------------------------------------------------------------------
Example scenario 898.3 without implementation - probably one of the best applications of NeuronRain IoT OS:
--------------------------------------------------------------------------------------------
- Autonomous Driverless Automobiles - a VIRGO driver for a vehicle which learns kernel analytics variables (driving directions) set by kernel_analytics driver and AsFer Machine Learning. A naive algorithm for Driverless Car (with some added modifications over A-Star and Motion planning algorithms):
        - AsFer analytics receives obstacle distance data 360+360 degrees (vertical and horizontal) around the vehicle (e.g ultrasound sensors) which is updated in a Spark DataFrame table with high frequency (100 times per second).
        - VIRGO Linux kernel on vehicle has two special drivers for Gear-Clutch-Break-Accelerator-Fuel(GCBAF) and Steering listening on some ports.
        - AsFer analytics with high frequency computes threshold variables for applying break, clutch, gear, velocity, direction, fuel changes which are written to kernel_analytics.conf realtime based on distance data from Spark table.
        - These analytics variables are continuously read by GCBAF and Steering drivers which autopilot the vehicle.

        - Above applies to Fly-by-wire aeronautics too with appropriate changes in analytics variables computed.
        - The crucial parameter is the response time in variable computation and table updates which requires a huge
computing power unless the vehicle is hooked onto a Spark cloud in motion by wireless which process the table and compute
analytic variables.

E.g. Autopilot in Tesla Cars processes Petabytes of information (Smooth-as-Silk algorithm) from sensors which are fed to
neural networks computed on a cloud - https://www.teslarati.com/tesla-firmware-v8-1-17-22-26-autopilot-2-0-smooth-silk-
update-video/.

-------------------------------------------------
References for Machine Learning + Linux Kernel
-------------------------------------------------
898.4 KernTune - http://repository.uwc.ac.za/xmlui/bitstream/handle/10566/53/Yi_KernTune(2007).pdf?sequence=3
898.5 Self-learning, Predictive Systems - https://icri-ci.technion.ac.il/projects/past-projects/machine-learning-for-
architecture-self-learning-predictive-computer-systems/
898.6 Linux Process Scheduling and Machine Learning - http://www.cs.ucr.edu/~kishore/papers/tencon.pdf
898.7 Network Latency and Machine Learning - https://users.soe.ucsc.edu/~slukin/rtt_paper.pdf
898.8 Machine Learning based Meta-Scheduler for Multicore processors - https://books.google.co.in/books?
id=1GWcHmCrl0QC&pg=PA528&lpg=PA528&dq=linux+kernel+machine+learning&source=bl&ots=zfJsq_uu5q&sig=mMIUZ-
oyJIwZXtYj4HntrQE8NSk&hl=en&sa=X&ved=0CCAQ6AEwATgKahUKEwjs9sqF9vPIAhVBFZQKHbNtA6A

899. A Symmetric Multi Processing subsystem Scheduler that virtualizes all nodes in cloud (probably this would involve
improving the loadbalancer into a scheduler with priority queues)

900. (FEATURE - ONGOING) Virgo is an effort to virtualize the cloud as a single machine - Here cloud is not limited to
servers and desktops but also mobile devices that run linux variants like Android, and other Mobile OSes. In the longterm,
Virgo may have to be ported or optimized for handheld devices. Boost::Python AsFer-VIRGO system call invocations
implemented in NeuronRain is framework for implementing python applications interfacing with kernel. If deployed on Mobile
processors (e.g by overlaying Android Kernel with VIRGO layer) there are IDEs like QPython to develop python apps for
Android.

901. (FEATURE - DONE) Memory Pooling Subsystem Driver - Virgo_malloc(), Virgo_set(), Virgo_get() and Virgo_free() system
calls and their Kernel Module Implementations. In addition to syscall path, telnet or userspace socket client interface is
also provided for both VIRGO CPU pooling(virgo_clone()) and VIRGO Memory Pooling Drivers.

902. (FEATURE - DONE) Virgo Cloud File System with virgo_cloud_open(), virgo_cloud_read() , virgo_cloud_write() and
virgo_cloud_close() commands invoked through telnet path has been implemented that transcends disk storage in all nodes in
the cloud. It is also fanciful feature addition that would make VIRGO a complete all-pervading cloud platform. The remote
telnet clients send the file path and the buf to be read or data to be written. The Virgo File System kernel driver service
creates a unique Virgo File Descriptor for each struct file* opened by filp_open() and is returned to client. Earlier
design option to use a hashmap (linux/hashmap.h) looked less attractive as file desciptor is an obvious unique description
for open file and also map becomes unscalable. The kernel upcall path has been implemented (paramIsExecutable=0) and may
not be necessary in most cases and all above cloudfs commands work in kernelspace using VFS calls.

903. (FEATURE - DONE) VIRGO Cloud File System commands through syscall paths - virgo_open(),virgo_close(),virgo_read() and virgo_write(). All the syscalls have been implemented with testcases and more bugs fixed. After fullbuild and testing, virgo_open() and virgo_read() work and copy_to_user() is working.

904. (FEATURE - DONE) VIRGO memory pooling feature is also a distributed key-value store similar to other prominent key-store software like BigTable implementations, Dynamo, memory caching tools etc., but with a difference that VIRGO mempool is implemented as part of Linux Kernel itself thus circumventing userspace latencies. Due to Kernel space VIRGO mempool has an added power to store and retrieve key-value pair in hardware devices directly which otherwise is difficult in userspace implementations.

905. VIRGO memory pooling can be improved with disk persistence for in-memory key-value store using virgo_malloc(),virgo_set(),virgo_get() and virgo_free() calls. Probably this might be just a set of invocations of read and write ops in disk driver or using sysfs. Probably this could be redundant as the VIRGO filesystem primitives have been implemented that write to a remote host's filesystem in kernelspace.

906. (FEATURE-DONE) Socket Debugging, Program Analysis and Verification features for user code that can find bugs statically. Socket skbuff debug utility and SATURN Program Analysis Software has been integrated into NEURONRAIN VIRGO Linux Kernel.

907. (FEATURE - DONE-Minimum Functionality) Operating System Logfile analysis using Machine Learning code in AstroInfer for finding patterns of processes execution and learn rules from the log. Kernel_Analytics VIRGO module reads /etc/virgo_kernel_analytics.conf config key-value pairs which are set by AsFer or other Machine Learning Software. At present an Apache Spark usecase that mines Uncomplicated Fire Wall logs in kern.log for most prominent source IP has been implemented in AsFer codebase : http://sourceforge.net/p/asfer/code/704/tree/python-src/SparkKernelLogMapReduceParser.py . This is set as a key-value config in /etc/virgo_kernel_analytics.conf read and exported by kernel_analytics module.

908. (USERSPACE C++ usecase implemented in GRAFIT course material - https://gitlab.com/shrinivaasanka/Grafit) Implementations of prototypical Software Transactional Memory and LockFree Datastructures for VIRGO memory pooling.

909. Scalability features for Multicore machines - references: (http://halobates.de/lk09-scalability.pdf, http://pdos.csail.mit.edu/papers/linux:osdi10.pdf)

910. (USERSPACE C++ usecase implemented in GRAFIT course material - https://gitlab.com/shrinivaasanka/Grafit) Read-Copy-Update algorithm implementation for VIRGO memory pooling that supports multiple simultaneous versions of memory for readers - widely used in redesigned Linux Kernel.

911. (FEATURE - SATURN integration - minimum functionality DONE) Program Comprehension features as an add-on described in : https://sites.google.com/site/kuja27/PhDThesisProposal.pdf?attredirects=0. SATURN program analysis has been integrated into VIRGO linux with a stub driver.

912. (FEATURE - DONE) Bakery Algorithm implementation - cloudsync kernel module

913. (FEATURE - minimal EventNet Logical Clock primitive implemented in AstroInfer and this section is an extended draft on respective topics in NeuronRain AstroInfer design -  https://github.com/shrinivaasanka/asfer-github-code/blob/master/asfer-

docs/AstroInferDesign.txt) Implementation of Distributed Systems primitives for VIRGO cloud viz., Logical Clocks, Termination Detection, Snapshots, Cache Coherency subsystem etc.,(as part of cloudsync driver module). Already a simple timestamp generation feature has been implemented for KingCobra requests with <ipaddress>:<localmachinetimestamp> format

914. (FEATURE - minimum functionality DONE - this section is an extended draft on respective packing/filling/tiling topics in NeuronRain AstroInfer design specific to kernel SLAB allocator -  https://github.com/shrinivaasanka/asfer-github-code/blob/master/asfer-docs/AstroInferDesign.txt) Enhancements to kmem if it makes sense, because it is better to rely on virgo_malloc() for per machine memory management and wrap it around with a cloudwide VIRGO Unique ID based address lookup implementation of which is already in place.
Kernel Malloc syscall kmalloc() internally works as follows:
        - kmem_cache_t object has pointers to 3 lists
        - These 3 lists are full objects SLAB list, partial objects SLAB list and free objects SLAB list - all are lists of objects of same size
 and cache_cache is the global list of all caches created thus far.
        - Any kmalloc() allocation searches partial objects SLAB list and allocates a memory block with kmem_cache_alloc() from the first SLAB available - returned to caller.
        - Any kfree() returns an object to a free SLAB list
        - Full SLABs are removed from partial SLAB list and appended to full SLAB list
        - SLABs are virtual memory pages created with kmem_cache_create
        - Each SLAB in SLABs list has blocks of similar sized objects (e.g. multiples of two). Closest matching block is returned and fragmentation is minimized (incidentally this is the knapsack and packing optimization LP problem and thus NP-complete).

KERNELSPACE:
VIRGO address translation table already implements a tree registry of vtables each of capacity 3000 that keep track of kmalloc() allocations across all cloud nodes. Implementation of SLAB allocator for kmalloc() creates a kmem_cache(s) of similar sized objects and kmem_cache_alloc() allocates from these caches. kmalloc() already does lot of per-machine optimizations. VIRGO vtable registry tree maintained in VIRGO memory syscall end combined with per-machine kmalloc() cache_cache already look sufficient. Instrumenting kmem_cache_create() with #ifdef SLAB_CLOUD_MALLOC flags to do RPC looks superfluous. Hence marking this action item as done. Any further optimization can be done on top of existing VIRGO address translation table struct - e.g bookkeeping flags, function pointers etc.,.
USERSPACE: sbrk() and brk() are no longer used internally in malloc() library routines. Instead mmap() has replaced it (http://web.eecs.utk.edu/courses/spring2012/cs360/360/notes/Malloc1/lecture.html, http://web.eecs.utk.edu/courses/spring2012/cs360/360/notes/Malloc1/diff.html).

915. (FEATURE - ONGOING) Cleanup the code and remove unnecessary comments.

916. (FEATURE - DONE) Documentation - This design document is also a documentation for commit notes and other build and debugging technical details. Doxygen html cross-reference documentation for AsFer, USBmd, VIRGO, KingCobra and Acadpdrafts has been created along with summed-up design document and committed to GitHub Repository at https://github.com/shrinivaasanka/Krishna_iResearch_DoxygenDocs

917. (FEATURE - DONE) Telnet path to virgo_cloud_malloc,virgo_cloud_set and virgo_cloud_get has been tested and working. This is similar to memcached but stores key-value in kernelspace (and hence has the ability to write to and retrieve from

any device driver memory viz., cards, handheld devices).An optional todo is to write a script or userspace socket client
that connects to VIRGO mempool driver for these commands.

918. Augment the Linux kernel workqueue implementation (http://lxr.free-electrons.com/source/kernel/workqueue.c) with disk
persistence if feasible and doesn't break other subsystems - this might require additional persistence flags in work_struct
and additional #ifdefs in most of the queue functions that write and read from the disk. Related to item 6 above.

919. (FEATURE - DONE) VIRGO queue driver with native userspace queue and kernel workqueue-handler framework that is
optionally used for KingCobra and is invoked through VIRGO cpupooling and memorypooling drivers. (Schematic in
http://sourceforge.net/p/kcobra/code-svn/HEAD/tree/KingCobraDesignNotes.txt and
http://sourceforge.net/p/acadpdrafts/code/ci/master/tree/Krishna_iResearch_opensourceproducts_archdiagram.pdf)

920. (FEATURE - DONE) KERNELSPACE EXECUTION ACROSS CLOUD NODES which geographically distribute userspace and kernelspace
execution creating
a logical abstraction for a cloudwide virtualized kernel:

```
        Remote Cloud Node Client
        (cpupooling, eventnet, memorypooling, cloudfs, queueing - telnet and syscalls clients)
                 |
                 |
  (Userspace)    |
                 |
                 |--------------------------------Kernel Sockets------------------------------------> Remote Cloud
  Node Service
                                                     (VIRGO cpupooling, memorypooling, cloudfs, queue,
  KingCobra drivers)
                                                                                                      |
                                                                                                      |
                                                                                                      |
                                                                                                      |
  (Kernelspace execution)
                                                                                                      |
                                                                                                      V
        <------------------------------------Kernel Sockets-------------------------------------------
                 |
                 |
                 |
                 |
  (Userspace)    |
```
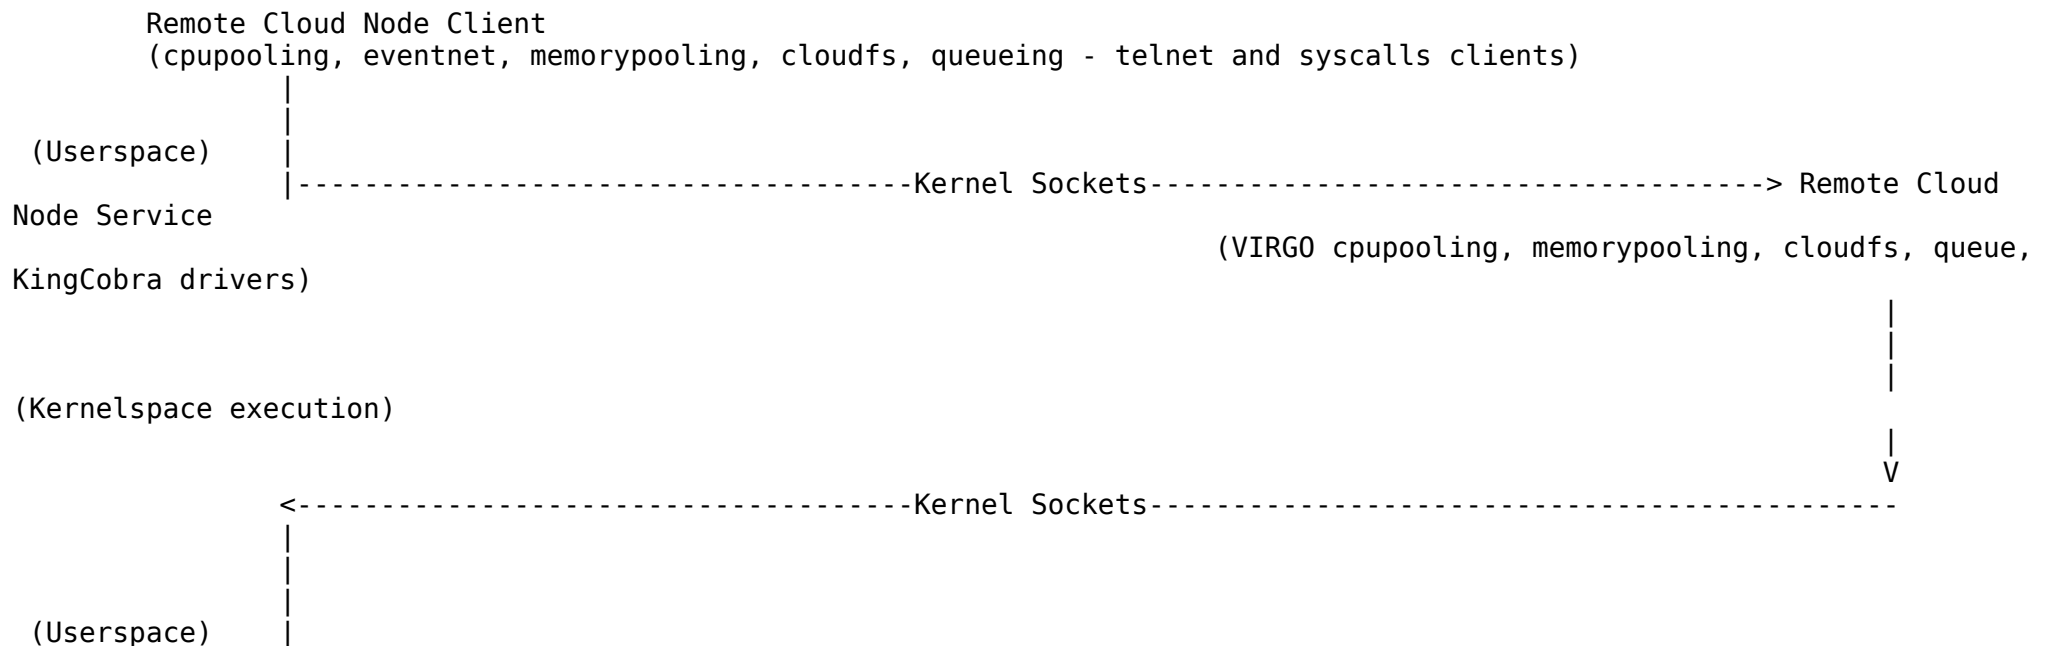
921. (FEATURE - DONE) VIRGO platform as on 5 May 2014 implements a minimum set of features and kernelsocket commands
required for a cloud OS kernel: CPU virtualization(virgo_clone), Memory
virtualization(virgo_malloc,virgo_get,virgo_set,virgo_free) and a distributed cloud file
system(virgo_open,virgo_close,virgo_read,virgo_write) on the cloud nodes and thus gives a logical view of one unified,
distributed linux kernel across all cloud nodes that splits userspace and kernelspace execution across cloud as above.

922. (FEATURE - DONE) VIRGO Queue standalone kernel service has been implemented in addition to paths in schematics above.
VIRGO Queue listens on hardcoded port 60000 and enqueues the incoming requests to VIRGO queue which is serviced by
KingCobra:

VIRGO Queue client(e.g telnet) ------> VIRGO Queue kernel service ---> Linux Workqueue handler ------> KingCobra

923. (FEATURE - DONE) EventNet kernel module service:
VIRGO eventnet client (telnet) -------> VIRGO EventNet kernel service -----> EventNet graph text files

924. (FEATURE - DONE) Related to point 22 - Reuse EventNet cloudwide logical time infinite graph in AsFer in place of
Logical clocks. At present the eventnet driver listens on port 20000 and writes the edges and vertices files in kernel
using vfs_read()/vfs_write(). These text files can then be read by the AsFer code to generate DOT files and visualize the
graph with graphviz.

925. (FEATURE - OPTIONAL) The kernel modules services listening on ports could return a JSON response when connected
instead of plaintext, conforming to REST protocol. Additional options for protocol buffers which are becoming a standard
data interchange format.

926. (FEATURE-Minimum Functionality DONE) Pointer Swizzling and Unswizzling of VIRGO addressspace pointers to/from VIRGO
Unique ID (VUID). Presently VIRGO memory system calls implement a basic minimal pointer address translation to unique kmem
location identifier.

*********************************************************************************************
                                CODE COMMIT RELATED NOTES
*********************************************************************************************

927. VIRGO code commits as on 16/05/2013
-----------------------------------
1. VIRGO cloudexec driver with a listener kernel thread service has been implemented and it listens on port 10000 on system
startup
through /etc/modules load-on-bootup facility

2. VIRGO cloudexec virgo_clone() system call has been implemented that would kernel_connect() to the VIRGO cloudexec
service listening at
port 10000

3. VIRGO cloudexec driver has been split into virgo.h (VIRGO typedefs), virgocloudexecsvc.h(VIRGO cloudexec service that is
invoked by
module_init() of VIRGO cloudexec driver) and virgo_cloudexec.c (with module ops definitions)

4. VIRGO does not implement SUN RPC interface anymore and now has its own virgo ops.

5. Lot of Kbuild related commits with commented lines for future use have been done viz., to integrate VIRGO to Kbuild,

KBUILD_EXTRA_SYMBOLS for cross-module symbol reference.

928. VIRGO code commits as on 20/05/2013
------------------------------------
1. test_virgo_clone.c testcase for sys_virgo_clone() system call works and connections are established to VIRGO cloudexec
kernel module.

2. Makefile for test_virgo_clone.c and updated buildscript.sh for headers_install for custom-built linux.

929. VIRGO code commits as on 6/6/2013
-------------------------------
1. Message header related bug fixes

930. VIRGO code commits as on 25/6/2013
---------------------------------
1.telnet to kernel service was tested and found working
2.GFP_KERNEL changed to GFP_ATOMIC in VIRGO cloudexec kernel service

931. VIRGO code commits as on 1/7/2013
---------------------------------
1. Instead of printing iovec, printing buffer correctly prints the messages
2. wake_up_process() added and function received from virgo_clone() syscall is executed with kernel_thread and results
returned to
virgo_clone() syscall client.


932. commit as on 03/07/2013
-----------------------
PRG loadbalancer preliminary code implemented. More work to be done

933. commit as on 10/07/2013
-----------------------
Tested PRG loadbalancer read config code through telnet and virgo_clone. VFS code to read from virgo_cloud.conf commented
for testing

934. commits as on 12/07/2013
------------------------
PRG loadbalancer prototype has been completed and tested with test_virgo_clone and telnet and symbol export errors and PRG
errors have been fixed

935. commits as on 16/07/2013
------------------------
read_virgo_config() and read_virgo_clone_config()(replica of read_virgo_config()) have been implemented and tested to read
the virgo_cloud.conf config parameters(at present the virgo_cloud.conf has comma separated list of ip addresses. Port is

hardcoded to 10000 for uniformity across
all nodes). Thus minimal cloud functionality with config file  support is in place. Todo things include function pointer
lookup in kernel service, more parameters to cloud config file if needed, individual configs for virgo_clone() and virgo
kernel service, kernel-to-userspace upcall and execution instead of kernel space, performance tuning etc.,

936. commits as on 17/07/2013
------------------------
moved read_virgo_config() to VIRGOcloudexec's module_init so that config is read at boot time and exported symbols are set
beforehand.
Also commented read_virgo_clone_config() as it is redundant

937. commits as on 23/07/2013
------------------------

Lack of reflection kind of facilities requires map of function_names to pointers_to_functions to be executed
on cloud has to be lookedup in the map to get pointer to function. This map is not scalable if number of functions are
in millions and size of the map increases linearly. Also having it in memory is both CPU and memory intensive.
Moreover this map has to be synchronized in all nodes for coherency and consistency which is another intensive task.
Thus name to pointer function table is at present not implemented. Suitable way to call a function by name of the function
is yet to be found out and references in this topic are scarce.

If parameterIsExecutable is set to 1 the data received from virgo_clone() is not a function but name of executable
This executable is then run on usermode using call_usermodehelper() which internally takes care of queueing the workstruct
and executes the binary as child of keventd and reaps silently. Thus workqueue component of kernel is indirectly made use
of.
This is sometimes more flexible alternative that executes a binary itself on cloud and
is preferable to clone()ing a function on cloud. Virgo_clone() syscall client or telnet needs to send the message with name
of binary.

If parameterIsExecutable is set to 0 then data received from virgo_clone() is name of a function and is executed in else
clause
using dlsym() lookup and pthread_create() in user space. This unifies both call_usermodehelper() and creating a userspace
thread
with a fixed binary which is same for any function. The dlsym lookup requires mangled function names which need to be sent
by
virgo_clone or telnet. This is far more efficient than a function pointer table.

call_usermodehelper() Kernel upcall to usermode to exec a fixed binary that would inturn execute the cloneFunction in
userspace
by spawning a pthread. cloneFunction is name of the function and not binary. This clone function will be dlsym()ed
and a pthread will be created by the fixed binary. Name of the fixed binary is hardcoded herein as
"virgo_kernelupcall_plugin". This fixed binary takes clone function as argument. For testing libvirgo.so has been created
from
virgo_cloud_test.c and separate build script to build the cloud function binaries has been added.

    - Ka.Shrinivaasan (alias) Shrinivas Kannan (alias) Srinivasan Kannan
      (https://sites.google.com/site/kuja27)


  938. commits as on 24/07/2013
  -----------------------

  test_virgo_clone unit test case updated with mangled function name to be sent to remote cloud node. Tested with
  test_virgo_clone
  end-to-end and all features are working. But sometimes kernel_connect hangs randomly (this was observed only today and
  looks similar
  to blocking vs non-blocking problem. Origin unknown).

  - Ka.Shrinivaasan (alias) Shrinivas Kannan (alias) Srinivasan Kannan
    (https://sites.google.com/site/kuja27)


  939. commits as on 29/07/2013
  -----------------------

  Added kernel mode execution in the clone_func and created a sample kernel_thread for a cloud function. Some File IO logging
  added to upcall
  binaries and parameterIsExecutable has been moved to virgo.h


  940. commits as on 30/07/2013
  -----------------------
  New usecase virgo_cloud_test_kernelspace.ko kernel module has been added. This exports a function
  virgo_cloud_test_kernelspace() and is
  accessed by virgo_cloudexec kernel service to spawn a kernel thread that is executed in kernel addresspace. This Kernel
  mode execution
  on cloud adds a unique ability to VIRGO cloud platform to seamlessly integrate hardware devices on to cloud and
  transparently send commands
  to them from a remote cloud node through virgo_clone().

  Thus above feature adds power to VIRGO cloud to make it act as a single "logical device driver" though devices are in
  geographically in a remote server.


  941. commits as on 01/08/2013 and 02/08/2013
  ------------------------------------------
  Added Bash shell commandline with -c option for call_usermodehelper upcall clauses to pass in remote virgo_clone command
  message as
  arguments to it. Also tried output redirection but it works some times that too with a fatal kernel panic.

  Ideal solutions are :
  1. either to do a copy_from_user() for message buffer from user address space (or)

2. somehow rebuild the kernel with fd_install() pointing stdout to a VFS file* struct. In older kernels like 2.6.x, there is an fd_install code
with in kmod.c (___call_usermodehelper()) which has been redesigned in kernel 3.x versions and fd_install has been removed in kmod.c .
3. Create a Netlink socket listener in userspace and send message up from kernel Netlink socket.

All the above are quite intensive and time consuming to implement.Moreover doing FileIO in usermode helper is strongly discouraged in kernel docs

Since Objective of VIRGO is to virtualize the cloud as single execution "machine", doing an upcall (which would run with root abilities) is
redundant often and kernel mode execution is sufficient. Kernel mode execution with intermodule function invocation can literally take over
the entire board in remote machine (since it can access PCI bus, RAM and all other device cards)

As a longterm design goal, VIRGO can be implemented as a separate protocol itself and sk_buff packet payload from remote machine
can be parsed by kernel service and kernel_thread can be created for the message.

942. commits as on 05/08/2013:
---------------------------
Major commits done for kernel upcall usermode output logging with fd_install redirection to a VFS file. With this it has become easy for user space to communicate runtime data to Kernel space. Also a new strip_control_M() function has been added to strip \r\n or " ".

943. 11 August 2013:
---------------
Open Source Design and Academic Research Notes uploaded to
http://sourceforge.net/projects/acadpdrafts/files/MiscellaneousOpenSourceDesignAndAcademicResearchNotes_2013-08-11.pdf/download


944. commits as on 23 August 2013
-----------------------------
New Multithreading Feature added for VIRGO Kernel Service - action item 5 in ToDo list above (virgo_cloudexec driver module). All dependent headers changed for kernel threadlocalizing global data.

945. commits as on 1 September 2013
-------------------------------
GNU Copyright license and Product Owner Profile (for identity of license issuer) have been committed. Also Virgo Memory Pooling - virgo_malloc() related initial design notes (handwritten scanned) have been committed(http://sourceforge.net/p/virgo-linux/code-0/HEAD/tree/trunk/virgo-docs/VIRGO_Memory_Pooling_virgomalloc_initial_design_notes.pdf)

946. commits as on 14 September 2013
--------------------------------
Updated virgo malloc design handwritten nodes on kmalloc() and malloc() usage in kernelspace and userspace execution mode
of virgo_cloudexec service (http://sourceforge.net/p/virgo-linux/code-0/HEAD/tree/trunk/virgo-
docs/VIRGO_Memory_Pooling_virgomalloc_design_notes_2_14September2013.pdf). As described in handwritten notes,
virgo_malloc() and related system calls might be needed when a large scale allocation of kernel memory is needed when in
kernel space execution mode and large scale userspace memory when in user modes (function and executable modes). Thus a
cloud memory pool both in user and kernel space is possible.


------------------------------------------
947. VIRGO virtual addressing
------------------------------------------
VIRGO virtual address is defined with the following datatype:

struct virgo_address
{
        int node_id;
        void* addr;
};

VIRGO address translation table is defined with following datatype:

struct virgo_addr_transtable
{
        int node_id;
        void* addr;
};


----------------------------------------------------
948. VIRGO memory pooling prototypical implementation
----------------------------------------------------
VIRGO memory pooling implementation as per the design notes committed as above is to be implemented as a prototype under
separate directory
under drivers/virgo/memorypooling and $LINUX_SRC_ROOT/virgo_malloc. But the underlying code is more or less similar to
drivers/virgo/cpupooling and $LINUX_SRC_ROOT/virgo_clone.

virgo_malloc() and related syscalls and virgo mempool driver connect to and listen on port different from cpupooling
driver. Though all these code can be within cpupooling itself, mempooling is implemented as separate driver and co-exists
with cpupooling on bootup (/etc/modules). This enables clear demarcation of functionalities for CPU and Memory
virtualization.

949. Commits as on 17 September 2013
--------------------------------
Initial untested prototype code - virgo_malloc and virgo mempool driver - for VIRGO Memory Pooling has been committed -

copied and modified from virgo_clone client and kernel driver service.

950. Commits as on 19 September 2013
--------------------------------
3.7.8 Kernel full build done and compilation errors in VIRGO malloc and mempool driver code and more functions code added

951. Commits as on 23 September 2013
--------------------------------
Updated virgo_malloc.c with two functions, int_to_str() and addr_to_str(), using kmalloc() with full kernel re-build.
(Rather a re-re-build because some source file updates in previous build got deleted somehow mysteriously. This could be
related to Cybercrime issues mentioned in https://sourceforge.net/p/usb-md/code-0/HEAD/tree/USBmd_notes.txt )

952. Commits as on 24 September 2013
--------------------------------
Updated syscall*.tbl files, staging.sh, Makefiles for virgo_malloc(),virgo_set(),virgo_get() and virgo_free() memory
pooling syscalls. New testcase test_virgo_malloc for virgo_malloc(), virgo_set(), virgo_get(), virgo_free() has been added
to repository. This testcase might have to be updated if return type and args to virgo_malloc+ syscalls are to be changed.

953. Commits as on 25 September 2013
--------------------------------
All build related errors fixed after kernel rebuild some changes made to function names to reflect their
names specific to memory pooling. Updated /etc/modules also has been committed to repository.

954. Commits as on 26 September 2013
--------------------------------
Circular dependency error in standalone build of cpu pooling and memory pooling drivers fixed and
datatypes and declarations for CPU pooling and Memory Pooling drivers have been segregated into respective header files
(virgo.h and
virgo_mempool.h with corresponding service header files) to avoid any dependency error.

955. Commits as on 27 September 2013
--------------------------------
Major commits for Memory Pooling Driver listen port change and parsing VIRGO memory pooling commands have been done.

956. Commits as on 30 September 2013
--------------------------------
New parser functions added for parameter parsing and initial testing on virgo_malloc() works with telnet client with logs
in test_logs/

957. Commits as on 1 October 2013
--------------------------------
Removed strcpy in virgo_malloc as ongoing bugfix for buffer truncation in syscall path.

958. Commits as on 7 October 2013

------------------------------
Fixed the buffer truncation error from virgo_malloc syscall to mempool driver service which was caused by
sizeof() for a char*. BUF_SIZE is now used for size in both syscall client and mempool kernel service.

959. Commits as on 9 October 2013 and 10 October 2013
----------------------------------------------------
Mempool driver kernelspace virgo mempool ops have been rewritten due to lack of facilities to return a
value from kernel thread function. Since mempool service already spawns a kthread, this seems to be sufficient. Also the
iov.iov_len in virgo_malloc has been changed from BUF_SIZE to strlen(buf) since BUF_SIZE
causes the kernel socket to block as it waits for more data to be sent.

960. Commits as on 11 October 2013
------------------------------
sscanf format error for virgo_cloud_malloc() return pointer address and sock_release() null pointer exception has been
rectified.
Added str_to_addr() utility function.

961. Commits as on 14 October 2013 and 15 October 2013
-----------------------------------------------------
Updated todo list.

Rewritten virgo_cloud_malloc() syscall with:
- mutexed virgo_cloud_malloc() loop
- redefined virgo address translation table in virgo_mempool.h
- str_to_addr(): removed (void**) cast due to null sscanf though it should have worked

962. Commits as on 18 October 2013
------------------------------
Continued debugging of null sscanf - added str_to_addr2() which uses simple_strtoll() kernel function
for scanning pointer as long long from string and casting it to void*. Also more %p qualifiers where
added in str_to_addr() for debugging.

Based on latest test_virgo_malloc run, simple_strtoll() correctly parses the address string into a long long base 16 and
then is reinterpret_cast to void*. Logs in test/

963. Commits as on 21 October 2013
------------------------------
Kern.log for testing after vtranstable addr fix with simple_strtoll() added to repository and still the other %p qualifiers
do not work and only simple_strtoll() parses the address correctly.

964. Commits as on 24 October 2013
------------------------------
Lot of bugfixes made to virgo_malloc.c for scanning address into VIRGO transtable and size computation. Testcase
test_virgo_malloc.c has also been modified to do reinterpret cast of long long into (struct virgo_address*) and

corresponding test logs have been added to repository under virgo_malloc/test.

Though the above sys_virgo_malloc() works, the return value is a kernel pointer if the virgo_malloc executes in the Kernel
mode which is more likely than User mode (call_usermodehelper which is circuitous). Moreover copy_from_user() or
copy_to_user() may not be directly useful here as this is an address allocation routine. The long long reinterpret cast
obfuscates the virgo_address(User or Kernel) as a large integer which is a unique id for the allocated memory on cloud.
Initial testing of sys_virgo_set() causes a Kernel Panic as usual probably due to direct access of struct virgo_address*.
Alternatives are to use only long long for allocation unique-id everywhere or do copy_to_user() or copy_from_user() of the
address on a user supplied buffer. Also vtranstable can be made into a bucketed hash table that maps each alloc_id to a
chained virgo malloc chunks than the present sequential addressing which is more similar to open addressing.

965. Commits as on 25 October 2013
------------------------------
virgo_malloc.c has been rewritten by adding a userspace __user pointer to virgo_get() and virgo_set() syscalls which are
internally copied with copy_from_user() and copy_to_user() kernel function to get and set userspace from kernelspace.Header
file syscalls.h has been updated with changed syscalls prototypes.Two functions have been added to map a VIRGO address to a
unique virgo identifier and viceversa for abstracting hardware addresses from userspace as mentioned in previous commit
notes. VIRGO cloud mempool kernelspace driver has been updated to use virgo_mempool_args* instead of void* and VIRGO
cloudexec mempool driverhas been updated accordingly during intermodule invocation.The virgo_malloc syscall client has been
updated to modified signatures and return types for all mempool alloc,get,set,free syscalls.

966. Commits as on 29 October 2013
------------------------------
Miscellaneous ongoing bugfixes for virgo_set() syscall error in copy_from_user().

967. Commits as on 2 November 2013
------------------------------
Due to an issue which corrupts the kernel memory, presently telnet path to VIRGO mempool driver has been
tested after commits on 31 October 2013 and 1 November 2013 and is working but again there is an issue in kstrtoul() that
returns the wrong address in virgo_cloud_mempool_kernelspace.ko that gives the address for
data to set.

968. Commits as on 6 November 2013
------------------------------
New parser function virgo_parse_integer() has been added to virgo_cloud_mempool_kernelspace driver module which is carried
over from
lib/kstrtox.c and modified locally to add an if clause to discard quotes and unquotes. With this the telnet path commands
for virgo_malloc()
and virgo_set() are working. Today's kern.log has been added to repository in test_logs/.

969. Commits as on 7 November 2013
------------------------------
In addition to virgo_malloc and virgo_set, virgo_get is also working through telnet path after today's commit for
"virgodata:" prefix in virgo_cloud_mempool_kernelspace.ko. This prefix is needed to differentiate data and address so that

toAddressString() can be invoked to sprintf() the address in virgo_cloudexec_mempool.ko. Also mempool command parser has been updated to strcmp() virgo_cloud_get command also.

970. Commits as on 11 November 2013
-------------------------------
More testing done on telnet path for virgo_malloc, virgo_set and virgo_get commands which work correctly. But there seem to be unrelated
kmem_cache_trace_alloc panics that follow each successful virgo command execution. kern.log for this has been added to repository.

971. Commits as on 22 November 2013
-------------------------------
More testing done on telnet path for virgo_malloc,virgo_set and virgo_set after commenting kernel socket shutdown code in the VIRGO cloudexec
mempool sendto code. Kernel panics do not occur after commenting kernel socket shutdown.

972. Commits as on 2 December 2013
-------------------------------
Lots of testing were done on telnet path and syscall path connection to VIRGO mempool driver and screenshots for working telnet path (virgo_malloc, virgo_set and virgo_get) have been committed to repository. Intriguingly, the syscall path is suddenly witnessing series of broken pipe erros, blocking errors etc., which are mostly Heisenbugs.

973. Commits as on 5 December 2013
-------------------------------
More testing on system call path done for virgo_malloc(), virgo_set() and virgo_get() system calls with test_virgo_malloc.c. All three syscalls work in syscall path after lot of bugfixes. Kern.log that has logs for allocating memory in remote cloud node with virgo_malloc, sets data "test_virgo_malloc_data" with virgo_set and retrieves data with virgo_get.


VIRGO version 12.0 tagged.

974. Commits as on 12 March 2014
-----------------------------
Initial VIRGO queueing driver implemented that flips between two internal queues: 1) a native queue implemented locally and 2) wrapper around linux kernel's workqueue facility 3) push_request() modified to pass on the request data to the workqueue handler using container_of on a wrapper
structure virgo_workqueue_request.

975. Commits as on 20 March 2014
-----------------------------
- VIRGO queue with additional boolean flags for its use as KingCobra queue
- KingCobra kernel space driver that is invoked by the VIRGO workqueue handler

976. Commits as on 30 March 2014
-----------------------------
- VIRGO mempool driver has been augmented with use_as_kingcobra_service flags in CPU pooling and Memory pooling drivers

977. Commits as on 6 April 2014
---------------------------
- VIRGO mempool driver recvfrom() function's if clause for KingCobra has been updated for REQUEST header formatting mentioned in KingCobra design notes

978. Commits as on 7 April 2014
---------------------------
- generate_logical_timestamp() function has been implemented in VIRGO mempool driver that generates timestamps based on 3 boolean flags. At present machine_timestamp is generated and prepended to the request to be pushed to VIRGO queue driver and then serviced by KingCobra.

979. Commits as on 25 April 2014
----------------------------
- client ip address in VIRGO mempool recvfrom KingCobra if clause is converted to host byte order from network byte order with ntohl()

980. Commits as on 5 May 2014
-------------------------
- Telnet path commands for VIRGO cloud file system - virgo_cloud_open(), virgo_cloud_read(), virgo_cloud_write(), virgo_cloud_close() has been implemented and test logs have been added to repository (drivers/virgo/cloudfs/ and cloudfs/testlogs) and kernel upcall path for paramIsExecutable=0

981. Commits as on 7 May 2014
-------------------------
- Bugfixes to tokenization in kernel upcall plugin with strsep() for args passed on to the userspace

982. Commits as on 8 May 2014
-------------------------
- Bugfixes to virgo_cloud_fs.c for kernel upcall (parameterIsExecutable=0) and with these the kernel to userspace upcall and writing to a file in userspace (virgofstest.txt) works. Logs and screenshots for this are added to repository in test_logs/

983. Commits as on 6 June 2014
--------------------------
- VIRGO File System Calls Path implementation has been committed. Lots of Linux Full Build compilation errors fixed and new integer parsing functionality added (similar to driver modules).  For the timebeing all syscalls invoke loadbalancer. This may be further optimized with a sticky flag to remember the first invocation which might be usually virgo_open syscall to get the VFS descriptor that is used in subsequent syscalls.

984. Commits as on 3 July 2014

```
---------------------------
- More testing and bugfixes for VIRGO File System syscalls have been done. virgo_write() causes kernel panic.
```

985. 7 July 2014 - virgo_write() kernel panic notes:
-----------------------------------------------
warning within http://lxr.free-electrons.com/source/arch/x86/kernel/smp.c#L121:

```
static void native_smp_send_reschedule(int cpu)
{
        if (unlikely(cpu_is_offline(cpu))) {
                WARN_ON(1);
                return;
        }
        apic->send_IPI_mask(cpumask_of(cpu), RESCHEDULE_VECTOR);
}
```

This is probably a fixed kernel bug in <3.7.8 but recurring in 3.7.8:
- http://lkml.iu.edu/hypermail/linux/kernel/1205.3/00653.html
- http://www.kernelhub.org/?p=3&msg=74473&body_id=72338
- http://lists.openwall.net/linux-kernel/2012/09/07/22
- https://bugzilla.kernel.org/show_bug.cgi?id=54331
- https://bbs.archlinux.org/viewtopic.php?id=156276


986. Commits as on 29 July 2014
---------------------------
All VIRGO drivers(cloudfs, queuing, cpupooling and memorypooling) have been built on 3.15.5 kernel with some Makefile
changes for ccflags and paths


--------------------------------------------------------------------------------
Commits as on 17 August 2014
--------------------------------------------------------------------------------
987. (FEATURE - DONE) VIRGO Kernel Modules and System Calls major rewrite for 3.15.5 kernel - 17 August 2014
--------------------------------------------------------------------------------
1. VIRGO config files have been split into /etc/virgo_client.conf and /etc/virgo_cloud.conf to delink the cloud client and
kernel service
config parameters reading and to do away with oft occurring symbol lookup errors and multiple definition errors for
num_cloud_nodes and
node_ip_addrs_in_cloud - these errors are frequent in 3.15.5 kernel than 3.7.8 kernel.

2. Each VIRGO module and system call now reads the config file independent of others - there is a
read_virgo_config_<module>_<client_or_service>() function variant for each driver and system call. Though at present smacks
of a replicated code, in future the config reads for each component (system call or module) might vary significantly
depending on necessities.

3. New kernel module config has been added in drivers/virgo. This is for future prospective use as a config export driver that can
be looked up by any other VIRGO module for config parameters.

4. include/linux/virgo_config.h has the declarations for all the config variables declared within each of the VIRGO kernel modules.

5. Config variables in each driver and system call have been named with prefix and suffix to differentiate the module and/or system call it serves.

6. In geographically distributed cloud virgo_client.conf has to be in client nodes and virgo_cloud.conf has to be in cloud nodes. For VIRGO Queue - KingCobra REQUEST-REPLY peer-to-peer messaging system same node can have virgo_client.conf and virgo_cloud.conf.

7. Above segregation largely simplifies the build process as each module and system call is independently built without need for a symbol to be exported from other module by pre-loading it.

8. VIRGO File system driver and system calls have been tested with above changes and the virgo_open(),virgo_read() and virgo_write() calls work with much less crashes and freeze problems compared to 3.7.8 (some crashes in VIRGO FS syscalls in 3.7.8 where already reported kernel bugs which seem to have been fixed in 3.15.5). Today's kern.log test logs have been committed to repository.

------------------------------------------------
988. Committed as on 23 August 2014
------------------------------------------------
Commenting use_as_kingcobra_service if clauses temporarily as disabling also doesnot work and only commenting the block works for VIRGO syscall path. Quite weird as to how this relates to the problem. As this is a heisenbug further testing is difficult and sufficient testing has been done with logs committed to repository. Probably a runtime symbol lookup for kingcobra
causes the freeze.
For forwarding messages to KingCobra and VIRGO queues, cpupooling driver is sufficient which also has the use_as_kingcobra_service clause.

------------------------------------------------
989. Committed as on 23 August 2014 and 24 August 2014
------------------------------------------------
As cpupooling driver has the same crash problem with kernel_accept() when KingCobra has benn enabled, KingCobra clauses have been commented in both cpupooling and memorypooling drivers. Instead queueing driver has been updated with a kernel service infrastructure to accept connections at port 60000. With this following paths are available for KingCobra requests:

        VIRGO cpupooling or memorypooling ====> VIRGO Queue =====> KingCobra

                                (or)

         VIRGO Queue kernel service ============================> KingCobra


 --------------------------------------------------
 990. Committed as on 26 August 2014
 --------------------------------------------------
 - all kmallocs have been made into GFP_ATOMIC instead of GFP_KERNEL
 - moved some kingcobra related header code before kernel_recvmsg()
 - some header file changes for set_fs()

 This code has been tested with modified code for KingCobra and the standalone
 kernel service that accepts requests from telnet directly at port 60000, pushes to virgo_queue
 and is handled to invoke KingCobra servicerequest kernelspace function, works
 (the kernel_recvmsg() crash was most probably due to Read-Only filesystem -errno printed is -30)


 -----------------------------------------------------------------
 VIRGO version 14.9.9 has been release tagged on 9 September 2014
 -----------------------------------------------------------------


 ---------------------------------------------------------
 991. Committed as on 26 November 2014
 ---------------------------------------------------------
 New kernel module cloudsync has been added to repository under drivers/virgo that can be used for synchronization(lock()
 and unlock()) necessities in VIRGO cloud. Presently Bakery Algorithm has been implemented.


 ---------------------------------------------------------
 992. Committed as on 27 November 2014
 ---------------------------------------------------------
 virgo_bakery.h bakery_lock() has been modified to take 2 parameters - thread_id and number of for loops (1 or 2)


 ---------------------------------------------------------
 993. Committed as on 2 December 2014
 ---------------------------------------------------------
 VIRGO bakery algorithm implementation has been rewritten with some bugfixes. Sometimes there are soft lockup errors due to
 looping in kernel time durations for which are kernel build configurable.


 -------------------------------------------------------------------
 994. Committed as on 17 December 2014
 -------------------------------------------------------------------
 Initial code commits for VIRGO EventNet kernel module service:
 -------------------------------------------------------------------
 1.EventNet Kernel Service listens on port 20000

 2.It receives eventnet log messages from VIRGO cloud nodes and writes the log messages
 after parsing into two text files /var/log/eventnet/EventNetEdges.txt and

/var/log/eventnet/EventNetVertices.txt by VFS calls

3.These text files can then be processed by the EventNet implementations in AsFer (python pygraph and
C++ boost::graph based)

4.Two new directories virgo/utils and virgo/eventnet have been added.

5.virgo/eventnet has the new VIRGO EventNet kernel module service implementation that listens on
port 20000.

6.virgo/utils is the new generic utilities driver that has a virgo_eventnet_log()
exported function which connects to EventNet kernel service and sends the vertex and edge eventnet
log messages which are parsed by kernel service and written to the two text files above.

7.EventNet log messages have two formats:
    - Edge message - "eventnet_edgemsg#<id>#<from_event>#<to_event>"
    - Vertex message - "eventnet_vertextmsg#<id>-<partakers csv>-<partaker conversations csv>"

8.The utilities driver Module.symvers have to be copied to any driver which are
then merged with the symbol files of the corresponding driver. Target clean has to be commented while
building the unified Module.symvers because it erases symvers carried over earlier.

9.virgo/utils driver can be populated with all necessary utility exported functions that might be needed
in other VIRGO drivers.

10.Calls to virgo_eventnet_log() have to be #ifdef guarded as this is quite network intensive.

--------------------------------------------------------------------
995. Commits as on 18 December 2014
--------------------------------------------------------------------
Miscellaneous bugfixes,logs and screenshot

- virgo_cloudexec_eventnet.c - eventnet messages parser errors and eventnet_func bugs fixed
- virgo_cloud_eventnet_kernelspace.c - filp_open() args updated due to vfs_write() kernel panics. The vertexmessage
vfs_write is done after looping through the vertice textfile and appending the conversation to the existing vertex.Some
more code has to be added.
- VIRGO EventNet build script updated for copying Module.symvers from utils driver for merging with eventnet Module.symvers
during Kbuild
- Other build generated sources and kernel objects
- new testlogs directory with screenshot for edgemsg sent to EventNet kernel service and kern.log with previous history for
vfs_write() panics due to permissions and the logs for working filp_open() fixed version
- vertex message update


--------------------------------------------------------------------

996. Commits as on 2,3,4 January 2015
---------------------------------------------------------------------
- fixes for virgo eventnet vertex and edge message text file vfs_write() errors
- kern.logs and screenshots


---------------------------------------------------------------------
VIRGO version 15.1.8 release tagged on 8 January 2015
---------------------------------------------------------------------


----------------------------------------------------------------------------------------
772. (FEATURE) Commits as on 3 March 2015 - Initial commits for Kernel Analytics Module which reads the
/etc/virgo_kernel_analytics.conf config (and) VIRGO memorypooling Key-Value Store Architecture Diagram - - this section is
an extended draft on respective topics in NeuronRain AstroInfer design -  https://github.com/shrinivaasanka/asfer-github-
code/blob/master/asfer-docs/AstroInferDesign.txt
----------------------------------------------------------------------------------------
- Architecture of Key-Value Store in memorypooling (virgo_malloc,virgo_get,virgo_set,virgo_free) has been
uploaded as a diagram at http://sourceforge.net/p/virgo-linux/code-0/HEAD/tree/trunk/virgo-
docs/VIRGOLinuxKernel_KeyValueStore_and_Modules_Interaction.jpg

- new kernel_analytics driver for AsFer <=> VIRGO+USBmd+KingCobra interface has been added.
- virgo_kernel_analytics.conf having csv(s) of key-value pairs of analytics variables is set by AsFer or any other Machine
Learning code.  With this VIRGO Linux Kernel is endowed with abilities to dynamically evolve than being just a platform for
user code. Implications are huge - for example, a config variable "MaxNetworkBandwidth=255" set by the ML miner in
userspace based on a Perceptron or Logistic Regression executed on network logs can be read by a kernel module that limits
the network traffic to 255Mbps. Thus kernel dynamically changes behaviour.
- kernel_analytics Driver build script has been added


-----------------------------------------------------------------
997. Commits as on 6 March 2015
-----------------------------------------------------------------
- code has been added in VIRGO config module to import EXPORTed kernel_analytics config key-pair array
set by Apache Spark (mined from Uncomplicated Fire Wall logs) and manually and write to kern.log.


-----------------------------------------------------------------
NeuronRain version 15.6.15 release tagged
-----------------------------------------------------------------


-----------------------------------------------------------------
998. Portability to linux kernel 4.0.5
-----------------------------------------------------------------
The VIRGO kernel module drivers are based on kernel 3.15.5. With kernel 4.0.5 kernel which is the latest following
compilation and LD errors occur - this is on cloudfs VIRGO File System driver :
- msghdr has to be user_msghdr for iov and iov_len as there is a segregation of msghdr
- modules_install throws an error in scripts/Makefile.modinst while overwriting already installed module

```
--------------------------------------------------------------------
999. Commits as on 9 July 2015
--------------------------------------------------------------------
VIRGO cpupooling driver has been ported to linux kernel 4.0.5 with msghdr changes as mentioned previously
with kern.log for VIRGO cpupooling driver invoked in parameterIsExecutable=2 (kernel module invocation)
added in testlogs


--------------------------------------------------------------------
1000. Commits as on 10,11 July 2015
--------------------------------------------------------------------
VIRGO Kernel Modules:
- memorypooling
- cloudfs
- utils
- config
- kernel_analytics
- cloudsync
- eventnet
- queuing
along with cpupooling have been ported to Linux Kernel 4.0.5 - Makefile and header files have been
updated wherever required.


--------------------------------------------------------------------
1001. Commits as on 20,21,22 July 2015
--------------------------------------------------------------------
Due to SourceForge Storage Disaster(http://sourceforge.net/blog/sourceforge-infrastructure-and-service-restoration/),
the github replica of VIRGO is urgently updated with some important changes for msg_iter,iovec
etc., in 4.0.5 kernel port specifically for KingCobra and VIRGO Queueing. These have to be committed to SourceForge
Krishna_iResearch
repository at http://sourceforge.net/users/ka_shrinivaasan once SourceForge repos are restored.
Time to move on to the manufacturing hub? GitHub ;-)
---------------------------------
1002. VIRGO Queueing Kernel Module Linux Kernel 4.0.5 port:
----------------------------------------------------
- msg_iter is used instead of user_msghdr
- kvec changed to iovec
- Miscellaneous BUF_SIZE related changes
- kern.logs for these have been added to testlogs
- Module.symvers has been recreated with KingCobra Module.symvers from 4.0.5 KingCobra build
- clean target commented in build script as it wipes out Module.symvers
- updated .ko and .mod.c
---------------------------------
1003. KingCobra Module Linux Kernel 4.0.5 port
```

```
--------------------------------------------------------
- vfs_write() has a problem in 4.0.5
- the filp_open() args and flags which were working in 3.15.5 cause a
kernel panic implicitly and nothing was written to logs
- It took a very long time to figure out the reason to be vfs_write and filp_open
- O_CREAT, O_RDWR and O_LARGEFILE cause the panic and only O_APPEND is working, but
does not do vfs_write(). All other VIRGO Queue + KingCobra functionalities work viz.,
enqueueing, workqueue handler invocation, dequeueing, invoking kingcobra kernelspace service
request function from VIRGO queue handler, timestamp, timestamp and IP parser, reply_to_publisher etc.,
- As mentioned in Greg Kroah Hartman's "Driving me nuts", persistence in Kernel space is
a bad idea but still seems to be a necessary stuff - yet only vfs calls are used which have to be safe
- Thus KingCobra has to be in-memory only in 4.0.5 if vfs_write() doesn't work
- Intriguingly cloudfs filesystems primitives - virgo_cloud_open, virgo_cloud_read, virgo_cloud_write etc.,
work perfectly and append to a file.
- kern.logs for these have been added to testlogs
- Module.symvers has been recreated for 4.0.5
- updated .ko and .mod.c


--------------------------------------------------------------------
Due to SourceForge outage and for a future code diversification
NeuronRain codebases (AsFer, USBmd, VIRGO, KingCobra)
in http://sourceforge.net/u/userid-769929/profile/ have been
replicated in GitHub also - https://github.com/shrinivaasanka
excluding some huge logs due to Large File Errors in GitHub.
--------------------------------------------------------------------


--------------------------------------------------------------------
1004. Commits as on 30 July 2015
--------------------------------------------------------------------
VIRGO system calls have been ported to Linux Kernel 4.0.5 with commented gcc option -Wimplicit-function-declaration,
msghdr and iovec changes similar to drivers mentioned in previous commit notes above. But Kernel 4.1.3 has some Makefile
and build issues.
The NeuronRain codebases in SourceForge and GitHub would henceforth be mostly and always out-of-sync and not guaranteed to
be replicas - might get diversified into different research and development directions (e.g one codebase might be more
focussed on IoT while the other towards enterprise bigdata analytics integration with kernel and training which is yet to
be designed- depend on lot of constraints)


--------------------------------------------------------------------
1005. Commits as on 2,3 August 2015
--------------------------------------------------------------------
- new .config file added which is created from menuconfig
- drivers/Kconfig has been updated with 4.0.5 drivers/Kconfig for trace event linker errors
Linux Kernel 4.0.5 - KConfig is drivers/ has been updated to resolve RAS driver trace event linker error. RAS was not
included in KConfig earlier.
```

    - link-vmlinux.sh has been replaced with 4.0.5 kernel version


    -------------------------------------------------------------------------
    1006. Commits as on 12 August 2015
    -------------------------------------------------------------------------
    VIRGO Linux Kernel 4.1.5 port - related code changes - some important notes:
    -------------------------------------------------------------------------
    - Linux Kernel 4.0.5 build suddenly had a serious root shell drop error in initramfs which was not resolved by:
            - adding rootdelay in grub
            - disabling uuid for block devices in grub config
            - mounting in read/write mode in recovery mode
            - no /dev/mapper related errors
            - repeated exits in root shell
            - delay before mount of root device in initrd scripts
    - mysteriously there were some firmware microcode bundle executions in ieucodetool
    - Above showed a serious grub corruption or /boot MBR bug or 4.0.5 VIRGO kernel build problem
    - Linux 4.0.x kernels are EOL-ed
    - Hence VIRGO is ported to 4.1.5 kernel released few days ago
    - Only minimum files have been changed as in commit log for Makefiles and syscall table and headers and a build script has
    been added
    for 4.1.5:
        Changed paths:
        A buildscript_4.1.5.sh
        M linux-kernel-extensions/Makefile
        M linux-kernel-extensions/arch/x86/syscalls/Makefile
        M linux-kernel-extensions/arch/x86/syscalls/syscall_32.tbl
        M linux-kernel-extensions/drivers/Makefile
        M linux-kernel-extensions/include/linux/syscalls.h


    - Above minimum changes were enough to build an overlay-ed Linux Kernel with VIRGO codebase


    -------------------------------------------------------------------------
    1007. Commits as on 14,15,16 August 2015
    -------------------------------------------------------------------------
    Executed the minimum end-end telnet path primitives in Linux kernel 4.1.5 VIRGO code:
    - cpu virtualization
    - memory virtualization
    - filesystem virtualization (updated filp_open flags)
    and committed logs and screenshots for the above.


    -------------------------------------------------------------------------
    1008. Commits as on 17 August 2015
    -------------------------------------------------------------------------
    VIRGO queue driver:

- Rebuilt Module.symvers
- kern.log for telnet request to VIRGO Queue + KingCobra queueing system in kernelspace


-----------------------------------------------------------------
1009. Commits as on 25,26 September 2015
-----------------------------------------------------------------
VIRGO Linux Kernel 4.1.5 - memory system calls:
----------------------------------------------
- updated testcases and added logs for syscalls invoked separately(malloc,set,get,free)
- The often observed unpredictable heisen kernel panics occur with 4.1.5 kernel too. The logs are 2.3G and
only grepped output is committed to repository.
- virgo_malloc.c has been updated with kstrdup() to copy the buf to iov.iov_base which was earlier
crashing in copy_from_iter() within tcp code. This problem did not happen in 3.15.5 kernel.
- But virgo_clone syscall code works without any changes to iov_base as above which does a strcpy()
 which is an internal memcpy() though. So what causes this crash in memory system calls alone
is a mystery.
- new insmod script has been added to load the VIRGO memory modules as necessary instead of at boot time.
- test_virgo_malloc.c and its Makefile has been updated.

VIRGO Linux Kernel 4.1.5 - filesystem calls- testcases and logs:
---------------------------------------------------------------
  - added insmod script for VIRGO filesystem drivers
  - test_virgo_filesystem.c has been updated for syscall numbers in 4.1.5 VIRGO kernel
  - virgo_fs.c syscalls code has been updated for iov.iov_base kstrdup() - without this there are kernel panics in
copy_from_iter(). kern.log
testlogs have been added, but there are heisen kernel panics. The virgo syscalls are executed but not written to kern.log
due to these crashes.
Thus execution logs are missing for VIRGO filesystem syscalls.


-----------------------------------------------------------------
1010. Commits as on 28,29 September 2015
-----------------------------------------------------------------


VIRGO Linux Kernel 4.1.5 filesystem syscalls:
---------------------------------------------
- Rewrote iov_base code with a separate iovbuf set to iov_base and strcpy()-ing the syscall command to iov_base similar to
VIRGO
memory syscalls
- Pleasantly the same iovbuf code that crashes in memory syscalls works for VIRGO FS without crash.Thus both virgo_clone
and virgo_filesystem
syscalls work without issues in 4.1.5 and virgo_malloc() works erratically in 4.1.5 which remains as issue.
- kern.log for VIRGO FS syscalls and virgofstest text file written by virgo_write() have been added to repository

VIRGO Linux 4.1.5 kernel memory syscalls:
------------------------------------
- rewrote the iov_base buffer code for all VIRGO memory syscalls by allocating separate iovbuf and copying the message to
it - this just replicates the virgo_clone syscall behaviour which works without any crashes mysteriously.
- did extensive repetitive tests that were frequented by numerous kernel panics and crashes
- The stability of syscalls code with 3.15.5 kernel appears to be completely absent in 4.1.5
- The telnet path works relatively better though
- Difference between virgo_clone and virgo_malloc syscalls despite having same kernel sockets code looks like a non-trivial
bug and a kernel stability issue.
- kernel OOPS traces are quite erratic.
- Makefile path in testcase has been updated


-----------------------------------------------------------------------
1011. Commits as on 4 October 2015
-----------------------------------------------------------------------
VIRGO Linux Kernel 4.1.5 - Memory System Calls:
-----------------------------------------------
- replaced copy_to_user() with a memcpy()
- updated the testcase with an example VUID hardcoded.
- str_to_addr2() is done on iov_base instead of buf which was causing NULL parsing
- kern.log with above resolutions and multiple VIRGO memory syscalls tests - malloc,get,set
- With above VIRGO malloc and set syscalls work relatively causing less number of random kernel panics
- return values of memory calls set to 0
- in virgo_get() syscall, memcpy() of iov_base is done to data_out userspace pointer
- kern.log with working logs for syscalls - virgo_malloc(), virgo_set(), virgo_get() but still there are random kernel
panics
- Abridged kern.log for VIRGO Memory System Calls with 4.1.5 Kernel - shows example logs for virgo_malloc(), virgo_set()
and virgo_get()


------------------------------------------------------------------
1012. Commits as on 14 October 2015
------------------------------------------------------------------
VIRGO Queue Workqueue handler usermode clause has been updated with 4.1.5 kernel paths and kingcobra in user mode is
enabled for invoking KingCobra Cloud Perfect Forwarding.


------------------------------------------------------------------
1013. Commits as on 15 October 2015
------------------------------------------------------------------
- Updated VIRGO Queue kernel binaries and build generated sources
- virgo_queue.h has been modified for call_usermodehelper() - set_ds() and fd_install() have been uncommented for output
redirection. Output redirection works but there are "alloc_fd: slot 1 not NULL!" errors at random (kern.log in kingcobra
testlogs) which seems to be a new feature in 4.1.5 kernel. This did not happen in 3.7.8-3.15.5 kernels


------------------------------------------------------------------

1014. Commits as on 3 November 2015
------------------------------------------------------------------
- kern.log for VIRGO kernel_analytics+config drivers which export the analytics variables from
/etc/virgo_kernel_analytics.conf kernel-wide and print them in config driver has been added to config/testlogs


------------------------------------------------------------------
1015. Commits as on 10 January 2016
------------------------------------------------------------------
NeuronRain VIRGO enterprise version 2016.1.10 released.


----------------------------------------------------------------------
NeuronRain - AsFer commits for VIRGO - C++ and C Python extensions
- Commits as on 29 January 2016
----------------------------------------------------------------------
--------------------------------------------------------------------------
1016. (FEATURE - DONE)  Python-C++-VIRGOKernel and Python-C-VIRGOKernel boost::python and cpython implementations:
--------------------------------------------------------------------------
- It is a known idiom that Linux Kernel and C++ are not compatible.
- In this commit an important feature to invoke VIRGO Linux Kernel from userspace python libraries via two alternatives
have been added.
- In one alternative, C++ boost::python extensions have been added to encapsulate access to VIRGO memory system calls -
virgo_malloc(), virgo_set(), virgo_get(), virgo_free(). Initial testing reveals that C++ and Kernel are not too
incompatible and all the VIRGO memory system calls work well though initially there were some errors because of config
issues.
- In the other alternative, C Python extensions have been added that replicate boost::python extensions above in C - C
Python with Linux kernel
works exceedingly well compared to boost::python.
- This functionality is required when there is a need to set kernel analytics configuration variables learnt by AsFer
Machine Learning Code
dynamically without re-reading /etc/virgo_kernel_analytics.conf.
- This completes a major integration step of NeuronRain suite - request travel roundtrip to-and-fro top level machine-
learning C++/python
code and rock-bottom C linux kernel - bull tamed ;-).
- This kind of python access to device drivers is available for Graphics Drivers already on linux (GPIO - for accessing
device states)
- logs for both C++ and C paths have been added in cpp_boost_python_extensions/ and cpython_extensions.
- top level python scripts to access VIRGO kernel system calls have been added in both directories:
        CPython - python cpython_extensions/asferpythonextensions.py
        C++ Boost::Python - python cpp_boost_python_extensions/asferpythonextensions.py
- .so, .o files with build commandlines(asferpythonextensions.build.out) for "python setup.py build" have been added
in build lib and temp directories.
- main implementations for C++ and C are in cpp_boost_python_extensions/asferpythonextensions.cpp and
cpython_extensions/asferpythonextensions.c

```
------------------------------------------------------------------------
Commits as on 12 February 2016
------------------------------------------------------------------------
1017. Commits for Telnet/System Call Interface to VIRGO CPUPooling -> VIRGO Queue -> KingCobra
------------------------------------------------------------------------
*) This was commented earlier for the past few years due to a serious kernel panic in previous kernel versions - <= 3.15.5
*) In 4.1.5 a deadlock between VIRGO CPUPooling and VIRGO queue driver init was causing following error in
"use_as_kingcobra_service" clause :
        - "gave up waiting for virgo_queue init, unknown symbol push_request()"
*) To address this a new boolean flag to selectively enable and disable VIRGO Queue kernel service mode
"virgo_queue_reactor_service_mode" has been added.
*) With this flag VIRGO Queue is both a kernel service driver and a standalone exporter of function symbols -
push_request/pop_request
*) Incoming request data from telnet/virgo_clone() system call into cpupooling kernel service reactor pattern (virgo
cpupooling listener loop) is treated as generic string and handed over to VIRGO queue and KingCobra which publishes it.
*) This resolves a long standing deadlock above between VIRGO cpupooling "use_as_kingcobra_service" clause and VIRGO queue
init.
*) This makes virgo_clone() systemcall/telnet both synchronous and asynchronous - requests from telnet client/virgo_clone()
system call can be either synchronous RPC functions executed on a remote cloud node in kernelspace (or) an asynchronous
invocation through "use_as_kingcobra_service" clause path to VIRGO Queue driver which enqueues the data in kernel workqueue
and subsequently popped by KingCobra.
*) Above saves an additional code implementation for virgo_queue syscall paths - virgo_clone() handles, based on config
selected, incoming data passed to it either as a remote procedure call or as a data that is pushed to VIRGO Queue/KingCobra
pub-sub kernelspace.
------------------------------------------------------------------------
Prerequisites:
--------------
- insmod kingcobra_main_kernelspace.ko
- insmod virgo_queue.ko compiled with flag virgo_queue_reactor_service_mode=1
        (when virgo_queue_reactor_service_mode=0, listens on port 60000 for direct telnet requests)
- insmod virgo_cloud_test_kernelspace.ko
- insmod virgo_cloudexec.ko (listens on port 10000)


------------------------------------------------------------------------
Schematic Diagram
------------------------------------------------------------------------
VIRGO clone system call/telnet client ---> VIRGO cpupooling(compiled with use_as_kingcobra_service=1) ------> VIRGO Queue
kernel service (compiled with virgo_queue_reactor_service_mode=1) ---> Linux Workqueue handler ------> KingCobra


------------------------------------------------------------------------
1018. Commits as on 15 February 2016 - Kernel Analytics - VIRGO Linux Kernelwide imports
------------------------------------------------------------------------
- Imported Kernel Analytics variables into CloudFS kernel module - printed in driver init()
- Module.symvers from kernel_analytics has been merged with CloudFS Module.symvers
```

- Logs for above has been added in cloudfs/test_logs/
- Makefile updated with correct fs path
- Copyleft notices updated


---------------------------------------------------------------------------
1019. Commits as on 15 February 2016 - Kernel Analytics - VIRGO Linux Kernelwide imports
---------------------------------------------------------------------------
- Kernel Analytics driver exported variables have been imported in CPU virtualization driver
- Module.symvers from kernel_analytics has been merged with Module.symvers in cpupooling
- kern.log for this import added to cpupooling/virgocloudexec/test_logs/


---------------------------------------------------------------------------
1020. Commits as on 15 February 2016 - Kernel Analytics - VIRGO Linux Kernelwide imports
---------------------------------------------------------------------------
- Imported kernel analytics variables into memory virtualization driver init() , exported from kernel_analytics driver
- build shell script updated
- logs added to test_logs/
- Module.symvers from kernel_analytics has been merged with memory driver Module.symvers
- Makefile updated


---------------------------------------------------------------------------
1021. Commits as on 15 February 2016 - Kernel Analytics - VIRGO Linux Kernelwide imports
---------------------------------------------------------------------------
- Imported kernel analytics variables into VIRGO Queueing Driver
- logs for this added in test_logs/
- Makefile updated
- Module.symvers from kernel_analytics has been merged with Queueing driver's Module.symvers
- .ko, .o and build generated sources


--------------------------------------------------------------------
Commits as on 16,17 February 2016
--------------------------------------------------------------------
1022. (FEATURE-DONE) Socket Buffer Debug Utility Function - uses linux skbuff facility
--------------------------------------------------------------------
- In this commit a multipurpose socket buffer debug utility function has been added in utils driver and exported
  kernelwide.
- It takes a socket as function argument does the following:
        - dereference the socket buffer head of skbuff per-socket transmit data queue
        - allocate skbuff with alloc_skb()
        - reserve head room with skb_reserve()
        - get a pointer to data payload with skb_put()
        - memcpy() an example const char* to skbuff data
        - Iterate through the linked list of skbuff queue in socket and print headroom and data pointers
        - This can be used as a packet sniffer anywhere within VIRGO linux network stack

- Any skb_*() functions can be plugged-in here as deemed necessary.
- kern.log(s) which print the socket internal skbuff data have been added to a new testlogs/ directory
- .cmd files generated by kbuild


--------------------------------------------------------------------
1023. (FEATURE-DONE) Commits as on 24 February 2016
--------------------------------------------------------------------
skbuff debug function in utils/ driver:
(*) Added an if clause to check NULLity of skbuff headroom before doing skb_alloc()
(*) kern.log for this commit has been added testlogs/
(*) Rebuilt kernel objects and sources


--------------------------------------------------------------------
Commits as on 29 February 2016
--------------------------------------------------------------------
------------------------------------------------------------------------
771. (FEATURE-DONE) Software Analytics - SATURN Program Analysis added to VIRGO Linux kernel drivers - this section is an
extended draft on respective topics in NeuronRain AstroInfer design -  https://github.com/shrinivaasanka/asfer-github-
code/blob/master/asfer-docs/AstroInferDesign.txt
------------------------------------------------------------------------
- SATURN (saturn.stanford.edu) Program Analysis and Verification software has been
integrated into VIRGO Kernel as a Verification+SoftwareAnalytics subsystem
- A sample driver that can invoke an exported function has been added in drivers - saturn_program_analysis
- Detailed document for an example null pointer analysis usecase has been created in virgo-
docs/VIRGO_SATURN_Program_Analysis_Integration.txt
- linux-kernel-extensions/drivers/virgo/saturn_program_analysis/saturn_program_analysis_trees/error.txt is the error report
from SATURN
- SATURN generated preproc and trees are in linux-kernel-extensions/drivers/virgo/saturn_program_analysis/preproc and
linux-kernel-extensions/drivers/virgo/saturn_program_analysis/saturn_program_analysis_trees/


--------------------------------------------------------------------
1024. Commits as on 10 March 2016
--------------------------------------------------------------------
- SATURN analysis databases (.db) for locking, memory and CFG analysis.
- DOT and PNG files for locking, memory and CFG analysis.
- new folder saturn_calypso_files/ has been added in saturn_program_analysis/ with new .clp files virgosaturncfg.clp and
virgosaturnmemory.clp
- SATURN alias analysis .db files


------------------------------------------------------------------------------
-----------------
1025.(FEATURE-DONE) NEURONRAIN - ASFER Commits for VIRGO - CloudFS systems calls integrated into Boost::Python C++ and
Python CAPI invocations
--------------------------------------------------------------------

```
---------------
------------------------------------------------------------------------------------------
AsFer Commits as on 30 May 2016
------------------------------------------------------------------------------------------
VIRGO CloudFS system calls have been added (invoked by unique number from syscall_32.tbl) for C++ Boost::Python interface
to
VIRGO Linux System Calls. Switch clause with a boolean flag has been introduced to select either VIRGO memory or filesystem
calls.
kern.log and CloudFS textfile Logs for VIRGO memory and filesystem invocations from AsFer python have been committed to
testlogs/

------------------------------------------------------------------------------------------
1026. AsFer Commits as on 31 May 2016
------------------------------------------------------------------------------------------
Python CAPI interface to NEURONRAIN VIRGO Linux System Calls has been updated to include File System open, read, write
primitives also.
Rebuilt extension binaries, kern.logs and example appended text file have been committed to testlogs/. This is exactly
similar to
commits done for Boost::Python C++ interface. Switch clause has been added to select memory or filesystem VIRGO syscalls.

------------------------------------------------------------------------------------------
(BUG - STABILITY ISSUES) Commits - 25 July 2016 - Static Analysis of VIRGO Linux kernel for investigating heisencrashes
------------------------------------------------------------------------------------------
Initial Documentation for Smatch and Coccinelle kernel static analyzers executed on VIRGO Linux kernel - to be updated
periodically with further analysis.

--------------------------------------------------------------------------------------------
--
(BUG - STABILITY ISSUES) Commits - 1 August 2016 - VIRGO Linux Stability Issues - Ongoing Random oops and panics
investigation
--------------------------------------------------------------------------------------------
--
1. GFP_KERNEL has been replaced with GFP_ATOMIC flags in kmem allocations.
2. NULL checks have been introduced in lot of places involving strcpy, strcat, strcmp etc., to circumvent
buffer overflows.
3. Though this has stabilized the driver to some extent, still there are OOPS in unrelated places deep
with in kernel where paging datastructures are accessed - kmalloc somehow corrupts paging
4. OOPS are debugged via gdb as:
        4.1 gdb ./vmlinux /proc/kcore
        or
        4.2 gdb <loadable_kernel_module>.o
    followed by
        4.3 l *(address+offset in OOPS dump)
5. kern.log(s) for the above have been committed in tar.gz format and have numerous OOPS occurred during repetitive telnet
```

and syscall
invocation(boost::python C++) invocations of virgo memory system calls.
6. Paging related OOPS look like an offshoot of set_fs() encompassing the filp_open VFS calls.


-----------------------------------------------------------------------------------------
(BUG-STABILITY ISSUES) Commits - 26 September 2016 - Ongoing Random Panic investigation
-----------------------------------------------------------------------------------------
Further analysis on direct VIRGO memory cache primitives telnet invocation - problems are similar
to Boost::Python AsFer VIRGO system calls invocations.


-----------------------------------------------------------------------------------------
(BUG-STABILITY ISSUES) Commits - 27 September 2016 - Ongoing Random Panic investigation
-----------------------------------------------------------------------------------------
Analysis of VIRGO memory cache primitives reveal more inconsistencies in cacheline flushes between CPU and GPU.


-----------------------------------------------------------------------------------------
1027. Commits - 20 March 2017 and 21 March 2017 - VIRGO Linux 64-bit build based on 4.10.3 kernel
-----------------------------------------------------------------------------------------
*) moved virgoeventnetclient_driver_build.sh to virgoutils_driver_build.sh in utils/ driver
*) Updated VIRGO Linux Build Steps for 4.10.3
*) New repository has been created for 64-bit VIRGO Linux kernel based on 4.10.3 mainline kernel in GitHub and imported in
SourceForge:
        https://github.com/shrinivaasanka/virgo64-linux-github-code
        https://sourceforge.net/p/virgo64-linux/
*) Though it could have been branched off from existing VIRGO repository (32-bit) which is based on 4.1.5 mainline kernel,
creating a
separate repository for 64-bit 4.10.3 VIRGO kernel code was simpler because:
        - there have been directory path changes for syscall entries in 4.10.3 and some other KBuild entities
        - Some script changes done for 4.1.5 in modpost and vmlinux phases are not required
        - having two VIRGO branches one with 4.1.5 code and 32-bit driver .ko binaries and other with 4.10.3 code and 64-
bit driver .ko
binaries could be unmanageable and commits could go into wrong branch
        - 4.10.3 64-bit VIRGO kernel build is still in experimental phase and it is not known if 64-bit 4.10.3 build solves
earlier panic
problems in 4.1.5
        - If necessary one of these two repositories could be made branch of the other later


-----------------------------------------------------------------------------------------
1028. Commits - 27 March 2017 Ongoing analysis of VIRGO 64 bit linux kernel based on 4.10.3 kernel mainline
-----------------------------------------------------------------------------------------
*) Prima facie, 64 bit kernel is quite finicky and importunate compared to 32 bit and 64 bit specific idiosyncrasies are to
the fore.
*) During the past 1 week, quite a few variants of kernel and drivers builds were tried with KASAN enabled and without
KASAN (Google Kernel Address Sanitizer)

 *) KASAN shows quite huge number of user memory accesses which later translate to panics.
 *) Most nagging of these was kernel_recvmsg() panic.
 *) Added and updated skbuff socket debug utility driver with a new debug function and to print more fields of skbuff
 *) KASAN was complaining about _asan_load8 (loading 8 userspace bytes)
 *) All erroneous return data types in VIRGO mempool ops structure have been corrected in VIRGO headers
 *) all type casts have been sanitized
 *) Changed all kernel stack allocations to kernel heap kzallocs
 *) This later caused a crash in inet_sendmsg in kernel_sendmsg()
 *) gdb64 disassemble showed a trapping instruction:
testb $0x6,0x91(%14) with corresponding source line:
sg = !!(sk->sk_route_caps & NETIF_F_SG)
in tcp_sendmsg() (net/ipv4/tcp.c)
 *) changed kernel_sendmsg() to sock->ops->sendmsg()
 *) These commits are still ongoing analysis only.
 *) Screenshots for these have been added to debug-info/


--------------------------------------------------------------------------------------------------
----
1029. Continued analysis of VIRGO 64-bit linux kernel built on 4.10.3 mainline - Commits - 30 March 2017
--------------------------------------------------------------------------------------------------
----
 *) Previous commit was crashing inside tcp_sendmsg()
 *) GDB64 disassembly shows NULL values for register R12 which is added with an offset 91 and is an operand in testb
 *) Protected all kernel_sendmsg() and kernel_recvmsg() in both system calls side and drivers side with
        oldfs=get_fs(), set_ds(KERNEL_DS) and set_fs(oldfs)
blocks without which there are random kernel_sendmsg and kernel_recvmsg hangs
 *) Removed init_net and sock_create_kern usage everywhere and replaced them with sock_create calls
 *) Tried MSG_FASTOPEN flags but it does not help much in resolving tcp_sendmsg() NULL pointer dereference issue.
MSG_FASTOPEN just
speedsup the message delivery by piggybacking the message payload before complete handshake is established(SYN, SYN-ACK,
SYN-ACK-ACK) in
SYN-ACK itself. But eventually it has to be enabled as fast open is becoming a standard.
 *) Kasan reports have been enabled.
 *) Added more debug code in skbuff debug utility functions in utils driver to check if sk->prot is a problem.
 *) Replaced kernel_sendmsg with a sock->ops->sendmsg() in mempool sendto function which otherwise crashes in tcp_sendmsg().
 *) With sock->ops->sendmsg() systemcalls <-----> drivers two-way request-reply works but still there are random -32 (broken
pipe) and -104
(Connection Reset by Peer) errors
 *) Logs for working sys_virgo_malloc() call with correctly returned VIRGO Unique ID for memory allocated has been committed
to test_logs in
virgocloudexecmempool
 *) sock->ops->sendmsg() in mempool driver sendto function requires a MSG_NOSIGNAL flag which prevents SIGPIPE signal though
not fully
 *) Reason for random broken pipe and connection reset by peer errors in mempool sendto is unknown. Both sides have

connections open and
there is no noticeable traffic.
*) While socket communications in 32 bit VIRGO kernel syscalls and drivers work with no issues, why 64-bit has so many
hurdles is puzzling.
Reasons could be 64 bit address alignment issues, 64 bit specific #ifdefs in kernel code flow, major changes from 4.1.5 to
4.10.3 kernels etc.,
*) NULL values for register R12 indicate already freed skbuff data which are accessed/double-freed. Kernel TCP engine has a
circular linked list of skbuff write queue which is iterated in skbuff utils driver debug functions.
*) TCP engine clones the head data of skbuff queue, transmits it and waits for an ACK or timeout. Data is freed only if ACK
or timeout occurs.
And head of the queue is advanced to next element in write queue and this continues till write queue is empty waiting for
more messages.
*) If ACK is not received, head data is cloned again and retransmitted by sequence number flow control.


-------------------------------------------------------------------------------
1030. Continued Analysis of VIRGO 64 bit based on 4.10.3 linux kernel - Commits - 1 April 2017, 3 April 2017
-------------------------------------------------------------------------------
*) kernel_sendmsg() has been replaced with sock->ops->sendmsg() because
kernel_sendmsg() is quite erratic in 4.10.3 64 bit
*) There were connection reset errors in system calls side for virgo_malloc/. This was probably because
sock->ops->sendmsg() requires MSG_DONTWAIT and MSG_NOSIGNAL flags and sendmsg does not block.
*) sock release happens and virgo_malloc syscalls receives -104 error
*) Temporarily sock_release has been commented. Rather socket timeout should be relied upon which should
do automatic release of socket resources
*) Similar flags have been applied in virgo_malloc syscalls too.
*) Logs with above changes do not have reset errors as earlier.
*) virgo set/get still crashes because 64 bit id is truncated which would require data type changes for
64 bit
*) test_virgo_malloc test case has been rebuilt with -m64 flag for invocation of 64 bit syscalls by
numeric ids


-------------------------------------------------------------------------------------
1031. Continued Analysis of VIRGO 64 bit 4.10.3 kernel - commits - 10 April 2017
-------------------------------------------------------------------------------------
*) There is something seriously wrong with 4.10.3 kernel sockets in 64 bit build VIRGO send/recv messages and even
accept/listen too.
*) All kernel socket functionalities which work well in 4.1.5 32 bit VIRGO , have random hangs, panics in 4.10.3 VIRGO64
build mostly
in inet_recvmsg/sendmsg code path
*) KASAN shows attempts to access user address which occurs despite set_fs(KERNEL_DS)
*) Crash stack is similar to previous crashes in tcp_sendmsg()
*) Tried different address and protocol families for kernel socket accept (TCP,UDP,RAW sockets)
*) With Datagram sockets, kernel_listen() mysteriously fails with -95 error in kernel_bind(operation not supported)
*) With RAW sockets, kernel_listen() fails with -93 error for AF_PACKET (protocol not supported)

*) tcpdump pcap sniffer doesn't show anything unruly.
*) This could either be a problem with kernel build (unlikely), Kbuild .config or could have extraneous reasons. But
.config for
4.1.5 and 4.10.3 are similar.
*) Only major difference between 4.10.3 and 4.1.5 is init_net added in sock_create_kern() internally
*) datatype of VIRGO Unique ID has been changed to unsigned long long (_u64)
*) tried with INADDR_LOOPBACK in place of INADDR_ANY
*) also tried with disabled multi(homing) in /etc/hosts.conf
*) Above random kernel socket hangs occur across all VIRGO system calls and drivers transport.
*) Utils kernel socket client to EventNet kernel service also has similar inet_recvmsg/inet_sendmsg panic problems.


--------------------------------------------------------------------------------
1032. Commits - 11 April 2017 - EventNet and Utils Drivers 64bit
--------------------------------------------------------------------------------
*) EventNet driver works in 64 bit VIRGO Linux
*) An example eventnet logging with utils virgo_eventnet_log() works now without tcp_sendmsg() related stalls in previous
builds
*) Return Datatypes for all EventNet operations have been sanitized (struct socket* was returned as int in 32 build and
reinterpret-cast to
struct socket*. This reinterpret cast does not work in 64 bit) in eventnet header.
*) utils eventnet log in init() has been updated with a meaningful edge update message
*) kern.log for this has been added to eventnet/testlogs


--------------------------------------------------------------------------------
1033. Commits - 17 April 2017 - VIRGO64 Memory, CPU, FileSystem, EventNet kernel module drivers
--------------------------------------------------------------------------------
*) telnet requests to VIRGO memory(kernelmemcache), cpu and filesystem modules work after resolving issues with return
value types
*) commented le32_to_cpu() and print_buffer() which was suppressing lot of log messages.
*) VIRGO <driver> ops structures have been updated with correct datatypes.
*) reinterpret cast of struct socket* to int has been completely done away with which could have caused 64bit specific
panics
*) lot of kern.log(s) and screen captures have been added for telnet requests in testlogs/ of respective <driver> directory
*) Prima facie 64bit telnet requests to VIRGO module listeners are relatively stabler than 32bit
*) Previous code changes should be relevant to 32 bit VIRGO kernel too.
*) tcp_sendmsg()/tcp_recvmsg() related hangs could be mostly related to corrupted skbuff queue within each socket.
*) This is because replacing kernel_<send/recv>msg() with sock_<send/recv>msg() causes return value to be 0 while
socket release crashes within skbuff related kernel functions.
*) To make socket state immutable, in VIRGO memory driver header files, client socket has been declared as const type.


--------------------------------------------------------------------------------
1034. Commits - KingCobra 64 bit and VIRGO Queue + KingCobra telnet requests - 17 April 2017
--------------------------------------------------------------------------------
*) Rebuilt KingCobra 64bit kernel module

*) telnet requests to VIRGO64 Queueing module listener driver are serviced by KingCobra servicerequest
*) Request_Reply queue persisted for this VIRGO Queue + KingCobra routing has been committed to c-src/testlogs.
*) kern.log for this routing has been committed in VIRGO64 queueing directory
*) Similar to other drivers struct socket* reinterpret cast to int has been removed and has been made const in queuesvc
kernel thread


-----------------------------------------------------------------------------------------------------------------
1035. Commits - VIRGO64 system calls - kernel module listeners - testcases and system calls updates - 18 April 2017
-----------------------------------------------------------------------------------------------------------------
*) All testcases have been rebuilt
*) VIRGO kernel memcache,cpu and filesystem system calls have been updated with set_fs()/get_fs() blocks for
kernel_sendmsg()
and kernel_recvmsg()
*) Of these virgo_clone() system call testcase (test_virgo_clone) works flawlessly and there are no
tcp_sendmsg()/tcp_recvmsg() related
kernel panics.
*) VIRGO memcache and filesystem system call testcases have usual tcp_sendmsg()/tcp_recvmsg() despite the kernel socket
code
being similar to VIRGO clone system call
*) Logs for VIRGO clone system call to CPU kernel driver module have been committed to virgo_clone/test/testlogs


-----------------------------------------------------------------------------------------------------------------
--
1036. Commits - VIRGO64 Kernel MemCache and FileSystem system calls to VIRGO Memory and FileSystem Drivers - 19 April 2017
-----------------------------------------------------------------------------------------------------------------
--
*) Changed iovec in virgo_clone.c to kvec
*) test_virgo_filesystem.c and test_virgo_malloc.c VIRGO system calls testcases have been changed with some additional
printf(s) in userspace
*) virgo_malloc.c has been updated with BUF_SIZE in iov_len and memset to zero initialize the buffer.
tcp_sendmsg()/tcp_recvmsg() pair were
getting stuck in copy_from_iter_full() memcpy with a NULL Dereference. memcpy() was reading past the buffer bound causing
an overrun. strlen()
didnot work for iov_len.
*) virgo_fs.c virgo_write() memcpy has been changed back to copy_from_user() thereby restoring status quo ante (commented
more than 3 years ago
because of a kernel panic in older versions of 32 bit VIRGO kernel)
*) Logs for VIRGO kmemcache and filesystem system calls have been committed to respective system call directories.
*) With this all VIRGO64 functionalities work in both telnet and system calls requests routes end-to-end from clients to
kernel modules with
kernel sockets issues resolved fully.
*) Major findings are:
- VIRGO 4.10.3 64 bit kernel is very much stable compared to 32 bit 4.1.5 kernel
- there are no i915 related errors which happened in VIRGO 32 bit 4.1.5 kernel

- Repetitive telnet and system calls requests to VIRGO modules are stable and there are no kernel panics like 4.1.5 32 bit
kernel.
- Google Kernel Address Sanitizer is quite helpful in finding stack overruns, null derefs, user memory accesses etc.,
- 64 bit kernel is visibly faster than 32 bit.
- Virgo Unique ID is now extended to 2^64 with unsigned long long.

------------------------------------------------------------------------------------------------------
1037. Commits - VIRGO64 memory and filesystem calls to memory and filesystem drivers requests routing - 20 April 2017
------------------------------------------------------------------------------------------------------
*) Changed return value of virgo_cloud_free_kernelspace() to a string literal "kernel memory freed"
*) Logs for VIRGO64 memory and filesystem calls to memory and filesystem drivers requests routing have been committed in
test_logs of
both driver directories


-----------------------------------------------------------------------------------------------
1038. Commits - 27 April 2017
-----------------------------------------------------------------------------------------------
Residual logs for VIRGO 64 bit 4.10.3 kernel committed.


-----------------------------------------------------------------------------------------------
1039. Commits - 25 May 2017
-----------------------------------------------------------------------------------------------
*) Changed LOOPBACK to INADDR_ANY for VIRGO64 kernel memcache listen port
*) All VIRGO64 RPC, kernel memcache, cloud filesystem primitives have been retested
*) VIRGO64 mempool binaries have been rebuilt


-------------------------------------------------------------------------------------------------
1040. Commits - 31 August 2017 - NeuronRain ReadTheDocs Documentation - VIRGO64 System calls and Drivers
(http://neuronrain-documentation.readthedocs.io/en/latest/)
-------------------------------------------------------------------------------------------------
(*) New directory systemcalls_drivers/ has been added to virgo-docs/ and representative VIRGO64
system calls and drivers functionality logs have been committed for demonstration purpose.
(*) VIRGO64 cloudfs driver has been rebuilt after changing virgofstest.txt file creation filp_open() call
(*) Screenshots and logs for VIRGO64 Clone, Kernel MemCache and Cloud FS SystemCalls-Drivers interaction, socket transport
debug messages and Kernel Address Sanitizer have been committed in virgo-docs/systemcalls_drivers


---------------------------------------------------------------------------------------------------------
------------
1041. Commits - 23 September 2017 - Major VIRGO mainline kernel version Upgrade for Kernel Transport Layer Security -
4.10.3 to 4.13.3
---------------------------------------------------------------------------------------------------------
------------
(*) Recently released mainline kernel version 4.13 integrates SSL/TLS into kernelspace- KTLS - for the first time.
(*) KTLS is a standalone kernel module af_ktls (https://github.com/ktls) implemented by RedHat and Facebook for optimizing

SSL handshake
within kernelspace itself and reduce userspace-kernelspace switches.
(*) sendfile() system call in linux which is used for file transmission (combining read+write) from one fd to another uses
this
KTLS optimization in kernelspace in af_ktls codebase (af_ktls tool)
(*) VIRGO Linux kernel fork-off requires this kernelspace TLS functionality to fully secure traffic from system call client
to remote
 cloud node's kernel module listeners
(*) Hence VIRGO64 linux kernel mainline base is urgently upgraded from 4.10.3 to 4.13.3
(*) All system calls and kernel module code in VIRGO64 now have #include(s) for tls.h and invoke kernel_setsockopt() on the
client-server
kernelspace sockets for SOL_TLS and TLS_TX options and have been rebuilt.
(*) VIRGO64 RPC clone/kmemcache/cloudfs system calls to kernel module listeners have been tested with this new KTLS socket
option,
on rebuilt VIRGO64 kernel overlay-ed on 4.13.3 64-bit linux kernel
(*) 4.13 mainline kernel also has SMB CIFS bug fixes for recent malware attacks (WannaCry etc.,) which further ensures
security of
VIRGO64 linux fork-off kernelspace traffic.
(*) New buildscript for 4.13.3 linux kernel has been committed
(*) testlogs for VIRGO64 system calls and driver listeners KTLS transport have been committed in virgo-
docs/systemcalls_drivers/kern.log.VIRGO64_SystemCalls_Drivers.4.13.3_KTLS_kernelsockets.22September2017
(*) After this upgrade, complete system calls to driver listener traffic is SSL enabled implicitly.
(*) Updated kernel object files for 4.13.3 build are part of this commit.

------------------------------------------------------------------------------------
1042. Commits - Remnant commits for 4.13.3 upgrade - 24 September 2017
------------------------------------------------------------------------------------
Updated init.h and syscalls.h headers for virgo system calls


--------------------------------------------------------------------------------------------------
----------------
1043. Commits - VIRGO64 4.13.3 KTLS Upgrade - System Calls-Driver Listeners End-to-End encrypted traffic testing - 25
September 2017
--------------------------------------------------------------------------------------------------
----------------
(*) VIRGO64 CPU/KMemCache/CloudFS system calls have been invoked by userspace testcases and all primitives work after
4.13.3 KTLS upgrade
(*) Some small modifications to system calls code have been made and rebuilt to remove redundant iovbuf variables in
printk(s)
(*) test_virgo_filesystem.c testcase has been updated and rebuilt
(*) kern.log(s) for CPU/KMemCache/CloudFS systemcalls to driver listeners invocations have been committed to respective
system call
directories
(*) virgofstest.txt written to by virgo_write() has also been committed. But a weird behaviour is still observed similar to

previous linux kernel versions (4.1.5 and 4.10.3): Repetitive invocations are required to flush the filesystem buffer to
force write the file.
(*) No DRM GEM i915 panics are observed and stability of VIRGO64 + 4.13.3 linux kernel is more or equal to VIRGO64 + 4.10.3
linux kernel


-----------------------------------------------------------------------------------------------------------------------
----------------
1044. Commits - VIRGO64 VIRGO_KTLS branch creation and rebase of master to previous commit - 30 September 2017
-----------------------------------------------------------------------------------------------------------------------
----------------
(*) New branch VIRGO_KTLS has been created after previous commit on 25 September 2017 and all 5 commits after 25 September
2017 till
28 September 2017 have been branched to VIRGO_KTLS (which has the #ifdef for crypto_info, reads from /etc/virgo_ktls.conf
and a new ktls
driver module)
(*) Following are the commit hashes and commandlines in GitHub and SourceForge:
                    git branch -b VIRGO_KTLS
                    git branch master
                    git rebase -i <SHA1_on_25September2017>
                    git rebase --continue
                    git commit --amend
                    git push --force
----------------------------------------------------------------------------------------------
1958  git checkout -b
1959  git checkout -b VIRGO_KTLS
1960  ls
1961  git checkout VIRGO_KTLS
1962  git push origin VIRGO_KTLS
1963  git status
1964  git checkout
1965  git checkout -b
1966  git branch
1967  git branch master
1968  git branch -h
1969  git branch
1970  git checkout master
1975  git checkout -b
1976  git checkout -b VIRGO_KTLS
1979  git push origin VIRGO_KTLS
1990  git rebase -i bb661e908cba2a5357414e89166f29086a28bdf0
1991  git status
1992  git rebase -i bb661e908cba2a5357414e89166f29086a28bdf0
1996  git rebase --continue
1997  git commit --amend

```
2019   git rebase -i bb661e908cba2a5357414e89166f29086a28bdf0
2029   git rebase --continue
2037   git push --force
2091   git rebase -i bb661e908cba2a5357414e89166f29086a28bdf0
2092   git rebase --continue
2093   git commit --amend
2094   git push --force
2110   git branch
2111   git branch master
2112   git checkout master
2113   git branch
2114   git rebase -i e76b4089633223f610fddc0e0eaff8c2cef8b9f1
2115   git commit --amend
2116   git rebase --continue
2117   git push --force
```
---------------------------------------------------------------------------------------
KTLS in 4.13.3 has support for only private symmetric encryption. It does not support Public Key Encryption yet. Since this might take a while
mainstream VIRGO64 code might change a lot for other features. Therefore, VIRGO_KTLS specific crypto_info code has been branched off and would parallelly evolve if PKI features are available on KTLS in next versions of Linux kernel.


------------------------------------------------------------------------------------
1045. Commits - 1 October 2017
------------------------------------------------------------------------------------
kern.log(s) for VIRGO64 systemcalls-driver 4.13.3 64-bit upgrade tests on master branch after reversal and rebase of master HEAD yesterday for
branching to VIRGO_KTLS. There is a weird General Protection Fault in intel atomic commit work not seen thus far. Also -32 and -107 socket errors are frequent after reversal though code remains same. But all virgo clone/kmemcache/fs systemcalls functionalities work in invocations after GPF.


------------------------------------------------------------------------------------
------
1046. Commits - VIRGO64 Utils and EventNet Drivers Update for tcp_sendmsg() stack out-of-bounds error - 3 October 2017
------------------------------------------------------------------------------------
------
(*) Utils Generic Socket Client function virgo_eventnet_log() for EventNet kernel module listener was repeatedly failing in kernel_connect()
emitting -32 and -107 errors.
(*) kernel_connect() was guarded by set_fs() and get_fs() memory segment routines to prevent any memory corruption. After this stack out-of-bounds error was reported by kernel address sanitizer in tcp_sendmsg() and copy_from_iter_full() implying the message buffer was not properly read within kernel.
(*) After replacing strlen(buf) by BUF_SIZE in msg flags before kernel_connect() stack out-of-bounds error has been remedied and Utils to EventNet virgo_eventnet_log() doesn't crash in tcp_sendmsg()
(*) kern.log for this has been committed in drivers/virgo/utils/testlogs/

(*) Both eventnet and utils drivers have been rebuilt


-------------------------------------------------------------------------------------------------
----------
1047. VIRGO64 system calls-drivers on linux kernel 4.13.3 - miscellaneous bugfixes - 5 October 2017
-------------------------------------------------------------------------------------------------
----------
(*) kernel_setsockopt() for KTLS has been commented in all system calls and drivers because KTLS functionality has been
branched to
VIRGO_KTLS
(*) In virgo_clone.c, iov.iov_len has been set to BUF_SIZE
(*) kernel_connect() has been guarded by set_fs()/get_ds() in VIRGO64 system call clients
(*) test_virgo_malloc.c testcase has been updated
(*) There was a weird problem in in4_pton(): sin_addr.saddr was not set correctly from string IP address and this was
randomly occurring only in virgo_set()
(*) in4_pton() is implemented in net/core/utils.c and reads the string IP address digits and sums up the ASCII values to
hex representation of the address. Bitwise operations within this functions were failing randomly.
(*) Repeated builds were done trying different possible fixes but didn't work e.g casting saddr to (u8*)
(*) There is an alternative in_aton() function which takes only String IP address and returns address as __be32
(*) After in_aton() in virgo_set() random faulty address conversion does not occur - in_aton() is differently implemented
compared to in4_pton()
(*) msg_hdr has been initialized to NULL in virgo_set()
(*) Lot of debug printk()s have been added
(*) kern.log (.tar.gz) for RPC clone/KMemCache/Filesystem systemcalls-driver has been committed to virgo-
docs/systemcalls_drivers
(*) VIRGO Linux build steps have been updated for example commandlines to overlay mainline kernel tree by VIRGO64 source

commit 4e6681ade4ddbf1bed17f7c115b59a5ebf884256
Author: K.Srinivasan <ka.shrinivaasan@gmail.com>
Date:   Fri Oct 6 11:36:15 2017 +0530


-------------------------------------------------------------------------------------------------
---------
1048. VIRGO64 Queueing Kernel Module Listener - KingCobra64 - 4.13.3 - 6 October 2017
-------------------------------------------------------------------------------------------------
---------
(*) telnet client connection to VIRGO64 Queue and a subsequent workqueue routing (pub/sub) to KingCobra64 has been tested
on 4.13.3
(*) TX_TLS socket option has not been disabled and is a no-op because it has no effect on the socket.
(*) REQUEST_REPLY.queue for this routing from VIRGO64 queue and persisted by KingCobra64 has been committed to KingCobra64
repositories in GitHub and SourceForge

commit d4e95b58474838d65da9c69944c6287acbdfe72c
Author: K.Srinivasan <ka.shrinivaasan@gmail.com>

```
Date:    Fri Oct 6 11:05:21 2017 +0530
```

----------------------------------------------------------------------------------------
-------
1049. VIRGO64 System Calls to Drivers and Telnet Client to Drivers on 4.13.3 linux kernel - master branch (after KTLS has been reverted
and branched to VIRGO_KTLS) - test case logs - 6 October 2017
----------------------------------------------------------------------------------------
-------
(*) VIRGO64 System calls - Clone, KMemCache and Filesystem system call primitives to Driver listeners invocations have been been tested
by respective test_<systemcall> unit testcases
(*) VIRGO64 Telnet Clients to Driver listeners invocations have been tested by telnet connections
(*) Master branches in SourceForge and GitHub VIRGO64 do not have KTLS provisions. Only VIRGO_KTLS branch has crypto_info and setsockopt()
 for TX_TLS for kernel sockets.
(*) It has been already mentioned in NeuronRain Documentation in https://neuronrain-documentation.readthedocs.io/en/latest/ on securing
VIRGO cloud nodes in the absence of KTLS - most obvious solution is to install VPN client-servers in all nodes which create Virtual IPs
 on a secure tunnel (e.g OpenVPN).
(*) VIRGO64 system call clients and driver listeners should read these Virtual IPs from /etc/virgo_client.conf and /etc/virgo_cloud.conf
  and cloud traffic is confined to the VPN tunnel.


----------------------------------------------------------------------------------------
----------
1050. VIRGO64 SystemCalls-Drivers endtoend invocations unit case tests - on 4.13.3 - VIRGO64 main branch - 11 October 2017
----------------------------------------------------------------------------------------
----------
(*) VIRGO64 systemcalls have been invoked from unit test cases (test_<system_call>) in a loop of few hundred iterations
(*) No DRM GEM i915 panics or random crashes are observed and stability is good
(*) This is probably the first loop iterative testing of VIRGO system calls and drivers.
(*) Kernel logs for this have been committed to virgo-docs/systemcalls_drivers directory.
(*) Note on concurrency: Presently mutexing within system calls have been commented because in past linux versions mutexing
within kernel was causing strange panic issues. As a design choice and feature-stability tradeoff (stability is more
important than introducing additional code) mutexing has been lifted up to userspace. It is upto the user applications
invoking the system calls to synchronize multiple user threads invoking VIRGO64 system calls i.e VIRGO64 system calls are
not re-entrant. This would allow just one kernel thread (mapped 1:1 to a user thread)
 to execute in kernel space. Mostly this is relevant only to kmemcache system calls which have global in-kernel-memory
address translation tables and next_id variable. VIRGO clone/filesystem calls do not have global in-kernel-memory
datastructures.


----------------------------------------------------------------------------------------

```
--------------
1051. VIRGO64 SystemCalls-Drivers concurrent invocations - 2 processes having shared mutex - 14 October 2017
--------------------------------------------------------------------------------------------------------
--------------
(*) VIRGO64 systemcalls are invoked in a function which is called from 2 processes concurrently
(*) Mutexes between the processes are PTHREAD_PROCESS_SHARED attribute set.
(*) test_virgo_malloc.c unit testcase has been enhanced to fork() a process and invoke systemcalls in a function for 100
iterations each
(*) Logs for the Virgo Unique IDs malloc/set/get/free in the systemcalls side and kern.logs for the drivers have been
committed to test/testlogs/
(*) No DRM GEM i915 crashes were observed
(*) test_virgo_malloc.c testcase demonstrates the coarse grained critical section lock/unlock for kmemcache systemcalls and
is a template
that should be followed for any userspace application.


--------------------------------------------------------------------------------------------------------
--------------
775. (FEATURE) VIRGO64 Kernel Analytics - Streaming Implementation - 13 December 2017 - this section is an extended draft
on respective topics in NeuronRain AstroInfer design -  https://github.com/shrinivaasanka/asfer-github-
code/blob/master/asfer-docs/AstroInferDesign.txt
--------------------------------------------------------------------------------------------------------
--------------
(#) Presently kernel analytics config have to be read from a file storage. This is a huge performance bottleneck when
frequency of
analytics variables written to is realtime. For example autonomous vehicles/drones write gigabytes of navigation data in
few minutes.
(#) Because of this /etc/virgo_kernel_analytics.conf grows tremendously. File I/O in linux kernel module is also fragile
and not recommended.
(#) Previous latency limitations necessitate an alternative high performance analytics config variable read algorithm
(#) This commit introduces new streaming kernel analytics config reading function - It connects to a kernel analytics
streaming IP address
on hardcoded port 64000 and reads analytics key-value pairs in an infinite loop.
(#) These read key-value pairs are stored in a kernel global ring buffer exported symbol (by modulus operator). Because of
circular buffer, older kernel analytics variables are overwritten repetitively.
(#) kernel socket message flags are set to MSG_MORE | MSG_FASTOPEN | MSG_NOSIGNAL for high response time. MSG_FASTOPEN
works with no panic
in 4.13.3 64-bit which was a problem in previous kernel versions.
(#) kern.log for this has been committed to kernel_analytics/testlogs/
(#) include/linux/virgo_kernel_analytics.h header file has been updated for declarations related to streaming analytics.
(#) Webserver used for this is netcat started on port 64000 as:
                nc -l 64000
                >k1=v1
                >k2=v2
                ...
```

------------------------------------------------------------------------------
1052. VIRGO64 Kernel Analytics - Reading Stream of Analytic Variables made a kernel thread - 13 December 2017
------------------------------------------------------------------------------
(#) This is sequel to previous commit for Stream reading Kernel Analytics variables over a network socket
(#) read_streaming_virgo_kernel_analytics_conf() function is invoked in a separate kernel thread because module init is
blocked
(#) VIRGO64 config module was loaded and exported kernel analytics variables read over socket by previous spin-off kernel
thread are
imported in VIRGO64 config init.
(#) kern.log for this has been committed to testlogs/
(#) Pre-requisite: Webservice serving kernel_analytics variables must be started before kernel_analytics module is loaded
in kernel.
(#) By this a minimum facility for live reading analytics anywhere on cloud (it can be userspace or kernelspace) and
exporting them
to modules on a local cloud node kernel has been achieved - ideal for cloud-analytics-driven IoT


------------------------------------------------------------------------------
---------
1053. VIRGO64 System Calls - Drivers - Kernel Analytics Streaming - on 4.13.3 kernel - 15 December 2017
------------------------------------------------------------------------------
---------
(#) VIRGO64 System Calls to Drivers invocations on 4.13.3 kernel have been executed after enabling streaming kernel
analytics
(#) VIRGO64 RPC/KMemCache/CloudFS Drivers import, streamed variable-value pairs exported from kernel_analytics read from
netcat webservice
(#) VIRGO64 KMemCache testcase has 2 concurrent processes invoking kememcache systemcalls in a loop.
(#) kern.log for this has been committed to virgo-docs/systemcalls_drivers
(#) virgofstest.txt written by CloudFS systemcalls-drivers invocation is also committed to virgo-docs/systemcalls_drivers


------------------------------------------------------------------------------
---------
1054. VIRGO64 system calls and drivers - Quadcore 64bit - Known Issues - documentation only - 26 May 2019
------------------------------------------------------------------------------
---------
1. VIRGO64 system calls to drivers interactions so far have been tested only on dual core 64 bit architecture.
2. In quad core 64 bit there have been random -32,-107, -101 errors in kernel_connect() from system call clients to driver
services in
almost all three types of VIRGO64 system calls - clone/kmemcache/filesystem - to respective drivers
3. These errors do not occur if following in4_pton() invocation is changed to in_aton() before kernel_connect() in system
call clients and kernel is rebuilt with following change before kernel_connect():

        /*in4_pton(vaddr->hstprt->hostip, strlen(vaddr->hstprt->hostip), (u8*)&sin.sin_addr.s_addr, '\0',NULL);*/
        sin.sin_addr.s_addr=in_aton(vaddr->hstprt->hostip);

4. Since this problem occurs erratically and only on quadcore 64-bit and reasons for these random -32,-101,-107 errors are
still unknown, no commit for this code change has been made and this issue is left as documented known issue.
5. Most likely the u8* cast causes client socket address corruption.
6. Because of random -32,-101,-107 errors, in quadcore, system calls sometimes do not transmit client side commandline
primitive strings to driver services.

--------------------------------------------------------------------------------------------------------------------------
-----------
773. (FEATURE) Linux Kernel 5.1.4, PXRC Drone/UAV/Flight Controller, UVC Video WebCam driver, Kernel Analytics, NeuronRain
Usecases - 28 May 2019 - this section is an extended draft on respective topics in NeuronRain AstroInfer design -
https://github.com/shrinivaasanka/asfer-github-code/blob/master/asfer-docs/AstroInferDesign.txt
--------------------------------------------------------------------------------------------------------------------------
-----------
1. Two analytics usecases mentioned in NeuronRainUsecases.txt in NeuronRain AsFer asfer-docs/ have been illustrated in this
commit by
2 example drivers for PXRC Drone controller Driver and UVC Video WebCam Driver.
2. PXRC Phoenix RC flight controller is part of linux kernel from 4.17 and kernel major version has been bumped to 5.x.x
recently.
3. Linux kernel 5.1.4 has been built by a new build script - buildscript_5.1.4.sh.
4. Linux kernel 5.1.4 has recent versions of PXRC drone controller driver and a UVC video webcam
driver (http://www.ideasonboard.org/uvc/faq/)
5. New directory linux-kernel-5.1.4-extensions/ has been created for VIRGO64 code built on kernel mainline version 5.1.4.
No branch is created
for version 5.1.4 because pxrc driver is part of kernel only from 4.17 while code in linux-kernel-extensions/ is based on
kernel 4.13.3 for
dual core 64-bit architecture.
6. New VIRGO64 build on 5.1.4 kernel is necessary only for PXRC, UVC and kernel_analytics drivers while other VIRGO64
drivers have not been
ported to 5.1.4 and are still on 4.13.3 kernel.
7. Two drivers for PXRC and UVC Webcam in 5.1.4 have been committed under linux-kernel-5.1.4-
extensions/drivers/media/usb/uvc and linux-kernel-5.1.4-extensions/drivers/input/joystick
8. VIRGO64 kernel_analytics driver for 5.1.4 has been committed under linux-kernel-5.1.4-
extensions/drivers/virgo/kernel_analytics.
9. Porting VIRGO64 kernel_analytics driver from 4.13.3 to 5.1.4 required changing vfs_read() of
/etc/virgo_kernel_analytics.conf to
kernel_read() in config file read.
10. Drivers code for PXRC is in drivers/input/joystick/pxrc.c has been instrumented with few printk() statements that print
the
virgo_kernel_analytics_conf array variable-value pairs exported by VIRGO64 kernel_analytics driver. Kernel analytics
variables are
imported by #include of virgo_config.h
11. Drivers code for UVC Video WebCam is in drivers/media/usb/uvc. File drivers/media/usb/uvc/uvc_video.c has been
instrumented with lot of

uvc_trace() statements which print kernel_analytics driver exported analytics variable "match" and its boolean value.
Kernel analytics variables are imported  by #include of virgo_config.h. Variable "match" has been set to "True" assuming a
face recognition or retinal scan match by userspace analytics and /etc/virgo_kernel_analytics.conf has been written to.
12. UVC Video WebCam traces are enabled by:
          echo 0xffff > /sys/module/uvcvideo/parameters/trace
13. Example kern.log(s) for PXRC and UVC drivers are committed under
drivers/input/joystick/testlogs/kern.log.pxrc.28May2019 and drivers/media/usb/uvc/kern.log.uvcvideo.28May2019 which show
UVC traces printing the imported kernel analytics variable "match=True" and PXRC driver registration. No other PXRC traces
are printed because of lack of drones and drone licensing dependency
14. Both UVC and PXRC drivers, VIRGO64 kernel_analytics driver and linux kernel 5.1.4 have been built on quadcore 64-bit
architecture.
15. This example import of VIRGO64 kernel analytics variables into PXRC drone and UVC webcam drivers demonstrate passing of
userspace analytics information to kernel for suitable action.
16. Driver build shell scripts have been committed to UVC and PXRC driver directories.


-----------------------------------------------------------------------------------------------------------------------
------------------
774. (FEATURE) Concurrent Managed Workqueue(CMWQ), VIRGO64 Queueing and KingCobra64 messaging - 12 June 2019 - this section
is an extended draft on respective topics in NeuronRain AstroInfer design and KingCobra design -
https://github.com/shrinivaasanka/asfer-github-code/blob/master/asfer-
docs/AstroInferDesign.txt,https://github.com/shrinivaasanka/kingcobra64-github-code/blob/master/KingCobraDesignNotes.txt
-----------------------------------------------------------------------------------------------------------------------
------------------
1. Existing workqueue underneath VIRGO64 queueing and requests routed by it to KingCobra64 messaging are old legacy
workqueues which
have been revamped to Concurrent Managed Workqueue which supports concurrent messaging and lot of other options in queue
creation.
2. create_workqueue() in VIRGO64 Queueing has been changed to alloc_workqueue() of Concurrent Managed Workqueue.
3. VIRGO64 Queueing request routing to KingCobra64 messaging has been tested with CMWQ and queueing log and kingcobra64
Request-Reply
Queue have been committed to respective testlogs of the drivers
4. reading from stream has been disabled in virgo_kernel_analytics.h
5. Reference - CMWQ documentation - https://www.kernel.org/doc/html/v4.11/core-api/workqueue.html
6. Byzantine Fault Tolerance in KingCobra64 persisted queue can be made available by performant CMWQ and routing to the
Replicas of
REQUEST_REPLY.queue by any of the practical BFT protocols available.
7. Most important application of CMWQ based VIRGO64-KingCobra64 is in the context of kernelspace hardware messaging in
IoT,Drones and other
analytics driven embedded systems.
8. An example usecase which is a mix of sync and async I/O in kernelspace:
          (*) Analytics Variables computed by userspace machine learning are read over socket stream by kernel_analytics
driver and
exported kernelwide
          (*) Some interested Drone driver in kernel (example PXRC) reads the analytics variables synchronously and sends

reply messages asynchronously to VIRGO64 Queuing driver over kernel sockets.
          (*) VIRGO Queuing routes the queued messages to KingCobra64 driver

 Srinivasan Kannan (alias) Ka.Shrinivaasan (alias) Shrinivas Kannan
 http://sites.google.com/site/kuja27