



()

[Articles \(/articles/\)](#)[Cloud Storage \(/cloud-storage/\)](#)[Business Model \(/business-model/\)](#)[Home \(/\)](#) / [Browse \(/directory/\)](#) / [virgo64 \(/p/virgo64-linux/\)](#) / [Code](#)

(/blog)

utm_source=sourceforge&utm_medium=navbar&utm_campaign=homepage) ▼



virgo64 (/p/virgo64-linux/)

64 bit VIRGO linux kernel - derived from virgo-linux-github-code

Brought to you by: ka_shrinivaasan (/u/userid-769929/)

[\[332956\] /p/virgo64-linux/code/ci/332956872938453739be889383006120ec8337620-762log/p/virgo-docs/VirgoDesign.txt](#)[Download this file \(?format=raw\)](#)

1622 lines (1349 with data), 137.1 kB

```
1 /*****
2 #-----
3 #NEURONRAIN VIRGO - Cloud, Machine Learning and Queue augmented Linux Kernel Fork-off
4 #This program is free software: you can redistribute it and/or modify
5 #it under the terms of the GNU General Public License as published by
6 #the Free Software Foundation, either version 3 of the License, or
7 #at your option) any later version.
8 #This program is distributed in the hope that it will be useful,
9 #but WITHOUT ANY WARRANTY; without even the implied warranty of
10 #MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
11 #GNU General Public License for more details.
12 #You should have received a copy of the GNU General Public License
13 #along with this program. If not, see <http://www.gnu.org/licenses/>.
14 #-----
15 #Copyleft (Copyright+):
16 #Srinivasan Kannan (alias) Ka.Shrinivaasan (alias) Shrinivas Kannan
17 #Ph: 9791499106, 9003082186
18 #Krishna iResearch Open Source Products Profiles:
19 #http://sourceforge.net/users/ka_shrinivaasan,
20 #https://github.com/shrinivaasanka,
21 #https://www.openhub.net/accounts/ka_shrinivaasan
22 #Personal website(research): https://sites.google.com/site/kuja27/
23 #emails: ka.shrinivaasan@gmail.com, shrinivas.kannan@gmail.com,
24 #kashrinivaasan@live.com
25 #-----
26 *****/
27
28 VIRGO is an operating system kernel forked off from Linux kernel mainline to add cloud functionalities (system calls,
29 Remote Device Invocation , which is an old terminology for Internet-Of-Things has already been experimented in SunRPC a
30
31 Memory pooling:
32 -----
33 Memory pooling is proposed to be implemented by a new virgo_malloc() system call that transparently allocates a block
34
35 CPU pooling or cloud ability in a system call:
36 -----
37 Clone() system call is linux specific and internally it invokes sys_clone(). All fork(),vfork() and clone() system cal
38
39 virgo_clone() is a wrapper over clone() that looks up a map of machines-to-loadfactor and get the host with least load
40
41 Kernel has support for kernel space sockets with kernel_accept(), kernel_bind(), kernel_connect(), kernel_sendmsg() an
42
43 Experimental Prototype
44 -----
45 virgo_clone() system call and a kernel module virgocloudexec which implements Sun RPC interface have been implemented.
46
47 VIRGO - loadbalancer to get the host:ip of the least loaded node
48 -----
49 Loadbalancer option 1 - Centralized loadbalancer registry that tracks load:
50 -----
51
52 Virgo_clone() system call needs to lookup a registry or map of host-to-load and get the least loaded host:ip from it.
53
54 Many application level userspace load monitoring tools are available but as virgo_clone() is in kernel space, it needs
55
56
57
```

(Design notes for LB option 1 handwritten by myself are at :<http://sourceforge.net/p/virgo-linux/code-0/HEAD/tree/trunk>)

Loadbalancer option 2 - Linux Psuedorandom number generator based load balancer(experimental) instead of centralized r

Each virgo_clone() client has a PRG which is queried (/dev/random or /dev/urandom) to get the id of the host to send t
Expected number of requests per node is derived as:

expected number of requests per node = summation(each_value_for_the_random_variable_for_number_of_requests * probabili
=expected number of requests per node = (math.pow(N, k+2) - k*math.pow(N,2) + k*math.pow(N,1) - 1) / (math.pow(N, k+3))

This loadbalancer is dependent on efficacy of the PRG and since each request is uniformly, identically, independently
would distribute requests evenly. This obviates the need for loadtracking and coherency of the load-to-host table.

(Design notes for LB option 2 handwritten by myself at :<http://sourceforge.net/p/virgo-linux/code-0/HEAD/tree/trunk/vi>)

(python script in virgo-python-src/)

Implemented VIRGO Linux components (as on 7 March 2016)

1. cpupooling virtualization - VIRGO_clone() system call and VIRGO cpupooling driver by which a remote procedure can b
2. memorypooling virtualization - VIRGO_malloc(), VIRGO_get(), VIRGO_set(), VIRGO_free() system calls and VIRGO memory
3. filesystem virtualization - VIRGO_open(), VIRGO_read(), VIRGO_write(), VIRGO_close() system calls and VIRGO cloud f
4. config - VIRGO config driver for configuration symbols export.
5. queueing - VIRGO Queueing driver kernel service for queueing incoming requests, handle them with workqueue and invoke
6. cloudsync - kernel module for synchronization primitives (Bakery algorithm etc.,) with exported symbols that can be
7. utils - utility driver that exports miscellaneous kernel functions that can be used across VIRGO Linux kernel
8. EventNet - eventnet kernel driver to vfs_read()/vfs_write() text files for EventNet vertex and edge messages (port:
9. Kernel_Analytics - kernel module that reads machine-learned config key-value pairs set in /etc/virgo_kernel_analytic
10. Testcases and kern.log testlogs for the above
11. SATURN program analysis wrapper driver.

Thus VIRGO Linux at present implements a minimum cloud OS (with cloud-wide cpu, memory and file system management) ove

VIRGO ToDo and NiceToHave Features (list is quite dynamic and might be rewritten depending on feasibility - longterm w

(FEATURE - DONE-minimum separate config file support in client and kernel service)1. More Sophisticated VIRGO config

(FEATURE - Special case implementation DONE) 2. Object Marshalling and Unmarshalling (Serialization) Features - Featur

(FEATURE - DONE) 3. Virgo_malloc(), virgo_set(), virgo_get() and virgo_free() syscalls that virtualize the physical me

Initial Design Handwritten notes committed at: <http://sourceforge.net/p/virgo-linux/code-0/210/tree/trunk/virgo-docs/v>

(FEATURE - DONE) 4. Integrated testing of AsFer-VIRGO Linux Kernel request roundtrip - invocation of VIRGO linux kerne

4.1 Schematic Diagram:

```

AsFer Python -----> Boost::Python C++ Extension -----> VIRGO memory system calls -----> VIRGO Linux Kernel
/\
|
|
-----<-----
AsFer Python -----> CPython Extensions -----> VIRGO memory system calls -----> VIRGO Linux Kernel Memory D
/\
|
|
-----<-----

```

(FEATURE - DONE)5. Multithreading of VIRGO clouddex kernel module (if not already done by kernel module subsystem int

(FEATURE - DONE) 6. Sophisticated queueing and persistence of CPU and Memory pooling requests in Kernel Side (by possib

(FEATURE - DONE-Minimum Functionality) 7. Integration of Asfer(AstroInfer) algorithm codes into VIRGO which would add

Example scenario 1 without implementation:

- Philips Hue IoT mobile app controlled bulb - <http://www2.meethue.com/en-xx/>
- kernel_analytics module learns key-value pairs from the AsFer code and exports it VIRGO kernel wide
- A driver function with in bulb embedded device driver can be invoked through VIRGO cpupooling (invoked from remote v

based on if-else clause of the kernel_analytics variable i.e remote_client invokes virgo_clone() with function argumen

Example scenario 2 without implementation:

- A swivel security camera driver is remotely invoked via virgo_clone() in the VIRGO cloud.
- The camera driver uses a machine learnt variable exported by kernel_analytics-and-AsFer to pan the camera by how muc

Example scenario 3 without implementation - probably one of the best applications of NeuronRain IoT OS:

- Autonomous Driverless Automobiles - a VIRGO driver for a vehicle which learns kernel analytics variables (driving di
- AsFer analytics receives obstacle distance data 360+360 degrees (vertical and horizontal) around the vehicle
- VIRGO Linux kernel on vehicle has two special drivers for Gear-Clutch-Break-Accelerator-Fuel(GCBAF) and Stee
- AsFer analytics with high frequency computes threshold variables for applying break, clutch, gear, velocity,
- These analytics variables are continuously read by GCBAF and Steering drivers which autopilot the vehicle.
- Above applies to Fly-by-wire aeronautics too with appropriate changes in analytics variables computed.
- The crucial parameter is the response time in variable computation and table updates which requires a huge c

E.g. Autopilot in Tesla Cars processes Petabytes of information (Smooth-as-Silk algorithm) from sensors which are fed

References for Machine Learning + Linux Kernel

- 7.1 KernTune - [http://repository.uwc.ac.za/xmlui/bitstream/handle/10566/53/Yi_KernTune\(2007\).pdf?sequence=3](http://repository.uwc.ac.za/xmlui/bitstream/handle/10566/53/Yi_KernTune(2007).pdf?sequence=3)
- 7.2 Self-learning, Predictive Systems - <https://icri-ci.technion.ac.il/projects/past-projects/machine-learning-for-arc>
- 7.3 Linux Process Scheduling and Machine Learning - <http://www.cs.ucr.edu/~kishore/papers/tencon.pdf>
- 7.4 Network Latency and Machine Learning - https://users.soe.ucsc.edu/~slukin/rtt_paper.pdf
- 7.5 Machine Learning based Meta-Scheduler for Multicore processors - <https://books.google.co.in/books?id=1GwCmCrl0QC&>

8. A Symmetric Multi Processing subsystem Scheduler that virtualizes all nodes in cloud (probably this would involve i

(FEATURE - ONGOING) 9. Virgo is an effort to virtualize the cloud as a single machine - Here cloud is not limited to s

(FEATURE - DONE) 10. Memory Pooling Subsystem Driver - Virgo_malloc(), Virgo_set(), Virgo_get() and Virgo_free() syste

(FEATURE - DONE) 11. Virgo Cloud File System with virgo_cloud_open(), virgo_cloud_read() , virgo_cloud_write() and vir

(FEATURE - DONE) 12. VIRGO Cloud File System commands through syscall paths - virgo_open(),virgo_close(),virgo_read()

(FEATURE - DONE) 13. VIRGO memory pooling feature is also a distributed key-value store similar to other prominent key

14. VIRGO memory pooling can be improved with disk persistence for in-memory key-value store using virgo_malloc(),virg

15. (FEATURE-DONE) Socket Debugging, Program Analysis and Verification features for user code that can find bugs stati

16(FEATURE - DONE-Minimum Functionality). Operating System Logfile analysis using Machine Learning code in AstroInfer

17. Implementations of prototypical Software Transactional Memory and LockFree Datastructures for VIRGO memory pooling

18. Scalability features for Multicore machines - references:

(<http://halobates.de/lk09-scalability.pdf>, <http://pdos.csail.mit.edu/papers/linux/osdi10.pdf>)

19. Read-Copy-Update algorithm implementation for VIRGO memory pooling that supports multiple simultaneous versions of

20. (FEATURE - SATURN integration - minimum functionality DONE) Program Comprehension features as an add-on described

21. (FEATURE - DONE) Bakery Algorithm implementation - cloudsync kernel module

22. (FEATURE - ONGOING) Implementation of Distributed Systems primitives for VIRGO cloud viz., Logical Clocks, Termina

23. (FEATURE - minimum functionality DONE) Enhancements to kmem if it makes sense, because it is better to rely on vir Kernel Malloc syscall kmalloc() internally works as follows:

- kmem_cache_t object has pointers to 3 lists
- These 3 lists are full objects SLAB list, partial objects SLAB list and free objects SLAB list - all are lis
- and cache_cache is the global list of all caches created thus far.
- Any kmalloc() allocation searches partial objects SLAB list and allocates a memory block with kmem_cache_all
- Any kfree() returns an object to a free SLAB list
- Full SLABs are removed from partial SLAB list and appended to full SLAB list
- SLABs are virtual memory pages created with kmem_cache_create
- Each SLAB in SLABs list has blocks of similar sized objects (e.g. multiples of two). Closest matching block

KERNELSPACE:

VIRGO address translation table already implements a tree registry of vtables each of capacity 3000 that keep track of USERSPACE: sbrk() and brk() are no longer used internally in malloc() library routines. Instead mmap() has replaced it

24.(FEATURE - ONGOING) Cleanup the code and remove unnecessary comments.

25.(FEATURE - DONE) Documentation - This design document is also a documentation for commit notes and other build and

26. (FEATURE - DONE) Telnet path to virgo_cloud_malloc,virgo_cloud_set and virgo_cloud_get has been tested and working

27. Augment the Linux kernel workqueue implementation (<http://lxr.free-electrons.com/source/kernel/workqueue.c>) with d

28.(FEATURE - DONE) VIRGO queue driver with native userspace queue and kernel workqueue-handler framework that is opti

29.(FEATURE - DONE) KERNELSPACE EXECUTION ACROSS CLOUD NODES which geographically distribute userspace and kernelspace a logical abstraction for a cloudwide virtualized kernel:

Remote Cloud Node Client

```

220         (cpupooling, eventnet, memorypooling, cloudfs, queueing - telnet and syscalls clients)
221         |
222         |
223 (Userspace) |
224         |-----Kernel Sockets-----> Remote Clou
225                                         (VIRGO cpupooling, memorypooling, cloudfs, que
226                                         |
227                                         |
228                                         | (Ke
229                                         |
230                                         V
231         <-----Kernel Sockets-----
232         |
233         |
234 (Userspace) |
235         |
236
237
238
239 30. (FEATURE - DONE) VIRGO platform as on 5 May 2014 implements a minimum set of features and kernelsocket commands re
240
241 31. (FEATURE - DONE) VIRGO Queue standalone kernel service has been implemented in addition to paths in schematics abo
242
243 VIRGO Queue client(e.g telnet) -----> VIRGO Queue kernel service ---> Linux Workqueue handler -----> KingCobra
244
245 32. (FEATURE - DONE) EventNet kernel module service:
246 VIRGO eventnet client (telnet) -----> VIRGO EventNet kernel service -----> EventNet graph text files
247
248 33. (FEATURE - DONE) Related to point 22 - Reuse EventNet cloudwide logical time infinite graph in AsFer in place of L
249
250 34. (FEATURE - OPTIONAL) The kernel modules services listening on ports could return a JSON response when connected in
251
252 35. (FEATURE-Minimum Functionality DONE) Pointer Swizzling and Unswizzling of VIRGO addressspace pointers to/from VIRG
253
254 *****
255                         CODE COMMIT RELATED NOTES
256 *****
257
258 VIRGO code commits as on 16/05/2013
259 -----
260 1. VIRGO clouddriver with a listener kernel thread service has been implemented and it listens on port 10000 on s
261 through /etc/modules load-on-bootup facility
262
263 2. VIRGO clouddriver virgo_clone() system call has been implemented that would kernel_connect() to the VIRGO clouddriver s
264 port 10000
265
266 3. VIRGO clouddriver has been split into virgo.h (VIRGO typedefs), virgoclouddriver.h(VIRGO clouddriver service th
267 module_init() of VIRGO clouddriver) and virgo_clouddriver.c (with module ops definitions)
268
269 4. VIRGO does not implement SUN RPC interface anymore and now has its own virgo ops.
270
271 5. Lot of Kbuild related commits with commented lines for future use have been done viz., to integrate VIRGO to Kbuild
272
273 VIRGO code commits as on 20/05/2013
274 -----
275 1. test_virgo_clone.c testcase for sys_virgo_clone() system call works and connections are established to VIRGO cloude
276
277 2. Makefile for test_virgo_clone.c and updated buildscript.sh for headers_install for custom-built linux.
278
279 VIRGO code commits as on 6/6/2013
280 -----
281 1. Message header related bug fixes
282
283 VIRGO code commits as on 25/6/2013
284 -----
285 1.telnet to kernel service was tested and found working
286 2.GFP_KERNEL changed to GFP_ATOMIC in VIRGO clouddriver kernel service
287
288 VIRGO code commits as on 1/7/2013
289 -----
290 1. Instead of printing iovec, printing buffer correctly prints the messages
291 2. wake_up_process() added and function received from virgo_clone() syscall is executed with kernel_thread and results
292 virgo_clone() syscall client.
293
294
295 commit as on 03/07/2013
296 -----
297 PRG loadbalancer preliminary code implemented. More work to be done
298
299 commit as on 10/07/2013
300 -----

```

```

301 Tested PRG loadbalancer read config code through telnet and virgo_clone. VFS code to read from virgo_cloud.conf commen
302
303 commits as on 12/07/2013
304 -----
305 PRG loadbalancer prototype has been completed and tested with test_virgo_clone and telnet and symbol export errors and
306
307 commits as on 16/07/2013
308 -----
309 read_virgo_config() and read_virgo_clone_config()(replica of read_virgo_config()) have been implemented and tested to
310 all nodes). Thus minimal cloud functionality with config file support is in place. Todo things include function point
311
312 commits as on 17/07/2013
313 -----
314 moved read_virgo_config() to VIRGOcloudexec's module_init so that config is read at boot time and exported symbols are
315 Also commented read_virgo_clone_config() as it is redundant
316
317 commits as on 23/07/2013
318 -----
319
320 Lack of reflection kind of facilities requires map of function_names to pointers_to_functions to be executed
321 on cloud has to be lookedup in the map to get pointer to function. This map is not scalable if number of functions are
322 in millions and size of the map increases linearly. Also having it in memory is both CPU and memory intensive.
323 Moreover this map has to be synchronized in all nodes for coherency and consistency which is another intensive task.
324 Thus name to pointer function table is at present not implemented. Suitable way to call a function by name of the func
325 is yet to be found out and references in this topic are scarce.
326
327 If parameterIsExecutable is set to 1 the data received from virgo_clone() is not a function but name of executable
328 This executable is then run on usermode using call_usermodehelper() which internally takes care of queueing the workst
329 and executes the binary as child of keventd and reaps silently. Thus workqueue component of kernel is indirectly made
330 This is sometimes more flexible alternative that executes a binary itself on cloud and
331 is preferable to clone()ing a function on cloud. Virgo_clone() syscall client or telnet needs to send the message with
332
333 If parameterIsExecutable is set to 0 then data received from virgo_clone() is name of a function and is executed in el
334 using dlsym() lookup and pthread_create() in user space. This unifies both call_usermodehelper() and creating a usersp
335 with a fixed binary which is same for any function. The dlsym lookup requires mangled function names which need to be
336 virgo_clone or telnet. This is far more efficient than a function pointer table.
337
338 call_usermodehelper() Kernel upcall to usermode to exec a fixed binary that would inturn execute the cloneFunction in
339 by spawning a pthread. cloneFunction is name of the function and not binary. This clone function will be dlsym()ed
340 and a pthread will be created by the fixed binary. Name of the fixed binary is hardcoded herein as
341 "virgo_kernelupcall_plugin". This fixed binary takes clone function as argument. For testing libvirgo.so has been crea
342 virgo_cloud_test.c and separate build script to build the cloud function binaries has been added.
343
344 - Ka.Shrinivaasan (alias) Shrinivas Kannan (alias) Srinivasan Kannan
345 (https://sites.google.com/site/kuja27)
346
347 commits as on 24/07/2013
348 -----
349
350 test_virgo_clone unit test case updated with mangled function name to be sent to remote cloud node. Tested with test_v
351 end-to-end and all features are working. But sometimes kernel_connect hangs randomly (this was observed only today and
352 to blocking vs non-blocking problem. Origin unknown).
353
354 - Ka.Shrinivaasan (alias) Shrinivas Kannan (alias) Srinivasan Kannan
355 (https://sites.google.com/site/kuja27)
356
357 commits as on 29/07/2013
358 -----
359
360 Added kernel mode execution in the clone_func and created a sample kernel_thread for a cloud function. Some File IO lo
361 binaries and parameterIsExecutable has been moved to virgo.h
362
363 commits as on 30/07/2013
364 -----
365
366 New usecase virgo_cloud_test_kernelspace.ko kernel module has been added. This exports a function virgo_cloud_test_ker
367 accessed by virgo_cloudeexec kernel service to spawn a kernel thread that is executed in kernel addressspace. This Kerne
368 on cloud adds a unique ability to VIRGO cloud platform to seamlessly integrate hardware devices on to cloud and transp
369 to them from a remote cloud node through virgo_clone().
370
371 Thus above feature adds power to VIRGO cloud to make it act as a single "logical device driver" though devices are in
372
373 commits as on 01/08/2013 and 02/08/2013
374 -----
375
376 Added Bash shell commandline with -c option for call_usermodehelper upcall clauses to pass in remote virgo_clone comma
377 arguments to it. Also tried output redirection but it works some times that too with a fatal kernel panic.
378
379 Ideal solutions are :
380 1. either to do a copy_from_user() for message buffer from user address space (or)
381 2. somehow rebuild the kernel with fd_install() pointing stdout to a VFS file* struct. In older kernels like 2.6.x, th
382 with in kmmod.c (___call_usermodehelper()) which has been redesigned in kernel 3.x versions and fd_install has been rem
383 3. Create a Netlink socket listener in userspace and send message up from kernel Netlink socket.

```

```

382
383 All the above are quite intensive and time consuming to implement. Moreover doing FileIO in usermode helper is strongly
384
385 Since Objective of VIRGO is to virtualize the cloud as single execution "machine", doing an upcall (which would run wi
386 redundant often and kernel mode execution is sufficient. Kernel mode execution with intermodule function invocation ca
387 the entire board in remote machine (since it can access PCI bus, RAM and all other device cards)
388
389 As a longterm design goal, VIRGO can be implemented as a separate protocol itself and sk_buff packet payload from remo
390 can be parsed by kernel service and kernel_thread can be created for the message.
391
392 commits as on 05/08/2013:
393 -----
394 Major commits done for kernel upcall usermode output logging with fd_install redirection to a VFS file. With this it h
395
396 11 August 2013:
397 -----
398 Open Source Design and Academic Research Notes uploaded to http://sourceforge.net/projects/acadpdrafts/files/Miscellan
399
400
401 commits as on 23 August 2013
402 -----
403 New Multithreading Feature added for VIRGO Kernel Service - action item 5 in ToDo list above (virgo_cloudexec driver m
404
405 commits as on 1 September 2013
406 -----
407 GNU Copyright license and Product Owner Profile (for identity of license issuer) have been committed. Also Virgo Memor
408
409 commits as on 14 September 2013
410 -----
411 Updated virgo malloc design handwritten nodes on kmalloc() and malloc() usage in kernelspace and userspace execution m
412
413 -----
414 VIRGO virtual addressing
415 -----
416 VIRGO virtual address is defined with the following datatype:
417
418 struct virgo_address
419 {
420     int node_id;
421     void* addr;
422 };
423
424 VIRGO address translation table is defined with following datatype:
425
426 struct virgo_addr_transtable
427 {
428     int node_id;
429     void* addr;
430 };
431
432 -----
433 VIRGO memory pooling prototypical implementation
434 -----
435 VIRGO memory pooling implementation as per the design notes committed as above is to be implemented as a prototype und
436 under drivers/virgo/memorypooling and $LINUX_SRC_ROOT/virgo_malloc. But the underlying code is more or less similar to
437
438 virgo_malloc() and related syscalls and virgo mempool driver connect to and listen on port different from cpupooling d
439
440 Commits as on 17 September 2013
441 -----
442 Initial untested prototype code - virgo_malloc and virgo mempool driver - for VIRGO Memory Pooling has been committed
443
444 Commits as on 19 September 2013
445 -----
446 3.7.8 Kernel full build done and compilation errors in VIRGO malloc and mempool driver code and more functions code ad
447
448 Commits as on 23 September 2013
449 -----
450 Updated virgo_malloc.c with two functions, int_to_str() and addr_to_str(), using kmalloc() with full kernel re-build.
451 (Rather a re-re-build because some source file updates in previous build got deleted somehow mysteriously. This could
452
453 Commits as on 24 September 2013
454 -----
455 Updated syscall*.tbl files, staging.sh, Makefiles for virgo_malloc(), virgo_set(), virgo_get() and virgo_free() memory p
456
457 Commits as on 25 September 2013
458 -----
459 All build related errors fixed after kernel rebuild some changes made to function names to reflect their
460 names specific to memory pooling. Updated /etc/modules also has been committed to repository.
461
462 Commits as on 26 September 2013

```



```
463 -----
464 Circular dependency error in standalone build of cpu pooling and memory pooling drivers fixed and
465 datatypes and declarations for CPU pooling and Memory Pooling drivers have been segregated into respective header file
466 virgo_mempool.h with corresponding service header files) to avoid any dependency error.
467
468 Commits as on 27 September 2013
469 -----
470 Major commits for Memory Pooling Driver listen port change and parsing VIRGO memory pooling commands have been done.
471
472 Commits as on 30 September 2013
473 -----
474 New parser functions added for parameter parsing and initial testing on virgo_malloc() works with telnet client with l
475
476 Commits as on 1 October 2013
477 -----
478 Removed strcpy in virgo_malloc as ongoing bugfix for buffer truncation in syscall path.
479
480 Commits as on 7 October 2013
481 -----
482 Fixed the buffer truncation error from virgo_malloc syscall to mempool driver service which was caused by
483 sizeof() for a char*. BUF_SIZE is now used for size in both syscall client and mempool kernel service.
484
485 Commits as on 9 October 2013 and 10 October 2013
486 -----
487 Mempool driver kernelspace virgo mempool ops have been rewritten due to lack of facilities to return a
488 value from kernel thread function. Since mempool service already spawns a kthread, this seems to be sufficient. Also t
489 causes the kernel socket to block as it waits for more data to be sent.
490
491 Commits as on 11 October 2013
492 -----
493 sscanf format error for virgo_cloud_malloc() return pointer address and sock_release() null pointer exception has been
494 Added str_to_addr() utility function.
495
496 Commits as on 14 October 2013 and 15 October 2013
497 -----
498 Updated todo list.
499
500 Rewritten virgo_cloud_malloc() syscall with:
501 - mutexed virgo_cloud_malloc() loop
502 - redefined virgo address translation table in virgo_mempool.h
503 - str_to_addr(): removed (void**) cast due to null sscanf though it should have worked
504
505 Commits as on 18 October 2013
506 -----
507 Continued debugging of null sscanf - added str_to_addr2() which uses simple_strtoll() kernel function
508 for scanning pointer as long long from string and casting it to void*. Also more %p qualifiers where
509 added in str_to_addr() for debugging.
510
511 Based on latest test_virgo_malloc run, simple_strtoll() correctly parses the address string into a long long base 16 a
512
513 Commits as on 21 October 2013
514 -----
515 Kern.log for testing after vtranstable addr fix with simple_strtoll() added to repository and still the other %p quali
516
517 Commits as on 24 October 2013
518 -----
519 Lot of bugfixes made to virgo_malloc.c for scanning address into VIRGO transtable and size computation. Testcase test_
520
521 Though the above sys_virgo_malloc() works, the return value is a kernel pointer if the virgo_malloc executes in the Ke
522
523 Commits as on 25 October 2013
524 -----
525 virgo_malloc.c has been rewritten by adding a userspace __user pointer to virgo_get() and virgo_set() syscalls which a
526
527 Commits as on 29 October 2013
528 -----
529 Miscellaneous ongoing bugfixes for virgo_set() syscall error in copy_from_user().
530
531 Commits as on 2 November 2013
532 -----
533 Due to an issue which corrupts the kernel memory, presently telnet path to VIRGO mempool driver has been
534 tested after commits on 31 October 2013 and 1 November 2013 and is working but again there is an issue in kstrtoul() t
535 data to set.
536
537 Commits as on 6 November 2013
538 -----
539 New parser function virgo_parse_integer() has been added to virgo_cloud_mempool_kernelspace driver module which is car
540 lib/kstrtoux.c and modified locally to add an if clause to discard quotes and unquotes. With this the telnet path comma
541 and virgo_set() are working. Today's kern.log has been added to repository in test_logs/.
542
543 Commits as on 7 November 2013
```

```

544 -----
545 In addition to virgo_malloc and virgo_set, virgo_get is also working through telnet path after today's commit for "vir
546
547 Commits as on 11 November 2013
548 -----
549 More testing done on telnet path for virgo_malloc, virgo_set and virgo_get commands which work correctly. But there se
550 kmem_cache_trace_alloc panics that follow each successful virgo command execution. kern.log for this has been added to
551
552 Commits as on 22 November 2013
553 -----
554 More testing done on telnet path for virgo_malloc, virgo_set and virgo_set after commenting kernel socket shutdown code
555 mempool sendto code. Kernel panics do not occur after commenting kernel socket shutdown.
556
557 Commits as on 2 December 2013
558 -----
559 Lots of testing were done on telnet path and syscall path connection to VIRGO mempool driver and screenshots for worki
560
561 Commits as on 5 December 2013
562 -----
563 More testing on system call path done for virgo_malloc(), virgo_set() and virgo_get() system calls with test_virgo_mal
564
565 VIRGO version 12.0 tagged.
566
567 Commits as on 12 March 2014
568 -----
569 Initial VIRGO queueing driver implemented that flips between two internal queues: 1) a native queue implemented locall
570 structure virgo_workqueue_request.
571
572 Commits as on 20 March 2014
573 -----
574 - VIRGO queue with additional boolean flags for its use as KingCobra queue
575 - KingCobra kernel space driver that is invoked by the VIRGO workqueue handler
576
577 Commits as on 30 March 2014
578 -----
579 - VIRGO mempool driver has been augmented with use_as_kingcobra_service flags in CPU pooling and Memory pooling driver
580
581 Commits as on 6 April 2014
582 -----
583 - VIRGO mempool driver recvfrom() function's if clause for KingCobra has been updated for REQUEST header formatting me
584
585 Commits as on 7 April 2014
586 -----
587 - generate_logical_timestamp() function has been implemented in VIRGO mempool driver that generates timestamps based o
588
589 Commits as on 25 April 2014
590 -----
591 - client ip address in VIRGO mempool recvfrom KingCobra if clause is converted to host byte order from network byte or
592
593 Commits as on 5 May 2014
594 -----
595 - Telnet path commands for VIRGO cloud file system - virgo_cloud_open(), virgo_cloud_read(), virgo_cloud_write(), virg
596
597 Commits as on 7 May 2014
598 -----
599 - Bugfixes to tokenization in kernel upcall plugin with strsep() for args passed on to the userspace
600
601 Commits as on 8 May 2014
602 -----
603 - Bugfixes to virgo_cloud_fs.c for kernel upcall (parameterIsExecutable=0) and with these the kernel to userspace upca
604
605 Commits as on 6 June 2014
606 -----
607 - VIRGO File System Calls Path implementation has been committed. Lots of Linux Full Build compilation errors fixed an
608
609 Commits as on 3 July 2014
610 -----
611 - More testing and bugfixes for VIRGO File System syscalls have been done. virgo_write() causes kernel panic.
612
613 7 July 2014 - virgo_write() kernel panic notes:
614 -----
615 warning within http://lxr.free-electrons.com/source/arch/x86/kernel/smp.c#L121:
616
617 static void native_smp_send_reschedule(int cpu)
618 {
619     if (unlikely(cpu_is_offline(cpu))) {
620         WARN_ON(1);
621         return;
622     }
623     apic->send_IPI_mask(cpumask_of(cpu), RESCHEDULE_VECTOR);
624

```



```

625 }
626
627 This is probably a fixed kernel bug in <3.7.8 but recurring in 3.7.8:
628 - http://lkml.iu.edu/hypermail/linux/kernel/1205.3/00653.html
629 - http://www.kernelhub.org/?p=3&msg=74473&body_id=72338
630 - http://lists.openwall.net/linux-kernel/2012/09/07/22
631 - https://bugzilla.kernel.org/show_bug.cgi?id=54331
632 - https://bbs.archlinux.org/viewtopic.php?id=156276
633
634
635 Commits as on 29 July 2014
636 -----
637 All VIRGO drivers(cloudfs, queuing, cpupooling and memorypooling) have been built on 3.15.5 kernel with some Makefile
638
639 -----
640 Commits as on 17 August 2014
641 -----
642 (FEATURE - DONE) VIRGO Kernel Modules and System Calls major rewrite for 3.15.5 kernel - 17 August 2014
643 -----
644 1. VIRGO config files have been split into /etc/virgo_client.conf and /etc/virgo_cloud.conf to delink the cloud client
645 config parameters reading and to do away with oft occurring symbol lookup errors and multiple definition errors for nu
646 node_ip_addrs_in_cloud - these errors are frequent in 3.15.5 kernel than 3.7.8 kernel.
647
648 2. Each VIRGO module and system call now reads the config file independent of others - there is a read_virgo_config_<m
649
650 3. New kernel module config has been added in drivers/virgo. This is for future prospective use as a config export dri
651 be looked up by any other VIRGO module for config parameters.
652
653 4. include/linux/virgo_config.h has the declarations for all the config variables declared within each of the VIRGO ke
654
655 5. Config variables in each driver and system call have been named with prefix and suffix to differentiate the module
656
657 6. In geographically distributed cloud virgo_client.conf has to be in client nodes and virgo_cloud.conf has to be in c
658
659 7. Above segregation largely simplifies the build process as each module and system call is independently built withou
660
661 8. VIRGO File system driver and system calls have been tested with above changes and the virgo_open(),virgo_read() and
662
663 -----
664 Committed as on 23 August 2014
665 -----
666 Commenting use_as_kingcobra_service if clauses temporarily as disabling also doesnot work and only commenting the bloc
667 works for VIRGO syscall path. Quite weird as to how this relates to the problem. As this is a heisenbug further testin
668 difficult and sufficient testing has been done with logs committed to repository. Probably a runtime symbol lookup for
669 causes the freeze.
670 For forwarding messages to KingCobra and VIRGO queues, cpupooling driver is sufficient which also has the use_as_king
671
672 -----
673 Committed as on 23 August 2014 and 24 August 2014
674 -----
675 As cpupooling driver has the same crash problem with kernel_accept() when KingCobra has benn enabled, KingCobra clause
676
677 VIRGO cpupooling or memorypooling ==> VIRGO Queue ==> KingCobra
678
679 (or)
680 VIRGO Queue kernel service ==> KingCobra
681
682 -----
683 Committed as on 26 August 2014
684 -----
685 - all kmallocs have been made into GFP_ATOMIC instead of GFP_KERNEL
686 - moved some kingcobra related header code before kernel_recvmg()
687 - some header file changes for set_fs()
688
689 This code has been tested with modified code for KingCobra and the standalone
690 kernel service that accepts requests from telnet directly at port 60000, pushes to virgo_queue
691 and is handled to invoke KingCobra servicerequest kernelspace function, works
692 (the kernel_recvmg() crash was most probably due to Read-Only filesystem -errno printed is -30)
693
694 -----
695 VIRGO version 14.9.9 has been release tagged on 9 September 2014
696 -----
697
698 -----
699 Committed as on 26 November 2014
700 -----
701 New kernel module cloudsnc has been added to repository under drivers/virgo that can be used for synchronization(lock
702
703 -----
704 Committed as on 27 November 2014
705 -----

```

```

706 virgo_bakery.h bakery_lock() has been modified to take 2 parameters - thread_id and number of for loops (1 or 2)
707
708 -----
709 Committed as on 2 December 2014
710 -----
711 VIRGO bakery algorithm implementation has been rewritten with some bugfixes. Sometimes there are soft lockup errors du
712
713 -----
714 Committed as on 17 December 2014
715 -----
716 Initial code commits for VIRGO EventNet kernel module service:
717 -----
718 1.EventNet Kernel Service listens on port 20000
719
720 2.It receives eventnet log messages from VIRGO cloud nodes and writes the log messages
721 after parsing into two text files /var/log/eventnet/EventNetEdges.txt and
722 /var/log/eventnet/EventNetVertices.txt by VFS calls
723
724 3.These text files can then be processed by the EventNet implementations in AsFer (python pygraph and
725 C++ boost::graph based)
726
727 4.Two new directories virgo/utils and virgo/eventnet have been added.
728
729 5.virgo/eventnet has the new VIRGO EventNet kernel module service implementation that listens on
730 port 20000.
731
732 6.virgo/utils is the new generic utilities driver that has a virgo_eventnet_log()
733 exported function which connects to EventNet kernel service and sends the vertex and edge eventnet
734 log messages which are parsed by kernel service and written to the two text files above.
735
736 7.EventNet log messages have two formats:
737   - Edge message - "eventnet_edgemsg#<id>#<from_event>#<to_event>"
738   - Vertex message - "eventnet_vertextmsg#<id>-<partakers csv>-<partaker conversations csv>"
739
740 8.The utilities driver Module.symvers have to be copied to any driver which are
741 then merged with the symbol files of the corresponding driver. Target clean has to be commented while
742 building the unified Module.symvers because it erases symvers carried over earlier.
743
744 9.virgo/utils driver can be populated with all necessary utility exported functions that might be needed
745 in other VIRGO drivers.
746
747 10.Calls to virgo_eventnet_log() have to be #ifdef guarded as this is quite network intensive.
748
749 -----
750 Commits as on 18 December 2014
751 -----
752 Miscellaneous bugfixes,logs and screenshot
753
754 - virgo_cloudeexec_eventnet.c - eventnet messages parser errors and eventnet_func bugs fixed
755 - virgo_cloud_eventnet_kernelspace.c - filp_open() args updated due to vfs_write() kernel panics. The vertexmessage vf
756 - VIRGO EventNet build script updated for copying Module.symvers from utils driver for merging with eventnet Module.sy
757 - Other build generated sources and kernel objects
758 - new testlogs directory with screenshot for edgemsg sent to EventNet kernel service and kern.log with previous histor
759 - vertex message update
760
761 -----
762 Commits as on 2,3,4 January 2015
763 -----
764 - fixes for virgo eventnet vertex and edge message text file vfs_write() errors
765 - kern.logs and screenshots
766
767 -----
768 VIRGO version 15.1.8 release tagged on 8 January 2015
769 -----
770
771 -----
772 Commits as on 3 March 2015 - Initial commits for Kernel Analytics Module which reads the /etc/virgo_kernel_analytics.c
773 -----
774 - Architecture of Key-Value Store in memorypooling (virgo_malloc,virgo_get,virgo_set,virgo_free) has been
775 uploaded as a diagram at http://sourceforge.net/p/virgo-linux/code-0/HEAD/tree/trunk/virgo-docs/VIRGOLinuxKernel\_KeyVa
776
777 - new kernel_analytics driver for AsFer <=> VIRGO+USBmd+KingCobra interface has been added.
778 - virgo_kernel_analytics.conf having csv(s) of key-value pairs of analytics variables is set by AsFer or any other Mac
779 - kernel_analytics Driver build script has been added
780
781 -----
782 Commits as on 6 March 2015
783 -----
784 - code has been added in VIRGO config module to import EXPORTed kernel_analytics config key-pair array
785 set by Apache Spark (mined from Uncomplicated Fire Wall logs) and manually and write to kern.log.
786

```

```

787 -----
788 NeuronRain version 15.6.15 release tagged
789 -----
790
791 -----
792 Portability to linux kernel 4.0.5
793 -----
794 The VIRGO kernel module drivers are based on kernel 3.15.5. With kernel 4.0.5 kernel which is the latest following
795 compilation and LD errors occur - this is on cloudfs VIRGO File System driver :
796 - msghdr has to be user_msghdr for iov and iov_len as there is a segregation of msghdr
797 - modules_install throws an error in scripts/Makefile.modinst while overwriting already installed module
798 -----
799
800 Commits as on 9 July 2015
801 -----
802 VIRGO cpupooling driver has been ported to linux kernel 4.0.5 with msghdr changes as mentioned previously
803 with kern.log for VIRGO cpupooling driver invoked in parameterIsExecutable=2 (kernel module invocation)
804 added in testlogs
805 -----
806
807 Commits as on 10,11 July 2015
808 -----
809 VIRGO Kernel Modules:
810 - memorypooling
811 - cloudfs
812 - utils
813 - config
814 - kernel_analytics
815 - cloudsync
816 - eventnet
817 - queuing
818 along with cpupooling have been ported to Linux Kernel 4.0.5 - Makefile and header files have been
819 updated wherever required.
820 -----
821
822 Commits as on 20,21,22 July 2015
823 -----
824 Due to SourceForge Storage Disaster(http://sourceforge.net/blog/sourceforge-infrastructure-and-service-restoration/),
825 the github replica of VIRGO is urgently updated with some important changes for msg_iter,iovec
826 etc., in 4.0.5 kernel port specifically for KingCobra and VIRGO Queueing. These have to be committed to SourceForge Kr
827 repository at http://sourceforge.net/users/ka\_shrinivaasan once SourceForge repos are restored.
828 Time to move on to the manufacturing hub? GitHub ;-))
829 -----
830 VIRGO Queueing Kernel Module Linux Kernel 4.0.5 port:
831 -----
832 - msg_iter is used instead of user_msghdr
833 - kvec changed to iovec
834 - Miscellaneous BUF_SIZE related changes
835 - kern.logs for these have been added to testlogs
836 - Module.symvers has been recreated with KingCobra Module.symvers from 4.0.5 KingCobra build
837 - clean target commented in build script as it wipes out Module.symvers
838 - updated .ko and .mod.c
839 -----
840 KingCobra Module Linux Kernel 4.0.5 port
841 -----
842 - vfs_write() has a problem in 4.0.5
843 - the filp_open() args and flags which were working in 3.15.5 cause a
844 kernel panic implicitly and nothing was written to logs
845 - It took a very long time to figure out the reason to be vfs_write and filp_open
846 - O_CREAT, O_RDWR and O_LARGEFILE cause the panic and only O_APPEND is working, but
847 does not do vfs_write(). All other VIRGO Queue + KingCobra functionalities work viz.,
848 enqueueing, workqueue handler invocation, dequeueing, invoking kingcobra kernelspace service
849 request function from VIRGO queue handler, timestamp, timestamp and IP parser, reply_to_publisher etc.,
850 - As mentioned in Greg Kroah Hartman's "Driving me nuts", persistence in Kernel space is
851 a bad idea but still seems to be a necessary stuff - yet only vfs calls are used which have to be safe
852 - Thus KingCobra has to be in-memory only in 4.0.5 if vfs_write() doesn't work
853 - Intriguingly cloudfs filesystems primitives - virgo_cloud_open, virgo_cloud_read, virgo_cloud_write etc.,
854 work perfectly and append to a file.
855 - kern.logs for these have been added to testlogs
856 - Module.symvers has been recreated for 4.0.5
857 - updated .ko and .mod.c
858 -----
859
860 Due to SourceForge outage and for a future code diversification
861 NeuronRain codebases (AsFer, USBmd, VIRGO, KingCobra)
862 in http://sourceforge.net/u/userid-769929/profile/ have been
863 replicated in GitHub also - https://github.com/shrinivaasanka
864 excluding some huge logs due to Large File Errors in GitHub.
865 -----
866
867 -----

```

Commits as on 30 July 2015

VIRGO system calls have been ported to Linux Kernel 4.0.5 with commented gcc option -Wimplicit-function-declaration, msghdr and iovec changes similar to drivers mentioned in previous commit notes above. But Kernel 4.1.3 has some Makefile The NeuronRain codebases in SourceForge and GitHub would henceforth be mostly and always out-of-sync and not guarantee

Commits as on 2,3 August 2015

- new .config file added which is created from menuconfig
- drivers/Kconfig has been updated with 4.0.5 drivers/Kconfig for trace event linker errors
- Linux Kernel 4.0.5 - KConfig is drivers/ has been updated to resolve RAS driver trace event linker error. RAS was not
- link-vmlinux.sh has been replaced with 4.0.5 kernel version

Commits as on 12 August 2015

VIRGO Linux Kernel 4.1.5 port - related code changes - some important notes:

- Linux Kernel 4.0.5 build suddenly had a serious root shell drop error in initramfs which was not resolved by:
 - adding rootdelay in grub
 - disabling uuid for block devices in grub config
 - mounting in read/write mode in recovery mode
 - no /dev/mapper related errors
 - repeated exits in root shell
 - delay before mount of root device in initrd scripts
- mysteriously there were some firmware microcode bundle executions in ieucodetool
- Above showed a serious grub corruption or /boot MBR bug or 4.0.5 VIRGO kernel build problem
- Linux 4.0.x kernels are EOL-ed
- Hence VIRGO is ported to 4.1.5 kernel released few days ago
- Only minimum files have been changed as in commit log for Makefiles and syscall table and headers and a build script for 4.1.5:
 - Changed paths:
 - A buildscript_4.1.5.sh
 - M linux-kernel-extensions/Makefile
 - M linux-kernel-extensions/arch/x86/syscalls/Makefile
 - M linux-kernel-extensions/arch/x86/syscalls/syscall_32.tbl
 - M linux-kernel-extensions/drivers/Makefile
 - M linux-kernel-extensions/include/linux/syscalls.h

- Above minimum changes were enough to build an overlay-ed Linux Kernel with VIRGO codebase

Commits as on 14,15,16 August 2015

Executed the minimum end-end telnet path primitives in Linux kernel 4.1.5 VIRGO code:

- cpu virtualization
- memory virtualization
- filesystem virtualization (updated filp_open flags)

and committed logs and screenshots for the above.

Commits as on 17 August 2015

VIRGO queue driver:

- Rebuilt Module.symvers
- kern.log for telnet request to VIRGO Queue + KingCobra queueing system in kernelspace

Commits as on 25,26 September 2015

VIRGO Linux Kernel 4.1.5 - memory system calls:

- updated testcases and added logs for syscalls invoked separately(malloc,set,get,free)
- The often observed unpredictable heisen kernel panics occur with 4.1.5 kernel too. The logs are 2.3G and only grepped output is committed to repository.
- virgo_malloc.c has been updated with kstrdup() to copy the buf to iov.iov_base which was earlier crashing in copy_from_iter() within tcp code. This problem did not happen in 3.15.5 kernel.
- But virgo_clone syscall code works without any changes to iov_base as above which does a strcpy() which is an internal memcpy() though. So what causes this crash in memory system calls alone is a mystery.
- new insmod script has been added to load the VIRGO memory modules as necessary instead of at boot time.
- test_virgo_malloc.c and its Makefile has been updated.

VIRGO Linux Kernel 4.1.5 - filesystem calls- testcases and logs:

- added insmod script for VIRGO filesystem drivers
- test_virgo_filesystem.c has been updated for syscall numbers in 4.1.5 VIRGO kernel
- virgo_fs.c syscalls code has been updated for iov.iov_base kstrdup() - without this there are kernel panics in cop

testlogs have been added, but there are heisen kernel panics. The virgo syscalls are executed but not written to kern. Thus execution logs are missing for VIRGO filesystem syscalls.

 Commits as on 28,29 September 2015

VIRGO Linux Kernel 4.1.5 filesystem syscalls:

- Rewrote iov_base code with a separate iovbuf set to iov_base and strcpy()-ing the syscall command to iov_base similar memory syscalls
- Pleasantly the same iovbuf code that crashes in memory syscalls works for VIRGO FS without crash. Thus both virgo_clone syscalls work without issues in 4.1.5 and virgo_malloc() works erratically in 4.1.5 which remains as issue.
- kern.log for VIRGO FS syscalls and virgofstest text file written by virgo_write() have been added to repository

VIRGO Linux 4.1.5 kernel memory syscalls:

- rewrote the iov_base buffer code for all VIRGO memory syscalls by allocating separate iovbuf and copying the message
- did extensive repetitive tests that were frequented by numerous kernel panics and crashes
- The stability of syscalls code with 3.15.5 kernel appears to be completely absent in 4.1.5
- The telnet path works relatively better though
- Difference between virgo_clone and virgo_malloc syscalls despite having same kernel sockets code looks like a non-trivial
- kernel OOPS traces are quite erratic.
- Makefile path in testcase has been updated

 Commits as on 4 October 2015

VIRGO Linux Kernel 4.1.5 - Memory System Calls:

- replaced copy_to_user() with a memcpy()
- updated the testcase with an example VUID hardcoded.
- str_to_addr2() is done on iov_base instead of buf which was causing NULL parsing
- kern.log with above resolutions and multiple VIRGO memory syscalls tests - malloc,get,set
- With above VIRGO malloc and set syscalls work relatively causing less number of random kernel panics
- return values of memory calls set to 0
- in virgo_get() syscall, memcpy() of iov_base is done to data_out userspace pointer
- kern.log with working logs for syscalls - virgo_malloc(), virgo_set(), virgo_get() but still there are random kernel
- Abridged kern.log for VIRGO Memory System Calls with 4.1.5 Kernel - shows example logs for virgo_malloc(), virgo_set

 Commits as on 14 October 2015

VIRGO Queue Workqueue handler usermode clause has been updated with 4.1.5 kernel paths and kingcobra in user mode is enabled

 Commits as on 15 October 2015

- Updated VIRGO Queue kernel binaries and build generated sources
- virgo_queue.h has been modified for call_usermodehelper() - set_ds() and fd_install() have been uncommented for output

 Commits as on 3 November 2015

- kern.log for VIRGO kernel_analytics+config drivers which export the analytics variables from /etc/virgo_kernel_analytics

 Commits as on 10 January 2016

NeuronRain VIRGO enterprise version 2016.1.10 released.

 NeuronRain - AsFer commits for VIRGO - C++ and C Python extensions

- Commits as on 29 January 2016

(FEATURE - DONE) Python-C++-VIRGOKernel and Python-C-VIRGOKernel boost::python and cpython implementations:

- It is a known idiom that Linux Kernel and C++ are not compatible.
- In this commit an important feature to invoke VIRGO Linux Kernel from userspace python libraries via two alternative
- In one alternative, C++ boost::python extensions have been added to encapsulate access to VIRGO memory system calls
- In the other alternative, C Python extensions have been added that replicate boost::python extensions above in C - C works exceedingly well compared to boost::python.
- This functionality is required when there is a need to set kernel analytics configuration variables learnt by AsFer dynamically without re-reading /etc/virgo_kernel_analytics.conf.
- This completes a major integration step of NeuronRain suite - request travel roundtrip to-and-fro top level machine-code and rock-bottom C linux kernel - bull tamed ;-).
- This kind of python access to device drivers is available for Graphics Drivers already on linux (GPIO - for access in
- logs for both C++ and C paths have been added in cpp_boost_python_extensions/ and cpython_extensions.
- top level python scripts to access VIRGO kernel system calls have been added in both directories:
- CPython - python cpython_extensions/asferpythonextensions.py
- C++ Boost::Python - python cpp_boost_python_extensions/asferpythonextensions.py

```

1030 - .so, .o files with build commandlines(asferpythonextensions.build.out) for "python setup.py build" have been added
1031 in build lib and temp directories.
1032 - main implementations for C++ and C are in cpp_boost_python_extensions/asferpythonextensions.cpp and cpython_extensio
1033
1034 -----
1035 Commits as on 12 February 2016
1036 -----
1037 Commits for Telnet/System Call Interface to VIRGO CPUPooling -> VIRGO Queue -> KingCobra
1038 -----
1039 *) This was commented earlier for the past few years due to a serious kernel panic in previous kernel versions - <= 3.
1040 *) In 4.1.5 a deadlock between VIRGO CPUPooling and VIRGO queue driver init was causing following error in "use_as_kin
1041 - "gave up waiting for virgo_queue init, unknown symbol push_request()"
1042 *) To address this a new boolean flag to selectively enable and disable VIRGO Queue kernel service mode "virgo_queue_r
1043 *) With this flag VIRGO Queue is both a kernel service driver and a standalone exporter of function symbols - push_req
1044 *) Incoming request data from telnet/virgo_clone() system call into cpupooling kernel service reactor pattern (virgo c
1045 *) This resolves a long standing deadlock above between VIRGO cpupooling "use_as_kingcobra_service" clause and VIRGO q
1046 *) This makes virgo_clone() syscall/telnet both synchronous and asynchronous - requests from telnet client/virgo_cl
1047 *) Above saves an additional code implementation for virgo_queue syscall paths - virgo_clone() handles, based on confi
1048 -----
1049 Prerequisites:
1050 -----
1051 - insmod kingcobra_main_kernelspace.ko
1052 - insmod virgo_queue.ko compiled with flag virgo_queue_reactor_service_mode=1
1053 (when virgo_queue_reactor_service_mode=0, listens on port 60000 for direct telnet requests)
1054 - insmod virgo_cloud_test_kernelspace.ko
1055 - insmod virgo_cloudexec.ko (listens on port 10000)
1056 -----
1057 Schematic Diagram
1058 -----
1059 VIRGO clone system call/telnet client ---> VIRGO cpupooling(compiled with use_as_kingcobra_service=1) -----> VIRGO Qu
1060 -----
1061 -----
1062 Commits as on 15 February 2016 - Kernel Analytics - VIRGO Linux Kernelwide imports
1063 -----
1064 - Imported Kernel Analytics variables into CloudFS kernel module - printed in driver init()
1065 - Module.symvers from kernel_analytics has been merged with CloudFS Module.symvers
1066 - Logs for above has been added in cloudfs/test_logs/
1067 - Makefile updated with correct fs path
1068 - Copley notices updated
1069 -----
1070 -----
1071 Commits as on 15 February 2016 - Kernel Analytics - VIRGO Linux Kernelwide imports
1072 -----
1073 - Kernel Analytics driver exported variables have been imported in CPU virtualization driver
1074 - Module.symvers from kernel_analytics has been merged with Module.symvers in cpupooling
1075 - kern.log for this import added to cpupooling/virgocloudexec/test_logs/
1076 -----
1077 -----
1078 Commits as on 15 February 2016 - Kernel Analytics - VIRGO Linux Kernelwide imports
1079 -----
1080 - Imported kernel analytics variables into memory virtualization driver init() , exported from kernel_analytics driver
1081 - build shell script updated
1082 - logs added to test_logs/
1083 - Module.symvers from kernel_analytics has been merged with memory driver Module.symvers
1084 - Makefile updated
1085 -----
1086 -----
1087 Commits as on 15 February 2016 - Kernel Analytics - VIRGO Linux Kernelwide imports
1088 -----
1089 - Imported kernel analytics variables into VIRGO Queueing Driver
1090 - logs for this added in test_logs/
1091 - Makefile updated
1092 - Module.symvers from kernel_analytics has been merged with Queueing driver's Module.symvers
1093 - .ko, .o and build generated sources
1094 -----
1095 -----
1096 Commits as on 16,17 February 2016
1097 -----
1098 (FEATURE-DONE) Socket Buffer Debug Utility Function - uses linux skbuff facility
1099 -----
1100 - In this commit a multipurpose socket buffer debug utility function has been added in utils driver and exported kerne
1101 - It takes a socket as function argument does the following:
1102 - dereference the socket buffer head of skbuff per-socket transmit data queue
1103 - allocate skbuff with alloc_skb()
1104 - reserve head room with skb_reserve()
1105 - get a pointer to data payload with skb_put()
1106 - memcpy() an example const char* to skbuff data
1107 - Iterate through the linked list of skbuff queue in socket and print headroom and data pointers
1108 - This can be used as a packet sniffer anywhere within VIRGO linux network stack
1109 - Any skb_*( ) functions can be plugged-in here as deemed necessary.
1110

```



```

1111 - kern.log(s) which print the socket internal skbuff data have been added to a new testlogs/ directory
1112 - .cmd files generated by kbuild
1113
1114 -----
1115 (FEATURE-DONE) Commits as on 24 February 2016
1116 -----
1117 skbuff debug function in utils/ driver:
1118 (*) Added an if clause to check NULLity of skbuff headroom before doing skb_alloc()
1119 (*) kern.log for this commit has been added testlogs/
1120 (*) Rebuilt kernel objects and sources
1121
1122 -----
1123 Commits as on 29 February 2016
1124 -----
1125 -----
1126 (FEATURE-DONE) Software Analytics - SATURN Program Analysis added to VIRGO Linux kernel drivers
1127 -----
1128 - SATURN (saturn.stanford.edu) Program Analysis and Verification software has been
1129 integrated into VIRGO Kernel as a Verification+SoftwareAnalytics subsystem
1130 - A sample driver that can invoke an exported function has been added in drivers - saturn_program_analysis
1131 - Detailed document for an example null pointer analysis usecase has been created in virgo-docs/VIRGO_SATURN_Program_A
1132 - linux-kernel-extensions/drivers/virgo/saturn_program_analysis/saturn_program_analysis_trees/error.txt is the error r
1133 - SATURN generated preproc and trees are in linux-kernel-extensions/drivers/virgo/saturn_program_analysis/preproc and
1134 linux-kernel-extensions/drivers/virgo/saturn_program_analysis/saturn_program_analysis_trees/
1135
1136 -----
1137 Commits as on 10 March 2016
1138 -----
1139 - SATURN analysis databases (.db) for locking, memory and CFG analysis.
1140 - DOT and PNG files for locking, memory and CFG analysis.
1141 - new folder saturn_calypso_files/ has been added in saturn_program_analysis/ with new .clp files virgosaturncfg.clp a
1142 - SATURN alias analysis .db files
1143
1144 -----
1145 (FEATURE-DONE) NEURONRAIN - ASFER Commits for VIRGO - CloudFS systems calls integrated into Boost::Python C++ and Pyth
1146 -----
1147 -----
1148 AsFer Commits as on 30 May 2016
1149 -----
1150 VIRGO CloudFS system calls have been added (invoked by unique number from syscall_32.tbl) for C++ Boost::Python interf
1151 VIRGO Linux System Calls. Switch clause with a boolean flag has been introduced to select either VIRGO memory or files
1152 kern.log and CloudFS textfile Logs for VIRGO memory and filesystem invocations from AsFer python have been committed t
1153
1154 -----
1155 AsFer Commits as on 31 May 2016
1156 -----
1157 Python CAPI interface to NEURONRAIN VIRGO Linux System Calls has been updated to include File System open, read, write
1158 Rebuilt extension binaries, kern.logs and example appended text file have been committed to testlogs/. This is exactly
1159 commits done for Boost::Python C++ interface. Switch clause has been added to select memory or filesystem VIRGO syscal
1160
1161 -----
1162 (BUG - STABILITY ISSUES) Commits - 25 July 2016 - Static Analysis of VIRGO Linux kernel for investigating heisencrashe
1163 -----
1164 Initial Documentation for Smatch and Coccinelle kernel static analyzers executed on VIRGO Linux kernel - to be updated
1165 periodically with further analysis.
1166
1167 -----
1168 (BUG - STABILITY ISSUES) Commits - 1 August 2016 - VIRGO Linux Stability Issues - Ongoing Random oops and panics inves
1169 -----
1170 1. GFP_KERNEL has been replaced with GFP_ATOMIC flags in kmem allocations.
1171 2. NULL checks have been introduced in lot of places involving strcpy, strcat, strcmp etc., to circumvent
1172 buffer overflows.
1173 3. Though this has stabilized the driver to some extent, still there are OOPS in unrelated places deep
1174 with in kernel where paging datastructures are accessed - kmalloc somehow corrupts paging
1175 4. OOPS are debugged via gdb as:
1176     4.1 gdb ./vmlinux /proc/kcore
1177     or
1178     4.2 gdb <loadable_kernel_module>.o
1179     followed by
1180     4.3 l *(address+offset in OOPS dump)
1181 5. kern.log(s) for the above have been committed in tar.gz format and have numerous OOPS occurred during repetitive te
1182 invocation(boost::python C++) invocations of virgo memory system calls.
1183 6. Paging related OOPS look like an offshoot of set_fs() encompassing the filp_open VFS calls.
1184
1185 -----
1186 (BUG-STABILITY ISSUES) Commits - 26 September 2016 - Ongoing Random Panic investigation
1187 -----
1188 Further analysis on direct VIRGO memory cache primitives telnet invocation - problems are similar
1189 to Boost::Python AsFer VIRGO system calls invocations.
1190
1191 -----

```

```

1192 (BUG-STABILITY ISSUES) Commits - 27 September 2016 - Ongoing Random Panic investigation
1193 -----
1194 Analysis of VIRGO memory cache primitives reveal more inconsistencies in cacheline flushes between CPU and GPU.
1195 -----
1196 Commits - 20 March 2017 and 21 March 2017 - VIRGO Linux 64-bit build based on 4.10.3 kernel
1197 -----
1198 *) moved virgoeventnetclient_driver_build.sh to virgoutils_driver_build.sh in utils/ driver
1199 *) Updated VIRGO Linux Build Steps for 4.10.3
1200 *) New repository has been created for 64-bit VIRGO Linux kernel based on 4.10.3 mainline kernel in GitHub and importe
1201 https://github.com/shrinivaasanka/virgo64-linux-github-code
1202 https://sourceforge.net/p/virgo64-linux/
1203 *) Though it could have been branched off from existing VIRGO repository (32-bit) which is based on 4.1.5 mainline ker
1204 separate repository for 64-bit 4.10.3 VIRGO kernel code was simpler because:
1205 - there have been directory path changes for syscall entries in 4.10.3 and some other KBuild entities
1206 - Some script changes done for 4.1.5 in modpost and vmlinux phases are not required
1207 - having two VIRGO branches one with 4.1.5 code and 32-bit driver .ko binaries and other with 4.10.3 code and
1208 binaries could be unmanageable and commits could go into wrong branch
1209 - 4.10.3 64-bit VIRGO kernel build is still in experimental phase and it is not known if 64-bit 4.10.3 build s
1210 problems in 4.1.5
1211 - If necessary one of these two repositories could be made branch of the other later
1212 -----
1213 Commits - 27 March 2017 Ongoing analysis of VIRGO 64 bit linux kernel based on 4.10.3 kernel mainline
1214 -----
1215 *) Prima facie, 64 bit kernel is quite finicky and importunate compared to 32 bit and 64 bit specific idiosyncrasies a
1216 *) During the past 1 week, quite a few variants of kernel and drivers builds were tried with KASAN enabled and without
1217 *) KASAN shows quite huge number of user memory accesses which later translate to panics.
1218 *) Most nagging of these was kernel_recvmmsg() panic.
1219 *) Added and updated skbuff socket debug utility driver with a new debug function and to print more fields of skbuff
1220 *) KASAN was complaining about _asan_load8 (loading 8 userspace bytes)
1221 *) All erroneous return data types in VIRGO mempool ops structure have been corrected in VIRGO headers
1222 *) all type casts have been sanitized
1223 *) Changed all kernel stack allocations to kernel heap kzallocs
1224 *) This later caused a crash in inet_sendmsg in kernel_sendmsg()
1225 *) gdb64 disassemble showed a trapping instruction:
1226 testb $0x6,0x91(%14) with corresponding source line:
1227 sg = !(sk->sk_route_caps & NETIF_F_SG)
1228 in tcp_sendmsg() (net/ipv4/tcp.c)
1229 *) changed kernel_sendmsg() to sock->ops->sendmsg()
1230 *) These commits are still ongoing analysis only.
1231 *) Screenshots for these have been added to debug-info/
1232 -----
1233 Continued analysis of VIRGO 64-bit linux kernel built on 4.10.3 mainline - Commits - 30 March 2017
1234 -----
1235 *) Previous commit was crashing inside tcp_sendmsg()
1236 *) GDB64 disassembly shows NULL values for register R12 which is added with an offset 91 and is an operand in testb
1237 *) Protected all kernel_sendmsg() and kernel_recvmmsg() in both system calls side and drivers side with
1238 oldfs=get_fs(), set_ds(KERNEL_DS) and set_fs(oldfs)
1239 blocks without which there are random kernel_sendmsg and kernel_recvmmsg hangs
1240 *) Removed init_net and sock_create_kern usage everywhere and replaced them with sock_create calls
1241 *) Tried MSG_FASTOPEN flags but it does not help much in resolving tcp_sendmsg() NULL pointer dereference issue. MSG_F
1242 speeds up the message delivery by piggybacking the message payload before complete handshake is established(SYN, SYN-AC
1243 SYN-ACK itself. But eventually it has to be enabled as fast open is becoming a standard.
1244 *) Kasan reports have been enabled.
1245 *) Added more debug code in skbuff debug utility functions in utils driver to check if sk->prot is a problem.
1246 *) Replaced kernel_sendmsg with a sock->ops->sendmsg() in mempool sendto function which otherwise crashes in tcp_sendm
1247 *) With sock->ops->sendmsg() systemcalls <-----> drivers two-way request-reply works but still there are random -32 (b
1248 (Connection Reset by Peer) errors
1249 *) Logs for working sys_virgo_malloc() call with correctly returned VIRGO Unique ID for memory allocated has been comm
1250 virgocloudexecmempool
1251 *) sock->ops->sendmsg() in mempool driver sendto function requires a MSG_NOSIGNAL flag which prevents SIGPIPE signal t
1252 *) Reason for random broken pipe and connection reset by peer errors in mempool sendto is unknown. Both sides have con
1253 there is no noticeable traffic.
1254 *) While socket communications in 32 bit VIRGO kernel syscalls and drivers work with no issues, why 64-bit has so many
1255 Reasons could be 64 bit address alignment issues, 64 bit specific #ifdefs in kernel code flow, major changes from 4.1.
1256 *) NULL values for register R12 indicate already freed skbuff data which are accessed/double-freed. Kernel TCP engine
1257 *) TCP engine clones the head data of skbuff queue, transmits it and waits for an ACK or timeout. Data is freed only i
1258 And head of the queue is advanced to next element in write queue and this continues till write queue is empty waiting
1259 *) If ACK is not received, head data is cloned again and retransmitted by sequence number flow control.
1260 -----
1261 Continued Analysis of VIRGO 64 bit based on 4.10.3 linux kernel - Commits - 1 April 2017, 3 April 2017
1262 -----
1263 *) kernel_sendmsg() has been replaced with sock->ops->sendmsg() because
1264 kernel_sendmsg() is quite erratic in 4.10.3 64 bit
1265 *) There were connection reset errors in system calls side for virgo_malloc/. This was probably because
1266 sock->ops->sendmsg() requires MSG_DONTWAIT and MSG_NOSIGNAL flags and sendmsg does not block.
1267 *) sock release happens and virgo_malloc syscalls receives -104 error
1268 *) Temporarily sock_release has been commented. Rather socket timeout should be relied upon which should

```

```

1273 do automatic release of socket resources
1274 *) Similar flags have been applied in virgo_malloc syscalls too.
1275 *) Logs with above changes do not have reset errors as earlier.
1276 *) virgo set/get still crashes because 64 bit id is truncated which would require data type changes for
1277 64 bit
1278 *) test_virgo_malloc test case has been rebuilt with -m64 flag for invocation of 64 bit syscalls by
1279 numeric ids
1280
1281 -----
1282 Continued Analysis of VIRGO 64 bit 4.10.3 kernel - commits - 10 April 2017
1283 -----
1284 *) There is something seriously wrong with 4.10.3 kernel sockets in 64 bit build VIRGO send/rcv messages and even acc
1285 *) All kernel socket functionalities which work well in 4.1.5 32 bit VIRGO , have random hangs, panics in 4.10.3 VIRGO
1286 in inet_rcvmsg/sendmsg code path
1287 *) KASAN shows attempts to access user address which occurs despite set_fs(KERNEL_DS)
1288 *) Crash stack is similar to previous crashes in tcp_sendmsg()
1289 *) Tried different address and protocol families for kernel socket accept (TCP,UDP,RAW sockets)
1290 *) With Datagram sockets, kernel_listen() mysteriously fails with -95 error in kernel_bind(operation not supported)
1291 *) With RAW sockets, kernel_listen() fails with -93 error for AF_PACKET (protocol not supported)
1292 *) tcpdump pcap sniffer doesn't show anything unruly.
1293 *) This could either be a problem with kernel build (unlikely), Kbuild .config or could have extraneous reasons. But .
1294 4.1.5 and 4.10.3 are similar.
1295 *) Only major difference between 4.10.3 and 4.1.5 is init_net added in sock_create_kern() internally
1296 *) datatype of VIRGO Unique ID has been changed to unsigned long long (_u64)
1297 *) tried with INADDR_LOOPBACK in place of INADDR_ANY
1298 *) also tried with disabled multi(homing) in /etc/hosts.conf
1299 *) Above random kernel socket hangs occur across all VIRGO system calls and drivers transport.
1300 *) Utils kernel socket client to EventNet kernel service also has similar inet_rcvmsg/inet_sendmsg panic problems.
1301
1302 -----
1303 Commits - 11 April 2017 - EventNet and Utils Drivers 64bit
1304 -----
1305 *) EventNet driver works in 64 bit VIRGO Linux
1306 *) An example eventnet logging with utils virgo_eventnet_log() works now without tcp_sendmsg() related stalls in previ
1307 *) Return Datatypes for all EventNet operations have been sanitized (struct socket* was returned as int in 32 build an
1308 struct socket*. This reinterpret cast does not work in 64 bit) in eventnet header.
1309 *) utils eventnet log in init() has been updated with a meaningful edge update message
1310 *) kern.log for this has been added to eventnet/testlogs
1311
1312 -----
1313 Commits - 17 April 2017 - VIRGO64 Memory, CPU, FileSystem, EventNet kernel module drivers
1314 -----
1315 *) telnet requests to VIRGO memory(kernelmemcache), cpu and filesystem modules work after resolving issues with return
1316 *) commented le32_to_cpu() and print_buffer() which was suppressing lot of log messages.
1317 *) VIRGO <driver> ops structures have been updated with correct datatypes.
1318 *) reinterpret cast of struct socket* to int has been completely done away with which could have caused 64bit specific
1319 *) lot of kern.log(s) and screen captures have been added for telnet requests in testlogs/ of respective <driver> dire
1320 *) Prima facie 64bit telnet requests to VIRGO module listeners are relatively stabler than 32bit
1321 *) Previous code changes should be relevant to 32 bit VIRGO kernel too.
1322 *) tcp_sendmsg()/tcp_rcvmsg() related hangs could be mostly related to corrupted skbuff queue within each socket.
1323 *) This is because replacing kernel_send/rcvmsg() with sock_send/rcvmsg() causes return value to be 0 while
1324 socket release crashes within skbuff related kernel functions.
1325 *) To make socket state immutable, in VIRGO memory driver header files, client socket has been declared as const type.
1326
1327 -----
1328 Commits - KingCobra 64 bit and VIRGO Queue + KingCobra telnet requests - 17 April 2017
1329 -----
1330 *) Rebuilt KingCobra 64bit kernel module
1331 *) telnet requests to VIRGO64 Queueing module listener driver are serviced by KingCobra servicerequest
1332 *) Request_Reply queue persisted for this VIRGO Queue + KingCobra routing has been committed to c-src/testlogs.
1333 *) kern.log for this routing has been committed in VIRGO64 queueing directory
1334 *) Similar to other drivers struct socket* reinterpret cast to int has been removed and has been made const in queuesv
1335
1336 -----
1337 Commits - VIRGO64 system calls - kernel module listeners - testcases and system calls updates - 18 April 2017
1338 -----
1339 *) All testcases have been rebuilt
1340 *) VIRGO kernel memcache,cpu and filesystem system calls have been updated with set_fs()/get_fs() blocks for kernel_se
1341 and kernel_rcvmsg()
1342 *) Of these virgo_clone() system call testcase (test_virgo_clone) works flawlessly and there are no tcp_sendmsg()/tcp_
1343 kernel panics.
1344 *) VIRGO memcache and filesystem system call testcases have usual tcp_sendmsg()/tcp_rcvmsg() despite the kernel socke
1345 being similar to VIRGO clone system call
1346 *) Logs for VIRGO clone system call to CPU kernel driver module have been committed to virgo_clone/test/testlogs
1347
1348 -----
1349 Commits - VIRGO64 Kernel MemCache and FileSystem system calls to VIRGO Memory and FileSystem Drivers - 19 April 2017
1350 -----
1351 *) Changed iovec in virgo_clone.c to kvec
1352 *) test_virgo_filesystem.c and test_virgo_malloc.c VIRGO system calls testcases have been changed with some additional
1353 *) virgo_malloc.c has been updated with BUF_SIZE in iov_len and memset to zero initialize the buffer. tcp_sendmsg()/tc

```

```

1354 getting stuck in copy_from_iter_full() memcpy with a NULL Dereference. memcpy() was reading past the buffer bound caus
1355 didnot work for iov_len.
1356 *) virgo_fs.c virgo_write() memcpy has been changed back to copy_from_user() thereby restoring status quo ante (commen
1357 because of a kernel panic in older versions of 32 bit VIRGO kernel)
1358 *) Logs for VIRGO kmemcache and filesystem system calls have been committed to respective system call directories.
1359 *) With this all VIRGO64 functionalities work in both telnet and system calls requests routes end-to-end from clients
1360 kernel sockets issues resolved fully.
1361 *) Major findings are:
1362 - VIRGO 4.10.3 64 bit kernel is very much stable compared to 32 bit 4.1.5 kernel
1363 - there are no i915 related errors which happened in VIRGO 32 bit 4.1.5 kernel
1364 - Repetitive telnet and system calls requests to VIRGO modules are stable and there are no kernel panics like 4.1.5 32
1365 - Google Kernel Address Sanitizer is quite helpful in finding stack overruns, null derefs, user memory accesses etc.,
1366 - 64 bit kernel is visibly faster than 32 bit.
1367 - Virgo Unique ID is now extended to 2^64 with unsigned long long.
1368 -----
1369
1370 Commits - VIRGO64 memory and filesystem calls to memory and filesystem drivers requests routing - 20 April 2017
1371 -----
1372 *) Changed return value of virgo_cloud_free_kernelspace() to a string literal "kernel memory freed"
1373 *) Logs for VIRGO64 memory and filesystem calls to memory and filesystem drivers requests routing have been committed
1374 both driver directories
1375 -----
1376
1377 Commits - 27 April 2017
1378 -----
1379 Residual logs for VIRGO 64 bit 4.10.3 kernel committed.
1380 -----
1381
1382 Commits - 25 May 2017
1383 -----
1384 *) Changed LOOPBACK to INADDR_ANY for VIRGO64 kernel memcache listen port
1385 *) All VIRGO64 RPC, kernel memcache, cloud filesystem primitives have been retested
1386 *) VIRGO64 mempool binaries have been rebuilt
1387 -----
1388
1389 Commits - 31 August 2017 - NeuronRain ReadTheDocs Documentation - VIRGO64 System calls and Drivers
1390 (http://neuronrain-documentation.readthedocs.io/en/latest/)
1391 -----
1392 (*) New directory systemcalls_drivers/ has been added to virgo-docs/ and representative VIRGO64
1393 system calls and drivers functionality logs have been committed for demonstration purpose.
1394 (*) VIRGO64 cloudfs driver has been rebuilt after changing virgofstest.txt file creation filp_open() call
1395 (*) Screenshots and logs for VIRGO64 Clone, Kernel MemCache and Cloud FS SystemCalls-Drivers interaction, socket trans
1396 -----
1397
1398 Commits - 23 September 2017 - Major VIRGO mainline kernel version Upgrade for Kernel Transport Layer Security - 4.10.3
1399 -----
1400 (*) Recently released mainline kernel version 4.13 integrates SSL/TLS into kernelspace- KTLS - for the first time.
1401 (*) KTLS is a standalone kernel module af_ktls (https://github.com/ktls) implemented by RedHat and Facebook for optimi
1402 within kernelspace itself and reduce userspace-kernelspace switches.
1403 (*) sendfile() system call in linux which is used for file transmission (combining read+write) from one fd to another
1404 KTLS optimization in kernelspace in af_ktls codebase (af_ktls tool)
1405 (*) VIRGO Linux kernel fork-off requires this kernelspace TLS functionality to fully secure traffic from system call c
1406 cloud node's kernel module listeners
1407 (*) Hence VIRGO64 linux kernel mainline base is urgently upgraded from 4.10.3 to 4.13.3
1408 (*) All system calls and kernel module code in VIRGO64 now have #include(s) for tls.h and invoke kernel_setsockopt() o
1409 kernelspace sockets for SOL_TLS and TLS_TX options and have been rebuilt.
1410 (*) VIRGO64 RPC clone/kmemcache/cloudfs system calls to kernel module listeners have been tested with this new KTLS so
1411 on rebuilt VIRGO64 kernel overlay-ed on 4.13.3 64-bit linux kernel
1412 (*) 4.13 mainline kernel also has SMB CIFS bug fixes for recent malware attacks (WannaCry etc.,) which further ensures
1413 VIRGO64 linux fork-off kernelspace traffic.
1414 (*) New buildscript for 4.13.3 linux kernel has been committed
1415 (*) testlogs for VIRGO64 system calls and driver listeners KTLS transport have been committed in virgo-docs/systemcall
1416 (*) After this upgrade, complete system calls to driver listener traffic is SSL enabled implicitly.
1417 (*) Updated kernel object files for 4.13.3 build are part of this commit.
1418 -----
1419
1420 Commits - Remnant commits for 4.13.3 upgrade - 24 September 2017
1421 -----
1422 Updated init.h and syscalls.h headers for virgo system calls
1423 -----
1424
1425 Commits - VIRGO64 4.13.3 KTLS Upgrade - System Calls-Driver Listeners End-to-End encrypted traffic testing - 25 Septem
1426 -----
1427 (*) VIRGO64 CPU/KMemCache/CloudFS system calls have been invoked by userspace testcases and all primitives work after
1428 (*) Some small modifications to system calls code have been made and rebuilt to remove redundant iovbuf variables in p
1429 (*) test_virgo_filesystem.c testcase has been updated and rebuilt
1430 (*) kern.Log(s) for CPU/KMemCache/CloudFS systemcalls to driver listeners invocations have been committed to respectiv
1431 directories
1432 (*) virgofstest.txt written to by virgo_write() has also been committed. But a weird behaviour is still observed simil
1433 (*) No DRM GEM i915 panics are observed and stability of VIRGO64 + 4.13.3 linux kernel is more or equal to VIRGO64 + 4
1434

```

 Commits - VIRGO64 VIRGO_KTLS branch creation and rebase of master to previous commit - 30 September 2017

(*) New branch VIRGO_KTLS has been created after previous commit on 25 September 2017 and all 5 commits after 25 September 2017 have been branched to VIRGO_KTLS (which has the *#ifdef for crypto_info, reads from /etc/virgo_ktls.c driver module*)

(*) Following are the commit hashes and commandlines in GitHub and SourceForge:

```
git branch -b VIRGO_KTLS
git branch master
git rebase -i <SHA1_on_25September2017>
git rebase --continue
git commit --amend
git push --force
```

```
1958 git checkout -b
1959 git checkout -b VIRGO_KTLS
1960 ls
1961 git checkout VIRGO_KTLS
1962 git push origin VIRGO_KTLS
1963 git status
1964 git checkout
1965 git checkout -b
1966 git branch
1967 git branch master
1968 git branch -h
1969 git branch
1970 git checkout master
1975 git checkout -b
1976 git checkout -b VIRGO_KTLS
1979 git push origin VIRGO_KTLS
1990 git rebase -i bb661e908cba2a5357414e89166f29086a28bdf0
1991 git status
1992 git rebase -i bb661e908cba2a5357414e89166f29086a28bdf0
1996 git rebase --continue
1997 git commit --amend
2019 git rebase -i bb661e908cba2a5357414e89166f29086a28bdf0
2029 git rebase --continue
2037 git push --force
2091 git rebase -i bb661e908cba2a5357414e89166f29086a28bdf0
2092 git rebase --continue
2093 git commit --amend
2094 git push --force
2110 git branch
2111 git branch master
2112 git checkout master
2113 git branch
2114 git rebase -i e76b4089633223f610fddc0e0eaff8c2cef8b9f1
2115 git commit --amend
2116 git rebase --continue
2117 git push --force
```

 KTLS in 4.13.3 has support for only private symmetric encryption. It does not support Public Key Encryption yet. Since mainstream VIRGO64 code might change a lot for other features. Therefore, VIRGO_KTLS specific crypto_info code has been

Commits - 1 October 2017

kern.log(s) for VIRGO64 systemcalls-driver 4.13.3 64-bit upgrade tests on master branch after reversal and rebase of master branching to VIRGO_KTLS. There is a weird General Protection Fault in intel atomic commit work not seen thus far. Also

Commits - VIRGO64 Utils and EventNet Drivers Update for tcp_sendmsg() stack out-of-bounds error - 3 October 2017

(*) Utils Generic Socket Client function virgo_eventnet_log() for EventNet kernel module listener was repeatedly failing emitting -32 and -107 errors.

(*) kernel_connect() was guarded by set_fs() and get_fs() memory segment routines to prevent any memory corruption. After

(*) After replacing strlen(buf) by BUF_SIZE in msg flags before kernel_connect() stack out-of-bounds error has been resolved

(*) kern.log for this has been committed in drivers/virgo/Utils/testlogs/

(*) Both eventnet and utils drivers have been rebuilt

VIRGO64 system calls-drivers on linux kernel 4.13.3 - miscellaneous bugfixes - 5 October 2017

(*) kernel_setsockopt() for KTLS has been commented in all system calls and drivers because KTLS functionality has been
 VIRGO_KTLS

(*) In virgo_clone.c, iov.iov_len has been set to BUF_SIZE

(*) kernel_connect() has been guarded by set_fs()/get_fs() in VIRGO64 system call clients

(*) test_virgo_malloc.c testcase has been updated

(*) There was a weird problem in in4_pton(): sin_addr.saddr was not set correctly from string IP address and this was

(*) in4_pton() is implemented in net/core/Utils.c and reads the string IP address digits and sums up the ASCII values

(*) Repeated builds were done trying different possible fixes but didn't work e.g casting saddr to (u8*)

```

1516  (*) There is an alternative in_aton() function which takes only String IP address and returns address as __be32
1517  (*) After in_aton() in virgo_set() random faulty address conversion does not occur - in_aton() is differently implemen
1518  (*) msg_hdr has been initialized to NULL in virgo_set()
1519  (*) Lot of debug printk()s have been added
1520  (*) kern.log (.tar.gz) for RPC clone/KMemCache/Filesystem systemcalls-driver has been committed to virgo-docs/systemca
1521  (*) VIRGO Linux build steps have been updated for example commandlines to overlay mainline kernel tree by VIRGO64 sour
1522

```

```

1523  commit 4e6681ade4ddb1bed17f7c115b59a5ebf884256

```

```

1524  Author: K.Srinivasan <ka.shrinivaasan@gmail.com>

```

```

1525  Date:   Fri Oct 6 11:36:15 2017 +0530

```

```

1526 -----
1527 VIRGO64 Queueing Kernel Module Listener - KingCobra64 - 4.13.3 - 6 October 2017
1528 -----

```

```

1529
1530  (*) telnet client connection to VIRGO64 Queue and a subsequent workqueue routing (pub/sub) to KingCobra64 has been tes
1531  (*) TX_TLS socket option has not been disabled and is a no-op because it has no effect on the socket.
1532  (*) REQUEST_REPLY.queue for this routing from VIRGO64 queue and persisted by KingCobra64 has been committed to KingCob
1533

```

```

1534  commit d4e95b58474838d65da9c69944c6287acbdfe72c

```

```

1535  Author: K.Srinivasan <ka.shrinivaasan@gmail.com>

```

```

1536  Date:   Fri Oct 6 11:05:21 2017 +0530

```

```

1537 -----
1538 VIRGO64 System Calls to Drivers and Telnet Client to Drivers on 4.13.3 linux kernel - master branch (after KTLS has be
1539 and branched to VIRGO_KTLS) - test case logs - 6 October 2017
1540 -----

```

```

1541
1542  (*) VIRGO64 System calls - Clone, KMemCache and Filesystem system call primitives to Driver listeners invocations have
1543  by respective test_<systemcall> unit testcases
1544  (*) VIRGO64 Telnet Clients to Driver listeners invocations have been tested by telnet connections
1545  (*) Master branches in SourceForge and GitHub VIRGO64 do not have KTLS provisions. Only VIRGO_KTLS branch has crypto_i
1546  for TX_TLS for kernel sockets.
1547  (*) It has been already mentioned in NeuronRain Documentation in https://neuronrain-documentation.readthedocs.io/en/la
1548  VIRGO cloud nodes in the absence of KTLS - most obvious solution is to install VPN client-servers in all nodes which c
1549  on a secure tunnel (e.g OpenVPN).
1550  (*) VIRGO64 system call clients and driver listeners should read these Virtual IPs from /etc/virgo_client.conf and /et
1551  and cloud traffic is confined to the VPN tunnel.
1552

```

```

1553 -----
1554 VIRGO64 SystemCalls-Drivers endtoend invocations unit case tests - on 4.13.3 - VIRGO64 main branch - 11 October 2017
1555 -----

```

```

1556  (*) VIRGO64 systemcalls have been invoked from unit test cases (test_<system_call>) in a loop of few hundred iteration
1557  (*) No DRM GEM i915 panics or random crashes are observed and stability is good
1558  (*) This is probably the first loop iterative testing of VIRGO system calls and drivers.
1559  (*) Kernel logs for this have been committed to virgo-docs/systemcalls_drivers directory.
1560  (*) Note on concurrency: Presently mutexing within system calls have been commented because in past linux versions mut
1561  to execute in kernel space. Mostly this is relevant only to kmemcache system calls which have global in-kernel-memory
1562

```

```

1563 -----
1564 VIRGO64 SystemCalls-Drivers concurrent invocations - 2 processes having shared mutex - 14 October 2017
1565 -----

```

```

1566  (*) VIRGO64 systemcalls are invoked in a function which is called from 2 processes concurrently
1567  (*) Mutexes between the processes are PTHREAD_PROCESS_SHARED attribute set.
1568  (*) test_virgo_malloc.c unit testcase has been enhanced to fork() a process and invoke systemcalls in a function for 1
1569  (*) Logs for the Virgo Unique IDs malloc/set/get/free in the systemcalls side and kern.logs for the drivers have been
1570  (*) No DRM GEM i915 crashes were observed
1571  (*) test_virgo_malloc.c testcase demonstrates the coarse grained critical section lock/unlock for kmemcache systemcall
1572  that should be followed for any userspace application.
1573

```

```

1574 -----
1575 VIRGO64 Kernel Analytics - Streaming Implementation - 13 December 2017
1576 -----

```

```

1577  (#) Presently kernel analytics config have to be read from a file storage. This is a huge performance bottleneck when
1578  analytics variables written to is realtime. For example autonomous vehicles/drones write gigabytes of navigation data
1579  (#) Because of this /etc/virgo_kernel_analytics.conf grows tremendously. File I/O in linux kernel module is also fragi
1580  (#) Previous latency limitations necessitate an alternative high performance analytics config variable read algorithm
1581  (#) This commit introduces new streaming kernel analytics config reading function - It connects to a kernel analytics
1582  on hardcoded port 64000 and reads analytics key-value pairs in an infinite loop.
1583  (#) These read key-value pairs are stored in a kernel global ring buffer exported symbol (by modulus operator). Becaus
1584  (#) kernel socket message flags are set to MSG_MORE | MSG_FASTOPEN | MSG_NOSIGNAL for high response time. MSG_FASTOPEN
1585  in 4.13.3 64-bit which was a problem in previous kernel versions.
1586  (#) kern.log for this has been committed to kernel_analytics/testlogs/
1587  (#) include/linux/virgo_kernel_analytics.h header file has been updated for declarations related to streaming analytic
1588  (#) Webserver used for this is netcat started on port 64000 as:

```

```

1589         nc -l 64000

```

```

1590         >k1=v1

```

```

1591         >k2=v2

```

```

1592         ...
1593

```

```

1594 -----
1595 VIRGO64 Kernel Analytics - Reading Stream of Analytic Variables made a kernel thread - 13 December 2017
1596 -----

```



```

1597  (#) This is sequel to previous commit for Stream reading Kernel Analytics variables over a network socket
1598  (#) read_streaming_virgo_kernel_analytics_conf() function is invoked in a separate kernel thread because module init i
1599  (#) VIRGO64 config module was loaded and exported kernel analytics variables read over socket by previous spin-off ker
1600  imported in VIRGO64 config init.
1601  (#) kern.log for this has been committed to testlogs/
1602  (#) Pre-requisite: Webservice serving kernel_analytics variables must be started before kernel_analytics module is loa
1603  (#) By this a minimum facility for live reading analytics anywhere on cloud (it can be userspace or kernelspace) and e
1604  to modules on a local cloud node kernel has been achieved - ideal for cloud-analytics-driven IoT
1605
1606  -----
1607  VIRGO64 System Calls - Drivers - Kernel Analytics Streaming - on 4.13.3 kernel - 15 December 2017
1608  -----
1609  (#) VIRGO64 System Calls to Drivers invocations on 4.13.3 kernel have been executed after enabling streaming kernel an
1610  (#) VIRGO64 RPC/KMemCache/CloudFS Drivers import, streamed variable-value pairs exported from kernel_analytics read fr
1611  (#) VIRGO64 KMemCache testcase has 2 concurrent processes invoking kememcache systemcalls in a loop.
1612  (#) kern.log for this has been committed to virgo-docs/systemcalls_drivers
1613  (#) virgofstest.txt written by CloudFS systemcalls-drivers invocation is also committed to virgo-docs/systemcalls_driv
1614
1615  Srinivasan Kannan (alias) Ka.Shrinivaasan (alias) Shrinivas Kannan
1616  http://sites.google.com/site/kuja27
1617

```

[About \(/about\)](#)[Create a Project](#)[Blog \(/blog/\)](#)[Articles \(/articles/\)](#)[Site Status](#)[\(/create\)](#)[@sourceforge](#)[Site Documentation](#)[\(/blog/category/sitestatus/\)](#)[Software Directory](#)[\(https://twitter.com/sourceforge\)](#)[\(https://p.sf.net/sourceforge/docs\)](#)[@sfnet_ops](#)[\(/directory/\)](#)[Resources](#)[Support Request](#)[\(https://twitter.com/sfnet_ops\)](#)[Top Downloaded](#)[\(https://library.slashdotmedia.com/support\)](#)[\(/\)](#)[Projects \(/top\)](#)

© 2018 Slashdot Media. All Rights Reserved.

[Terms \(http://slashdotmedia.com/terms-of-use\)](#)[Privacy \(http://slashdotmedia.com/privacy-statement/\)](#)[Opt Out \(http://slashdotmedia.com/opt-out-choices\)](#)[Advertise \(http://slashdotmedia.com/\)](#)